

# **CV Assignment 1**

**Vikranth Udandarao**  
**2022570**

**Theory**

$$1-a. H(p, q) = \mathbb{E}_p \left[ -\log q_{\pi}(u) \right]$$

$$H(y, q) = -\sum_{i=1}^K y_i \underbrace{\log q_i}_{\text{one hot encoded vector}} \quad \text{smooth label}$$

$$= -\log q_{\pi}$$

Label smoothing  $= (y \rightarrow \tilde{y})$  smooth label

$$\tilde{y}_{i=c} = 1 - \varepsilon + \frac{\varepsilon}{K}, \quad \begin{cases} i \neq c & \text{if } c \text{ is correct class} \\ i = c & \end{cases}$$

$$\tilde{y}_{i \neq c} = \frac{\varepsilon}{K}, \quad \text{if it is incorrect class } (i \neq c)$$

$$H(\tilde{y}, q) = -\sum_{i=1}^K \tilde{y}_i \log q_i$$

$$= -\varepsilon \cdot \frac{\varepsilon}{K} \log q_{\pi} + (1 - \varepsilon + \frac{\varepsilon}{K}) \log q_c$$

$$= -\left( \varepsilon \sum_{i \neq c} \frac{\varepsilon}{K} \log q_{\pi} + \log q_{\pi} - \varepsilon \log q_c \right)$$

$$+ \frac{\varepsilon}{K} \log q_c \right)_{i=c}$$

$$= -\left( \log q_{\pi} - \varepsilon \log q_c + \sum_{i=1}^K \frac{\varepsilon}{K} \log q_{\pi} \right)$$

$$= -\left( (1 - \varepsilon) \log q_c + \frac{\varepsilon}{K} \sum_{i=1}^K \log q_{\pi} \right)$$

$$= -\left( \log q_c \right) (1 - \varepsilon) - \frac{\varepsilon}{K} \sum_{i=1}^K \log q_{\pi}$$

$$= - (1-\varepsilon) (-H(y, p)) - \frac{\varepsilon}{K} \sum_{i=1}^K \log q_i -$$

$\downarrow$  one hot encoded

$$= \varepsilon (1-\varepsilon) H(y, p) - \frac{\varepsilon}{K} \sum_{i=1}^K \log q_i -$$

b) Label smoothing: Modifies the loss by introducing a factor of  $(1-\varepsilon)$  which balances the true labels and encourages consecutive predictions. This regularization technique.

This regularization discourages overfitting in predictions by preventing the model from assuming probabilities too close to 0 or 1. It is reducing the overconfidence in predictions indirectly.

It helps in reducing overfitting, the loss becomes softer thus giving more stable gradients / avoiding gradient explosion / vanishing and also gives better ambiguity or noise in labels, promoting smoother decision boundaries.

$$2- P(u) = N(M_p, \sigma_p^2), q(u) = N(M_q, \sigma_q^2)$$

a) we know  $M(p, q) = \mathbb{E}_{p(u)} [\log q(u)]$

from  $\bullet$  Q1

$$q(u) = \frac{1}{\sqrt{2\pi}\sigma_q} e^{-\left(\frac{1}{2} \frac{(u-u)^2}{\sigma_q^2}\right)}$$

log on both sides →

$$\log(q(u)) = -\frac{1}{2} \log 2\pi \sigma_q^2 - \frac{1}{2} \frac{(u-u)^2}{\sigma_q^2}$$

Substitute in ①,

$$H(p, q) = \frac{1}{2} \log (2\pi \sigma_q^2) + \frac{1}{2\sigma_q^2} E_p(u) [u - u_q]^2$$

b- from ②  $\rightarrow H(p, q) = \dots - \frac{1}{2\sigma_q^2} E_p(u) [u - u_q]^2$

$$E_{p(u)} [(u - u_q)^2] = E_{p(u)} [(u - u_p + u_p - u_q)^2]$$

$$= E_{p(u)} [(u - u_p) + (u_p - u_q)]^2$$

$$= E_{p(u)} [(u - u_p)^2] + E_{p(u)} [(u_p - u_q)^2]$$

$$+ E_{p(u)} [(u_p - u_q)(u - u_p)]$$

$$= \sigma_p^2 + (u_p - u_q)^2 + 2(u_p - u_q)$$

$$\neq 0 \quad (u \neq u_p)$$

$$= \sigma_p^2 + (u_p - u_q)^2$$

$$\therefore H(p, q) = \frac{1}{2} \log (2\pi \sigma^2) + \frac{1}{2\sigma^2} (\sigma_p^2 + (u_p - u_q)^2)$$

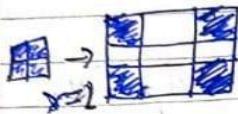
$$\leftarrow \sigma_p = \sigma_q = \sigma$$

$$\rightarrow H(p, q) = \frac{1}{2} \log (2\pi \sigma^2) + \frac{\sigma^2 + (u_p - u_q)^2}{2\sigma^2}$$

$$= \frac{1}{2} \log(2\pi\sigma^2) + \frac{1}{2} + \left(\frac{\mu_p - \mu_q}{\sigma}\right)^2$$

$$= \frac{1}{2} \left[ (\log(2\pi\sigma^2) + 1) + \left(\frac{\mu_p - \mu_q}{\sigma}\right)^2 \right] -$$

3-a- In strong convolution  $\alpha = 1 \rightarrow$  normal convolution,  
 $\alpha = 2 \rightarrow$  elements are spaced 1 pixel apart  
 and so on.



, effective kernel size (kernel size)

$$\rightarrow (\text{kernel size} - 1) \alpha + 1$$

$$= (K-1) \alpha + 1$$

when using 1 such layer -

$$\text{receptive field} \rightarrow R_C = \sum_{i=1}^{L-1} (K-1)r_i + 1$$

It gives that the dilation rate  $\alpha$  increases exponentially which means  $r_i = 2^{i-1}$ .

(geometric

mean)

$$\rightarrow R_C = \sum_{i=1}^{L-1} (K-1) 2^{i-1} + 1$$

$$= 1 + (K-1)(2^{L-1})$$

Since, it's a function of  $L$  of  $\propto n^{2^L}$ ,

we can say that the receptive field grows exponentially.

b-  $R = 1 + (K-1) (2^L - 1)$

$R_{\text{area}} = R_{\text{height}} \times R_{\text{width}}$   
 $= [1 + (K-1) (2^L - 1)]^2$   
(Since  $H=W$ )

$K \times K \rightarrow \text{square}$

now  $R_{\text{area}}$  changes as  $f(l) \approx 4^L$ , the area of receptive field grows quadratically exponentially.

c- for standard  $K \times K$  convolution-

~~Let~~ Let Input =  $m \times n$  (ignoring channels for now)

No. of equations =  $m \times n + k^2$

for dilated  $K \times K$  convolution-

Let Input =  $m \times n$

No. of equations =  $m \times n + k^2$  (we don't include holes)

→ We observe that dilation of the kernel in std convolution requires the same computation power as normal convolution but it covers a larger receptive field thus giving better global context for future iterations.

# Coding

2. 1. a.

```
# Custom Dataset Class

class WildlifeDataset(Dataset):
    def __init__(self, img_dir, class_mapping, transform=None, train=False):
        self.img_dir = img_dir
        self.transform = transform
        self.class_mapping = class_mapping
        self.train = train # Flag to determine if this is training set
        self.img_labels = []

        # # Data augmentation for training
        # self.train_transform = transforms.Compose([
        #     transforms.ToPILImage(),
        #     transforms.RandomHorizontalFlip(p=0.5), # Randomly flip image horizontally
        #     transforms.RandomRotation(degrees=15), # Randomly rotate image up to 15 degrees
        #     transforms.ColorJitter(brightness=0.2, contrast=0.2, saturation=0.2), # Adjust color properties
        #     transforms.Resize((224, 224)),
        #     transforms.ToTensor(),
        #     transforms.Normalize(mean=[0.485, 0.456, 0.406],
        #                         std=[0.229, 0.224, 0.225])
        # ])

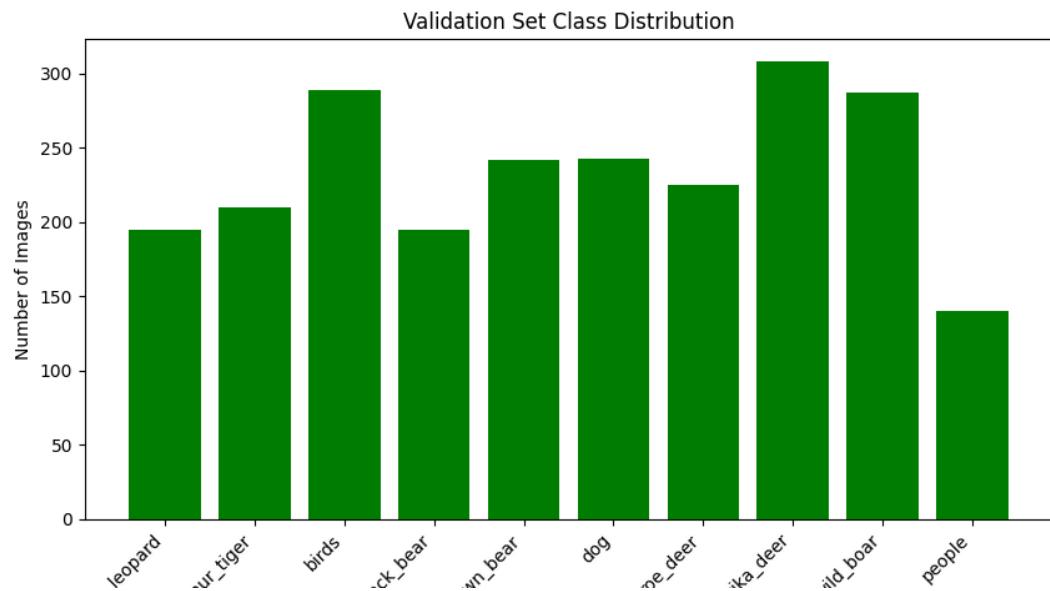
        # Transform for validation (no augmentation)
        self.val_transform = transforms.Compose([
            transforms.ToPILImage(),
            transforms.Resize((224, 224)),
            transforms.ToTensor(),
            transforms.Normalize(mean=[0.485, 0.456, 0.406],
                                std=[0.229, 0.224, 0.225])
        ])

    # Collect image paths and labels
    for class_name in os.listdir(img_dir):
        class_path = os.path.join(img_dir, class_name)
        if os.path.isdir(class_path):
            for img_name in os.listdir(class_path):
                img_path = os.path.join(class_path, img_name)
                self.img_labels.append((img_path, class_mapping[class_name]))
```

b. None

```
# Create dataloaders
batch_size = 64
train_dataloader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
val_dataloader = DataLoader(val_dataset, batch_size=batch_size, shuffle=False)
```

c. None



2. a.

```
class ConvNet(nn.Module):
    def __init__(self, num_classes=10):
        super(ConvNet, self).__init__()

        # First Convolutional Block
        self.conv1 = nn.Conv2d(
            in_channels=3,           # RGB input
            out_channels=32,         # 32 feature maps
            kernel_size=3,           # 3x3 kernel
            stride=1,
            padding=1
        )
        self.relu1 = nn.ReLU()
        self.pool1 = nn.MaxPool2d(kernel_size=4, stride=4)

        # Second Convolutional Block
        self.conv2 = nn.Conv2d(
            in_channels=32,          # Input from previous layer
            out_channels=64,         # 64 feature maps
            kernel_size=3,
            stride=1,
            padding=1
        )
        self.relu2 = nn.ReLU()
        self.pool2 = nn.MaxPool2d(kernel_size=2, stride=2)

        # Third Convolutional Block
        self.conv3 = nn.Conv2d(
            in_channels=64,          # Input from previous layer
            out_channels=128,         # 128 feature maps
            kernel_size=3,
            stride=1,
            padding=1
        )
        self.relu3 = nn.ReLU()
        self.pool3 = nn.MaxPool2d(kernel_size=2, stride=2)

        # Classification head
        self.fc1 = nn.Linear(128 * 14 * 14, 256) # Adjust size based on pooling
        self.relu4 = nn.ReLU()
        self.fc2 = nn.Linear(256, num_classes) # Output layer
```

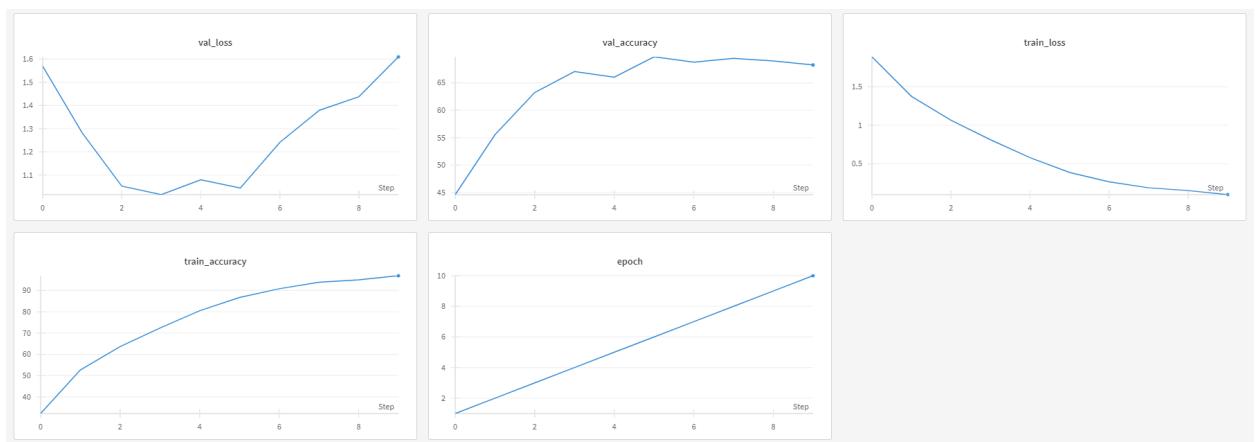
b.

The screenshot shows the WandB interface with the following details:

- Overview:** Shows a tree view of the project structure.
- Misclassified Examples:** A section showing images of misclassified examples.
- Charts:** Five line charts showing metrics over 10 steps.
  - val\_loss:** Training loss decreasing from ~1.6 to ~1.05.
  - val\_accuracy:** Validation accuracy increasing from ~45% to ~68%.
  - train\_loss:** Training loss decreasing from ~1.6 to ~0.4.
  - train\_accuracy:** Training accuracy increasing from ~40% to ~92%.
  - epoch:** Epoch number increasing from 0 to 10.
- Windows PowerShell:** Terminal output showing the training process and metrics for 10 epochs.

```
wandb: Syncing run convnet-baseline
wandb: View project at https://wandb.ai/vikranth2764-na/russian-wildlife-classification
wandb: View run at https://wandb.ai/vikranth2764-na/russian-wildlife-classification/runs/2uc5g73h
Epoch [1/10]
Train Loss: 1.8853, Train Acc: 32.22%
Val Loss: 1.5697, Val Acc: 44.60%
Epoch [2/10]
Train Loss: 1.3730, Train Acc: 52.68%
Val Loss: 1.2820, Val Acc: 55.53%
Epoch [3/10]
Train Loss: 1.0632, Train Acc: 63.64%
Val Loss: 1.0526, Val Acc: 63.24%
Epoch [4/10]
Train Loss: 0.8122, Train Acc: 72.35%
Val Loss: 1.0156, Val Acc: 67.05%
Epoch [5/10]
Train Loss: 0.5802, Train Acc: 80.48%
Val Loss: 1.0797, Val Acc: 66.02%
Epoch [6/10]
Train Loss: 0.3887, Train Acc: 86.73%
Val Loss: 1.0443, Val Acc: 69.75%
Epoch [7/10]
Train Loss: 0.2679, Train Acc: 90.82%
Val Loss: 1.2410, Val Acc: 68.77%
Epoch [8/10]
Train Loss: 0.1899, Train Acc: 93.88%
Val Loss: 1.3793, Val Acc: 69.45%
Epoch [9/10]
Train Loss: 0.1541, Train Acc: 95.00%
Val Loss: 1.4379, Val Acc: 68.98%
Epoch [10/10]
Train Loss: 0.1014, Train Acc: 96.96%
Val Loss: 1.6102, Val Acc: 68.25%
Training finished!
```

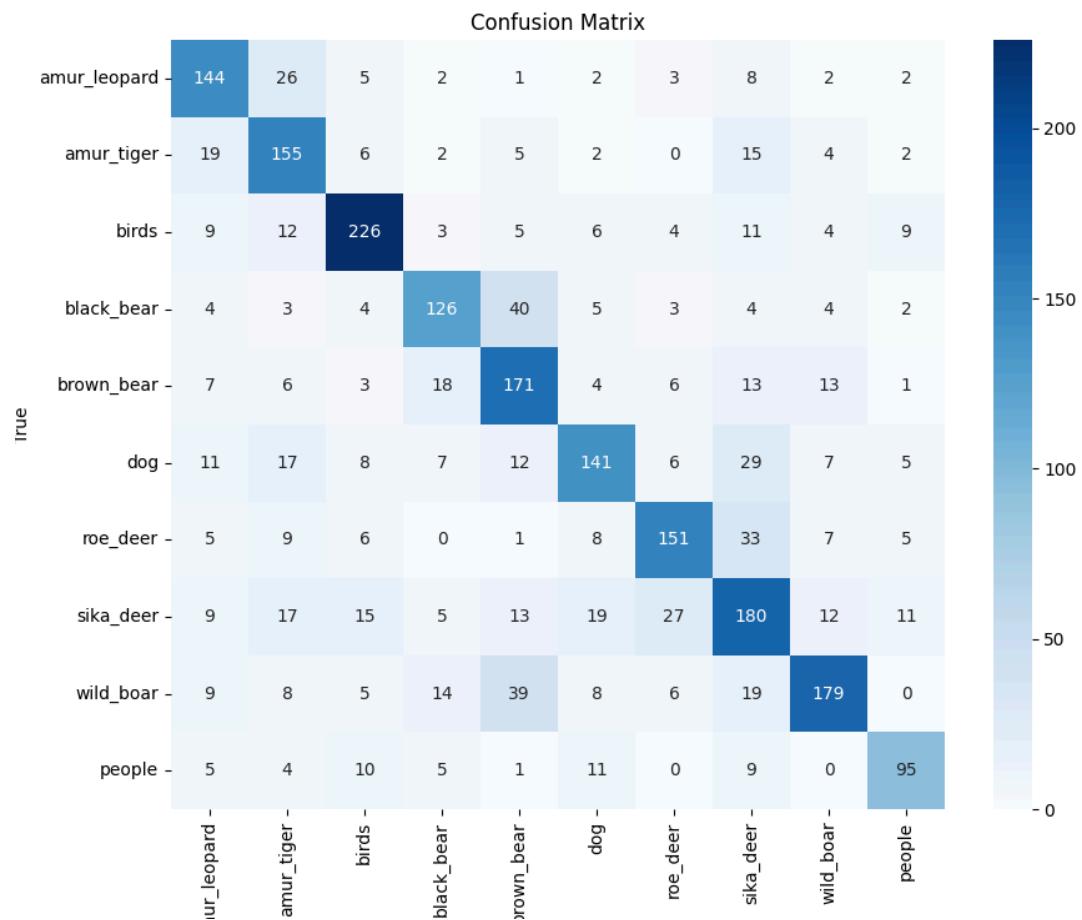
C.



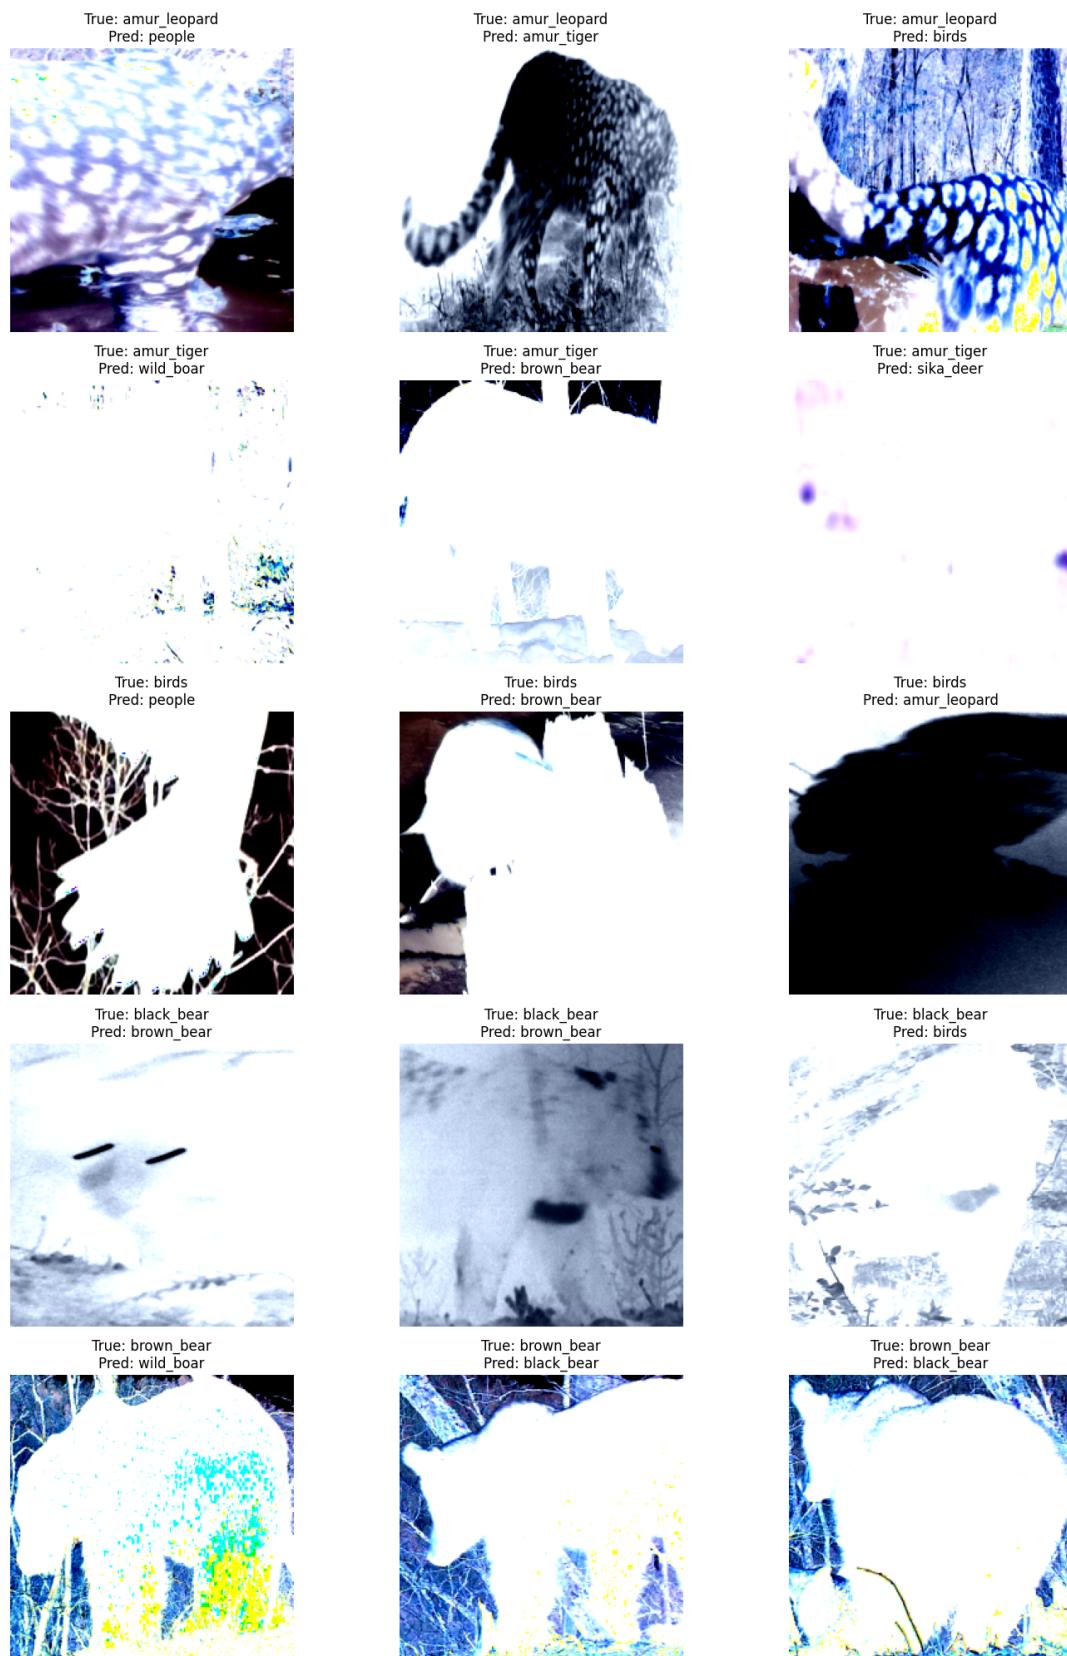
Yes model is overfitting as Train Accuracy is 96.96% whereas Validation Accuracy is only 68.25%.

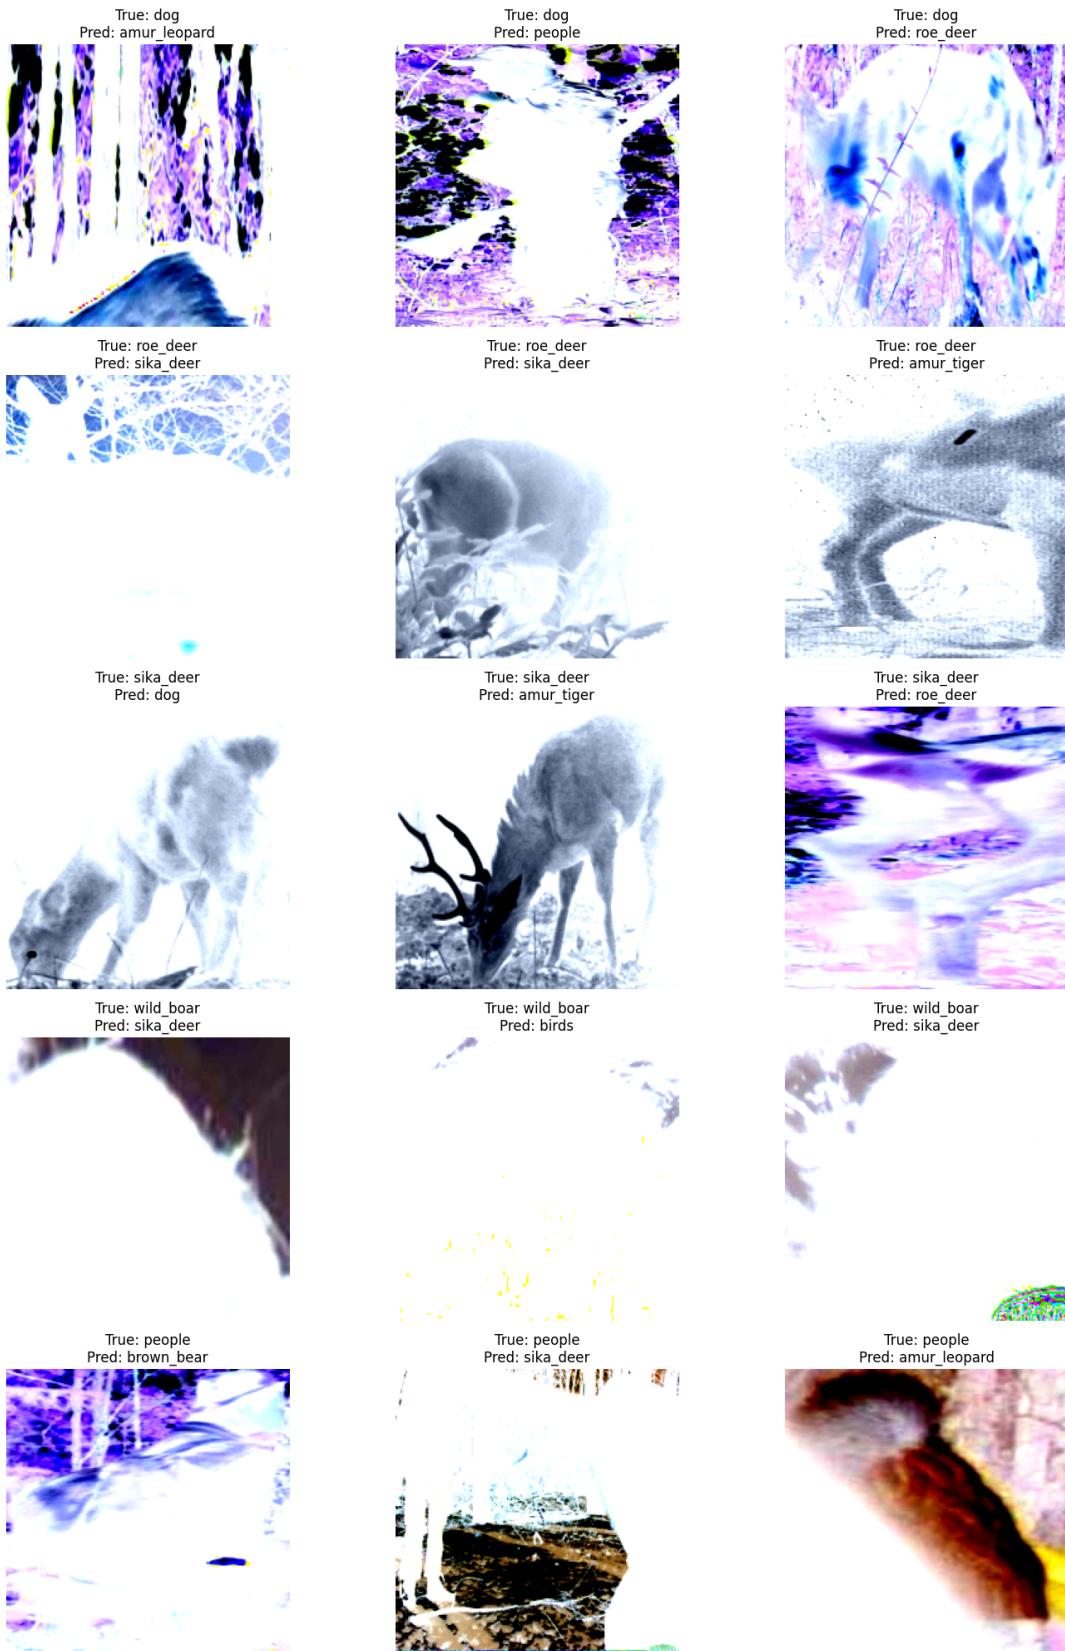
**Validation Metrics:**  
**Accuracy: 0.6718**  
**F1-Score: 0.6724**

d.



e.





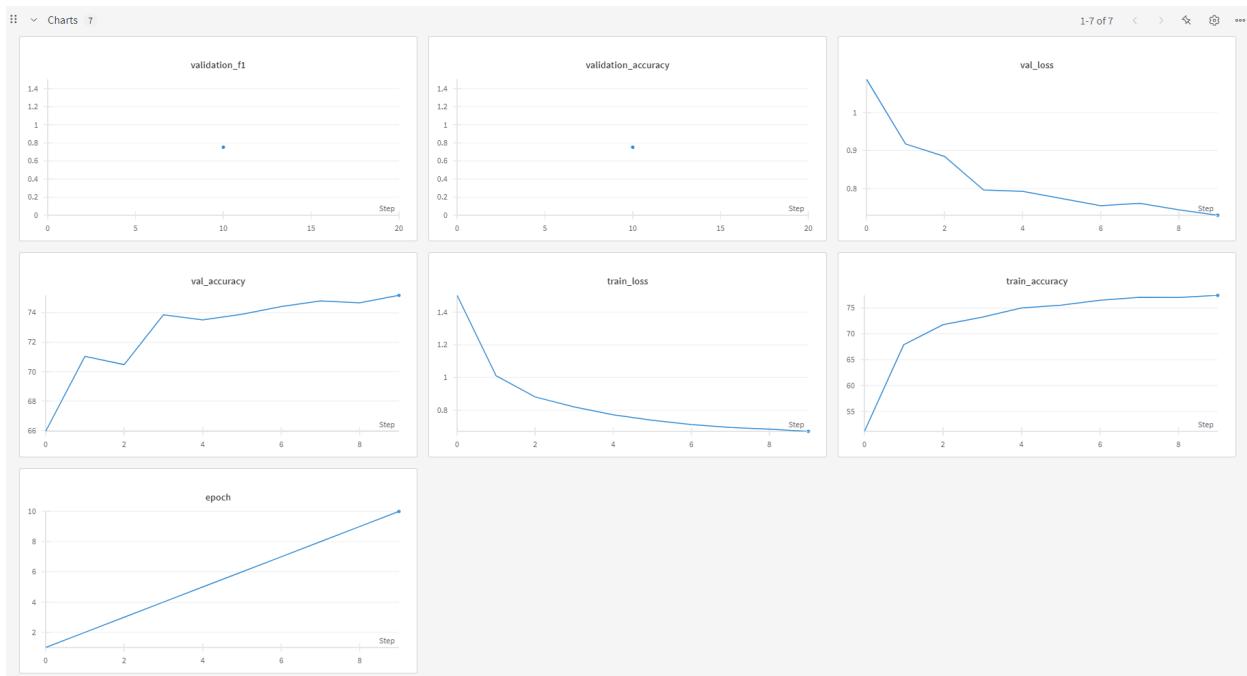
**Visual Similarity:** Many misclassified images belong to visually similar categories (e.g., "amur\_leopard" vs. "amur\_tiger" and "black\_bear" vs. "brown\_bear").

**Partial Occlusion:** In certain cases, only a part of the object is visible, leading the model to misinterpret it.

**Data Augmentation:** We can try to introduce transformations such as contrast adjustments, brightness corrections, and different lighting conditions.

**Feature Engineering:** Extract more discriminative features like texture patterns specific to each class

### 3. a.



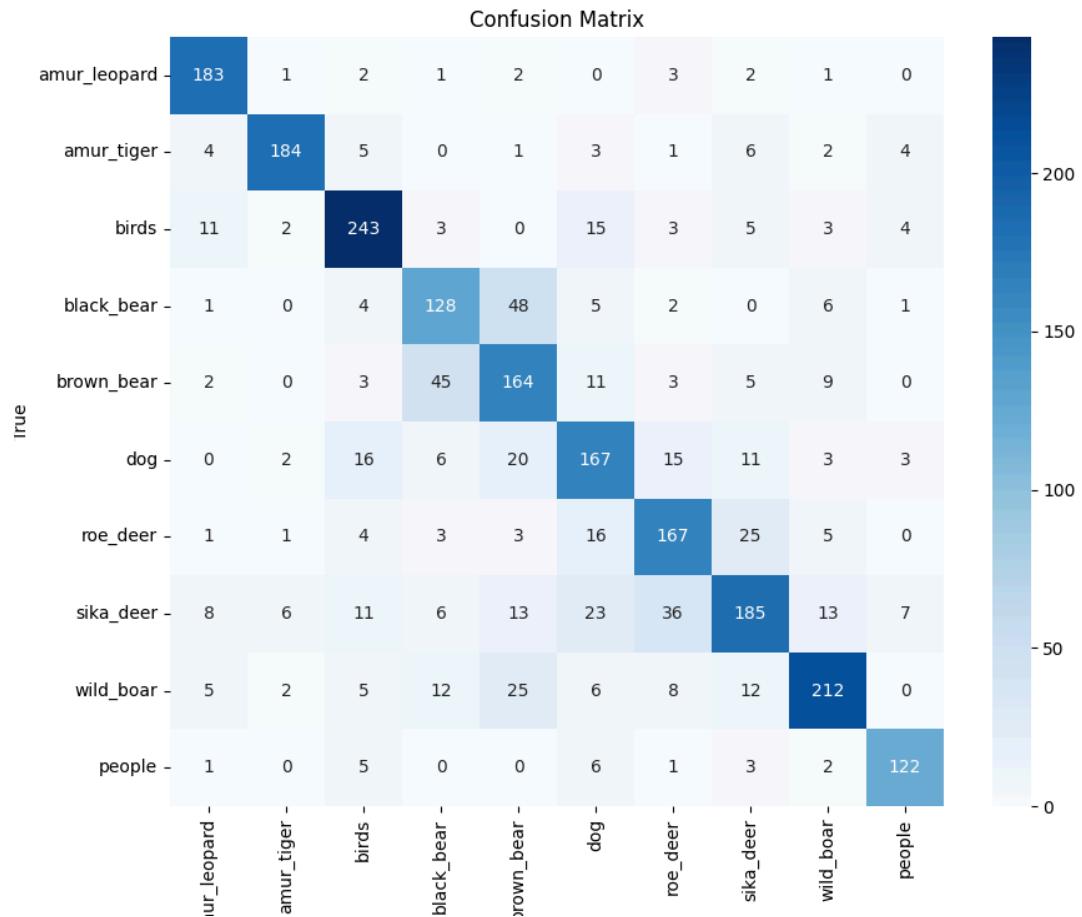
b.

```
Epoch [1/10]
Train Loss: 1.5037, Train Acc: 51.12%
Val Loss: 1.0882, Val Acc: 65.94%
Epoch [2/10]
Train Loss: 1.0110, Train Acc: 67.88%
Val Loss: 0.9175, Val Acc: 71.04%
Epoch [3/10]
Train Loss: 0.8819, Train Acc: 71.77%
Val Loss: 0.8846, Val Acc: 70.48%
Epoch [4/10]
Train Loss: 0.8208, Train Acc: 73.22%
Val Loss: 0.7960, Val Acc: 73.86%
Epoch [5/10]
Train Loss: 0.7730, Train Acc: 75.02%
Val Loss: 0.7925, Val Acc: 73.52%
Epoch [6/10]
Train Loss: 0.7394, Train Acc: 75.54%
Val Loss: 0.7737, Val Acc: 73.91%
Epoch [7/10]
Train Loss: 0.7131, Train Acc: 76.49%
Val Loss: 0.7549, Val Acc: 74.42%
Epoch [8/10]
Train Loss: 0.6949, Train Acc: 77.07%
Val Loss: 0.7608, Val Acc: 74.81%
Epoch [9/10]
Train Loss: 0.6851, Train Acc: 77.03%
Val Loss: 0.7439, Val Acc: 74.68%
Epoch [10/10]
Train Loss: 0.6713, Train Acc: 77.46%
Val Loss: 0.7295, Val Acc: 75.19%
```

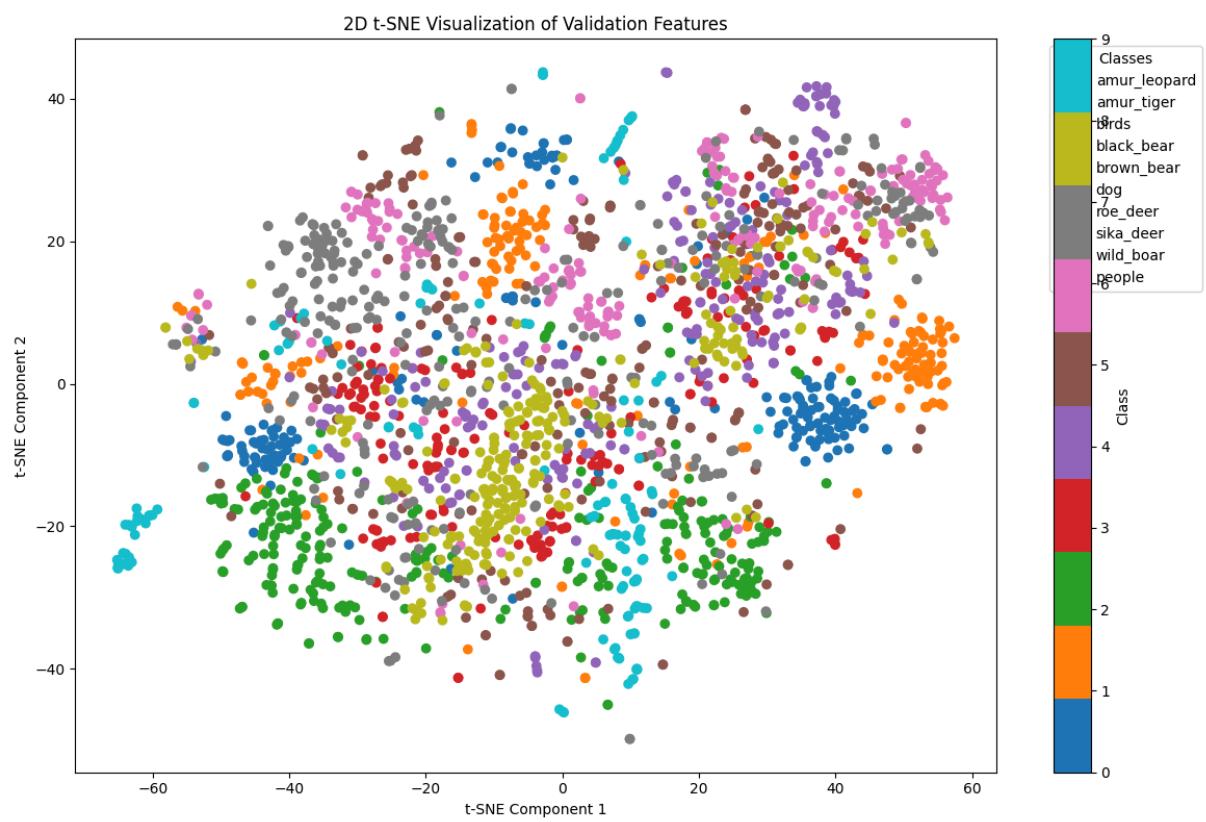
No, the model is not overfitting as training ad validation accuracy is almost similar with not much difference between them and training accuracy is not very very higher compared to validation accuracy.

```
Validation Metrics:  
Accuracy: 0.7519  
F1-Score: 0.7520  
Training finished!
```

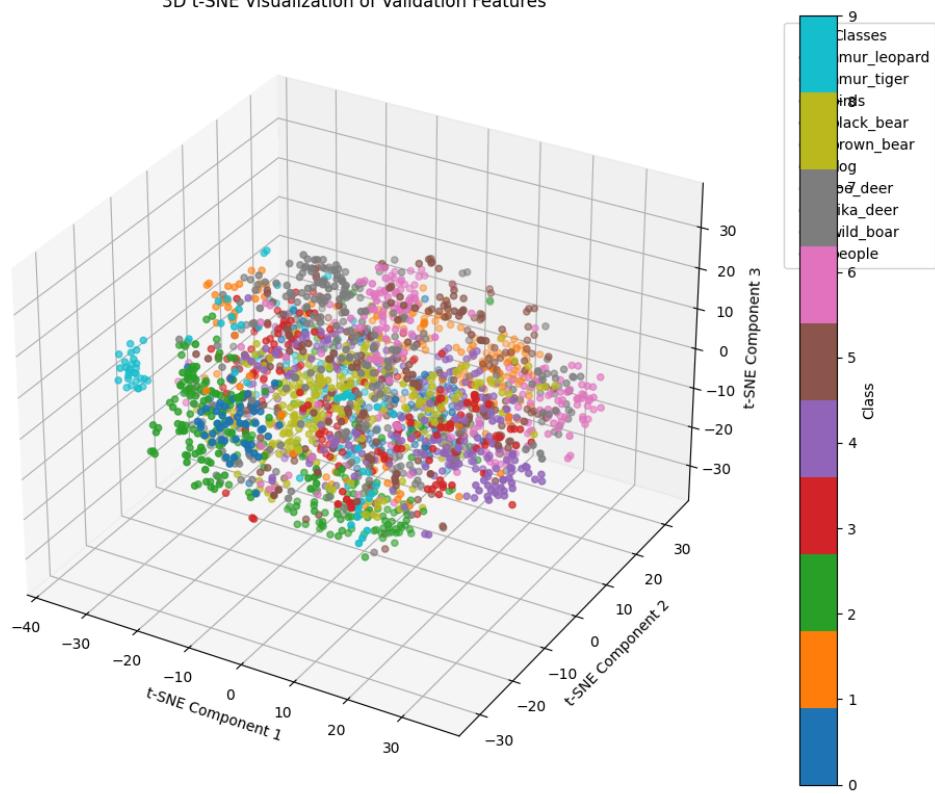
C.



d.

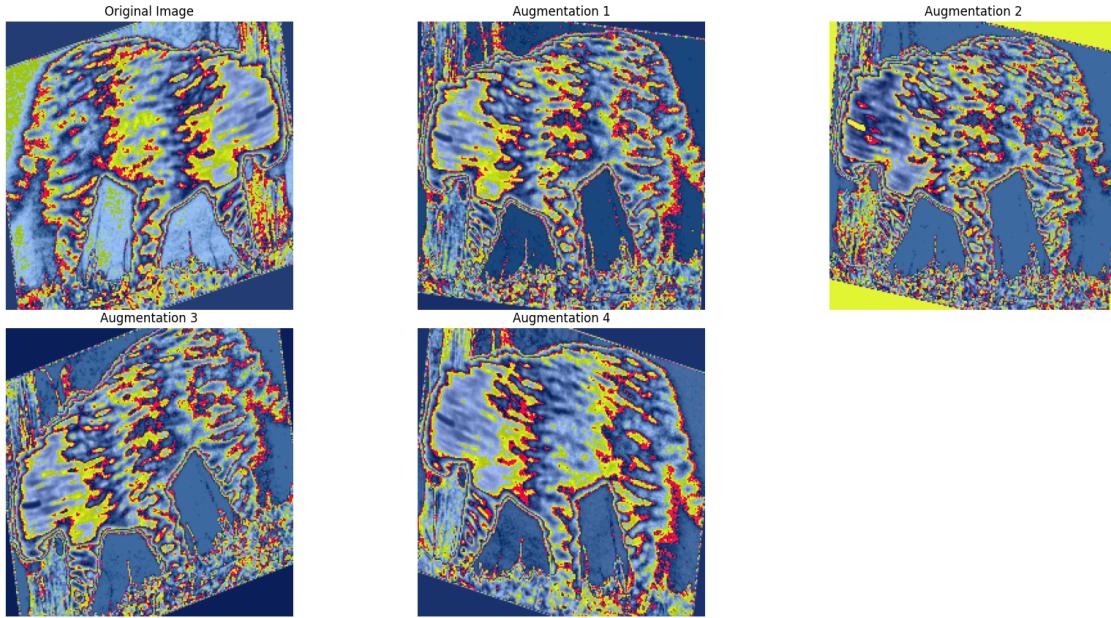


3D t-SNE Visualization of Validation Features

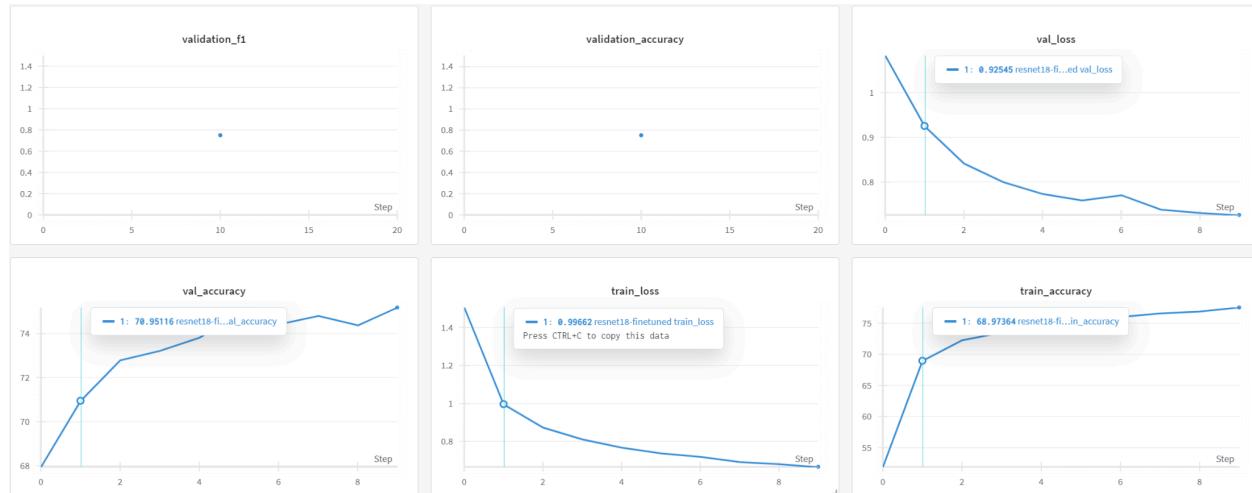


#### 4. a.

```
# Use this for data augmentation for resnet augmentation
transform = transforms.Compose([
    transforms.ToPILImage(),
    transforms.RandomHorizontalFlip(p=0.5), # Flip images horizontally with 50% probability
    transforms.RandomRotation(degrees=15), # Rotate images randomly within ±15 degrees
    transforms.ColorJitter(brightness=0.2, contrast=0.2, saturation=0.2, hue=0.1), # can adjust brightness, contrast, saturation but as of now I putt 0.2 for all
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406],
                        std=[0.229, 0.224, 0.225]) # ImageNet normalization
])
```



#### b.



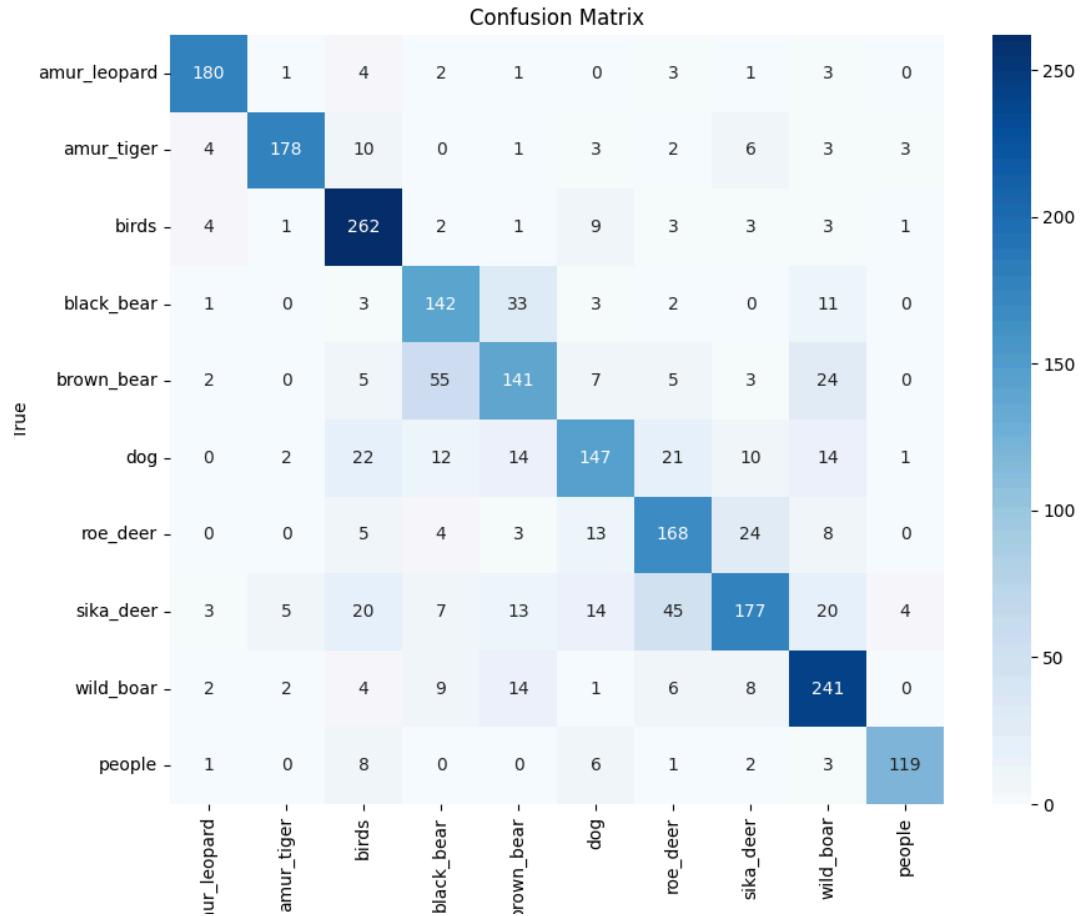
c.

```
Epoch [1/10]
Train Loss: 1.5065, Train Acc: 51.92%
Val Loss: 1.0830, Val Acc: 67.95%
Epoch [2/10]
Train Loss: 0.9966, Train Acc: 68.97%
Val Loss: 0.9254, Val Acc: 70.95%
Epoch [3/10]
Train Loss: 0.8748, Train Acc: 72.28%
Val Loss: 0.8420, Val Acc: 72.79%
Epoch [4/10]
Train Loss: 0.8127, Train Acc: 73.52%
Val Loss: 0.8003, Val Acc: 73.22%
Epoch [5/10]
Train Loss: 0.7690, Train Acc: 74.56%
Val Loss: 0.7740, Val Acc: 73.82%
Epoch [6/10]
Train Loss: 0.7391, Train Acc: 75.48%
Val Loss: 0.7596, Val Acc: 74.85%
Epoch [7/10]
Train Loss: 0.7209, Train Acc: 76.02%
Val Loss: 0.7711, Val Acc: 74.42%
Epoch [8/10]
Train Loss: 0.6935, Train Acc: 76.59%
Val Loss: 0.7389, Val Acc: 74.81%
Epoch [9/10]
Train Loss: 0.6825, Train Acc: 76.89%
Val Loss: 0.7317, Val Acc: 74.38%
Epoch [10/10]
Train Loss: 0.6663, Train Acc: 77.56%
Val Loss: 0.7262, Val Acc: 75.19%
```

Yes, now the model is not overfitting, there is not high variance like how Convnet had.

**Validation Metrics:**  
**Accuracy: 0.7519**  
**F1-Score: 0.7497**  
**Training finished!**

d.



5. Convnet model is highly overfitting and has high variance and low bias and only has 68% validation accuracy.

Resnet model does not overfit and has low variance and both training and validation accuracy is around 75% accuracy.

Similarly for resnet model with data augmentation techniques, I did not get much different results from resnet.

Maybe this is because of the data augmentation techniques I am doing as mentioned in 4.a. I am using 0.2 for brightness, contrast, saturation and only rotating the images by 15°.

### 3. 1. a.

```

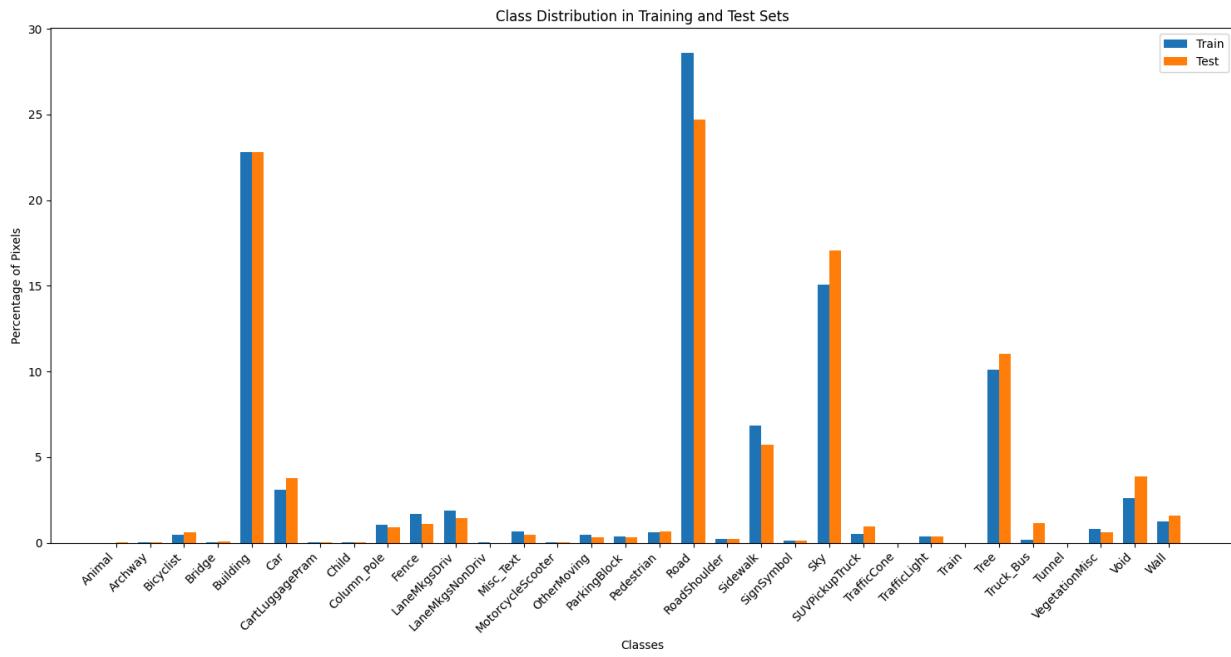
class SegmentationDataset(Dataset):
    def __init__(self, data_dir, split='train', transform=None):
        self.data_dir = data_dir
        self.split = split
        self.transform = transform
        self.class_mapping = CLASS_MAPPING

        # Get all image and mask files based on split
        if split == 'train':
            self.images_dir = os.path.join(data_dir, 'train')
            self.masks_dir = os.path.join(data_dir, 'train_labels')
        else: # test
            self.images_dir = os.path.join(data_dir, 'test_images')
            self.masks_dir = os.path.join(data_dir, 'test_labels')

        self.image_files = sorted([f for f in os.listdir(self.images_dir)])
        self.mask_files = sorted([f for f in os.listdir(self.masks_dir)])

        # Transformation according to what is said in 1.a. for both image and mask
        self.resize = transforms.Resize((360, 480), interpolation=transforms.InterpolationMode.BILINEAR)
        self.resize_mask = transforms.Resize((360, 480), interpolation=transforms.InterpolationMode.NEAREST)
        self.normalize = transforms.Normalize(
            mean=[0.485, 0.456, 0.406],
            std=[0.229, 0.224, 0.225]
        )
    
```

### b.



C.  
Car:

Class: Car



Pedestrian:

Class: Pedestrian



Road:

Class: Road



## 2. a.



b.

Test Set Evaluation Results:			
=====			
Mean IoU (mIoU): 0.1356			
Class-wise Results:			
Class	Pixel Acc	Dice	IoU
Animal	0.9999	0.0000	0.0000
Archway	0.9998	0.0000	0.0000
Bicyclist	0.9934	0.0854	0.0479
Bridge	0.9994	0.0000	0.0000
Building	0.8693	0.7306	0.6057
Car	0.9789	0.6772	0.5418
CartLuggagePram	0.9997	0.0000	0.0000
Child	0.9997	0.0000	0.0000
Column_Pole	0.9899	0.0606	0.0321
Fence	0.9816	0.1496	0.0959
LaneMkgsDriv	0.9868	0.1705	0.0999
LaneMkgsNonDriv	1.0000	0.0000	0.0000
Misc_Text	0.9953	0.0000	0.0000
MotorcycleScooter	0.9998	0.0000	0.0000
OtherMoving	0.9970	0.0070	0.0036
ParkingBlock	0.9965	0.0000	0.0000
Pedestrian	0.9925	0.0342	0.0176
Road	0.9580	0.9081	0.8358
RoadShoulder	0.9979	0.0000	0.0000
Sidewalk	0.9684	0.6980	0.5642
SignSymbol	0.9990	0.0000	0.0000
Sky	0.9532	0.8451	0.7524
SUVPickupTruck	0.9904	0.0000	0.0000
TrafficCone	1.0000	0.0000	0.0000
TrafficLight	0.9963	0.0000	0.0000
Train	1.0000	0.0000	0.0000
Tree	0.9516	0.6265	0.5007
Truck_Bus	0.9885	0.0000	0.0000
Tunnel	1.0000	0.0000	0.0000
VegetationMisc	0.9941	0.0017	0.0008
Void	0.9525	0.1184	0.0634
Wall	0.9854	0.2719	0.1777

Precision and Recall at different IoU thresholds:

Class	Threshold	Precision	Recall
Animal	0.0	0.0000	0.0000
	0.1	0.0000	0.0000
	0.2	0.0000	0.0000
	0.3	0.0000	0.0000
	0.4	0.0000	0.0000
	0.5	0.0000	0.0000
	0.6	0.0000	0.0000
	0.7	0.0000	0.0000
	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
Archway	1.0	0.0000	0.0000
	0.0	0.0000	0.0000
	0.1	0.0000	0.0000
	0.2	0.0000	0.0000
	0.3	0.0000	0.0000
	0.4	0.0000	0.0000
	0.5	0.0000	0.0000
	0.6	0.0000	0.0000
	0.7	0.0000	0.0000
	0.8	0.0000	0.0000
Bicyclist	0.9	0.0000	0.0000
	1.0	0.0000	0.0000
	0.0	0.2387	0.0597
	0.1	0.1419	0.0377
	0.2	0.0000	0.0000
	0.3	0.0000	0.0000
	0.4	0.0000	0.0000
	0.5	0.0000	0.0000
	0.6	0.0000	0.0000
	0.7	0.0000	0.0000
Bridge	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000
	0.0	0.0000	0.0000
	0.1	0.0000	0.0000
	0.2	0.0000	0.0000
	0.3	0.0000	0.0000
	0.4	0.0000	0.0000
	0.5	0.0000	0.0000
	0.6	0.0000	0.0000
Building	0.7	0.0000	0.0000
	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000
	0.0	0.6196	0.9521
	0.1	0.6196	0.9521
	0.2	0.6085	0.8919
	0.3	0.5996	0.8626
	0.4	0.5757	0.7984
	0.5	0.5456	0.7325

Car	0.0	0.6526	0.7394
	0.1	0.6526	0.7394
	0.2	0.6469	0.7248
	0.3	0.6263	0.6799
	0.4	0.5999	0.6230
	0.5	0.5373	0.5509
	0.6	0.3643	0.3788
	0.7	0.1733	0.1839
	0.8	0.0319	0.0319
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000
CartLuggagePram	0.0	0.0000	0.0000
	0.1	0.0000	0.0000
	0.2	0.0000	0.0000
	0.3	0.0000	0.0000
	0.4	0.0000	0.0000
	0.5	0.0000	0.0000
	0.6	0.0000	0.0000
	0.7	0.0000	0.0000
	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000
Child	0.0	0.0000	0.0000
	0.1	0.0000	0.0000
	0.2	0.0000	0.0000
	0.3	0.0000	0.0000
	0.4	0.0000	0.0000
	0.5	0.0000	0.0000
	0.6	0.0000	0.0000
	0.7	0.0000	0.0000
	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000
Column_Pole	0.0	0.1869	0.0379
	0.1	0.0145	0.0049
	0.2	0.0000	0.0000
	0.3	0.0000	0.0000
	0.4	0.0000	0.0000
	0.5	0.0000	0.0000
	0.6	0.0000	0.0000
	0.7	0.0000	0.0000
	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000
Fence	0.0	0.2040	0.1412
	0.1	0.1674	0.1222
	0.2	0.1033	0.1066
	0.3	0.0612	0.0652
	0.4	0.0612	0.0652
	0.5	0.0216	0.0249
	0.6	0.0000	0.0000
	0.7	0.0000	0.0000
	0.8	0.0000	0.0000

Fence	0.0	0.2040	0.1412
	0.1	0.1674	0.1222
	0.2	0.1033	0.1066
	0.3	0.0612	0.0652
	0.4	0.0612	0.0652
	0.5	0.0216	0.0249
	0.6	0.0000	0.0000
	0.7	0.0000	0.0000
	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000
LaneMkgsDriv	0.0	0.6487	0.1031
	0.1	0.3927	0.0844
	0.2	0.1227	0.0392
	0.3	0.0283	0.0112
	0.4	0.0000	0.0000
	0.5	0.0000	0.0000
	0.6	0.0000	0.0000
	0.7	0.0000	0.0000
	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000
LaneMkgsNonDriv	0.0	0.0000	0.0000
	0.1	0.0000	0.0000
	0.2	0.0000	0.0000
	0.3	0.0000	0.0000
	0.4	0.0000	0.0000
	0.5	0.0000	0.0000
	0.6	0.0000	0.0000
	0.7	0.0000	0.0000
	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000
Misc_Text	0.0	0.0000	0.0000
	0.1	0.0000	0.0000
	0.2	0.0000	0.0000
	0.3	0.0000	0.0000
	0.4	0.0000	0.0000
	0.5	0.0000	0.0000
	0.6	0.0000	0.0000
	0.7	0.0000	0.0000
	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000
MotorcycleScooter	0.0	0.0000	0.0000
	0.1	0.0000	0.0000
	0.2	0.0000	0.0000
	0.3	0.0000	0.0000
	0.4	0.0000	0.0000
	0.5	0.0000	0.0000
	0.6	0.0000	0.0000
	0.7	0.0000	0.0000
	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000
OtherMoving	0.0	0.1164	0.0037
	0.1	0.0000	0.0000
	0.2	0.0000	0.0000
	0.3	0.0000	0.0000
	0.4	0.0000	0.0000
	0.5	0.0000	0.0000
	0.6	0.0000	0.0000
	0.7	0.0000	0.0000
	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000

OtherMoving	0.0	0.1164	0.0037
	0.1	0.0000	0.0000
	0.2	0.0000	0.0000
	0.3	0.0000	0.0000
	0.4	0.0000	0.0000
	0.5	0.0000	0.0000
	0.6	0.0000	0.0000
	0.7	0.0000	0.0000
	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000
ParkingBlock	0.0	0.0000	0.0000
	0.1	0.0000	0.0000
	0.2	0.0000	0.0000
	0.3	0.0000	0.0000
	0.4	0.0000	0.0000
	0.5	0.0000	0.0000
	0.6	0.0000	0.0000
	0.7	0.0000	0.0000
	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000
Pedestrian	0.0	0.1527	0.0214
	0.1	0.0000	0.0000
	0.2	0.0000	0.0000
	0.3	0.0000	0.0000
	0.4	0.0000	0.0000
	0.5	0.0000	0.0000
	0.6	0.0000	0.0000
	0.7	0.0000	0.0000
	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000
Road	0.0	0.8687	0.9531
	0.1	0.8687	0.9531
	0.2	0.8687	0.9531
	0.3	0.8687	0.9531
	0.4	0.8687	0.9531
	0.5	0.8687	0.9531
	0.6	0.8480	0.9239
	0.7	0.8214	0.8941
	0.8	0.6522	0.7028
	0.9	0.2592	0.2692
	1.0	0.0000	0.0000
RoadShoulder	0.0	0.0000	0.0000
	0.1	0.0000	0.0000
	0.2	0.0000	0.0000
	0.3	0.0000	0.0000
	0.4	0.0000	0.0000
	0.5	0.0000	0.0000
	0.6	0.0000	0.0000
	0.7	0.0000	0.0000
	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000
Sidewalk	0.0	0.6204	0.8399
	0.1	0.6204	0.8399
	0.2	0.6138	0.8254
	0.3	0.6059	0.8011
	0.4	0.5350	0.6730
	0.5	0.4105	0.4763
	0.6	0.3879	0.4464
	0.7	0.2653	0.2876
	0.8	0.0923	0.1000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000

Sidewalk	0.0	0.6204	0.8399
	0.1	0.6204	0.8399
	0.2	0.6138	0.8254
	0.3	0.6059	0.8011
	0.4	0.5350	0.6730
	0.5	0.4105	0.4763
	0.6	0.3879	0.4464
	0.7	0.2653	0.2876
	0.8	0.0923	0.1000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000
SignSymbol	0.0	0.0000	0.0000
	0.1	0.0000	0.0000
	0.2	0.0000	0.0000
	0.3	0.0000	0.0000
	0.4	0.0000	0.0000
	0.5	0.0000	0.0000
	0.6	0.0000	0.0000
	0.7	0.0000	0.0000
	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000
Sky	0.0	0.9295	0.8035
	0.1	0.9295	0.8035
	0.2	0.9295	0.8035
	0.3	0.8971	0.7955
	0.4	0.8971	0.7955
	0.5	0.8635	0.7811
	0.6	0.7320	0.7024
	0.7	0.6047	0.6111
	0.8	0.5125	0.5235
	0.9	0.1333	0.1325
	1.0	0.0000	0.0000
SUVPickupTruck	0.0	0.0345	0.0000
	0.1	0.0000	0.0000
	0.2	0.0000	0.0000
	0.3	0.0000	0.0000
	0.4	0.0000	0.0000
	0.5	0.0000	0.0000
	0.6	0.0000	0.0000
	0.7	0.0000	0.0000
	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000
TrafficCone	0.0	0.0000	0.0000
	0.1	0.0000	0.0000
	0.2	0.0000	0.0000
	0.3	0.0000	0.0000
	0.4	0.0000	0.0000
	0.5	0.0000	0.0000
	0.6	0.0000	0.0000
	0.7	0.0000	0.0000
	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000
TrafficLight	0.0	0.0690	0.0000
	0.1	0.0000	0.0000
	0.2	0.0000	0.0000
	0.3	0.0000	0.0000
	0.4	0.0000	0.0000
	0.5	0.0000	0.0000
	0.6	0.0000	0.0000
	0.7	0.0000	0.0000
	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000

	1.0	0.0000	0.0000
TrafficLight	0.0	0.0690	0.0000
	0.1	0.0000	0.0000
	0.2	0.0000	0.0000
	0.3	0.0000	0.0000
	0.4	0.0000	0.0000
	0.5	0.0000	0.0000
	0.6	0.0000	0.0000
	0.7	0.0000	0.0000
	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000
Train	0.0	0.0000	0.0000
	0.1	0.0000	0.0000
	0.2	0.0000	0.0000
	0.3	0.0000	0.0000
	0.4	0.0000	0.0000
	0.5	0.0000	0.0000
	0.6	0.0000	0.0000
	0.7	0.0000	0.0000
	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000
Tree	0.0	0.7688	0.5576
	0.1	0.7512	0.5568
	0.2	0.6993	0.5369
	0.3	0.6530	0.5169
	0.4	0.6067	0.4845
	0.5	0.4805	0.4181
	0.6	0.3436	0.3031
	0.7	0.2238	0.2031
	0.8	0.1007	0.0910
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000
Truck_Bus	0.0	0.0000	0.0000
	0.1	0.0000	0.0000
	0.2	0.0000	0.0000
	0.3	0.0000	0.0000
	0.4	0.0000	0.0000
	0.5	0.0000	0.0000
	0.6	0.0000	0.0000
	0.7	0.0000	0.0000
	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000
Tunnel	0.0	0.0000	0.0000
	0.1	0.0000	0.0000
	0.2	0.0000	0.0000
	0.3	0.0000	0.0000
	0.4	0.0000	0.0000
	0.5	0.0000	0.0000
	0.6	0.0000	0.0000
	0.7	0.0000	0.0000
	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000

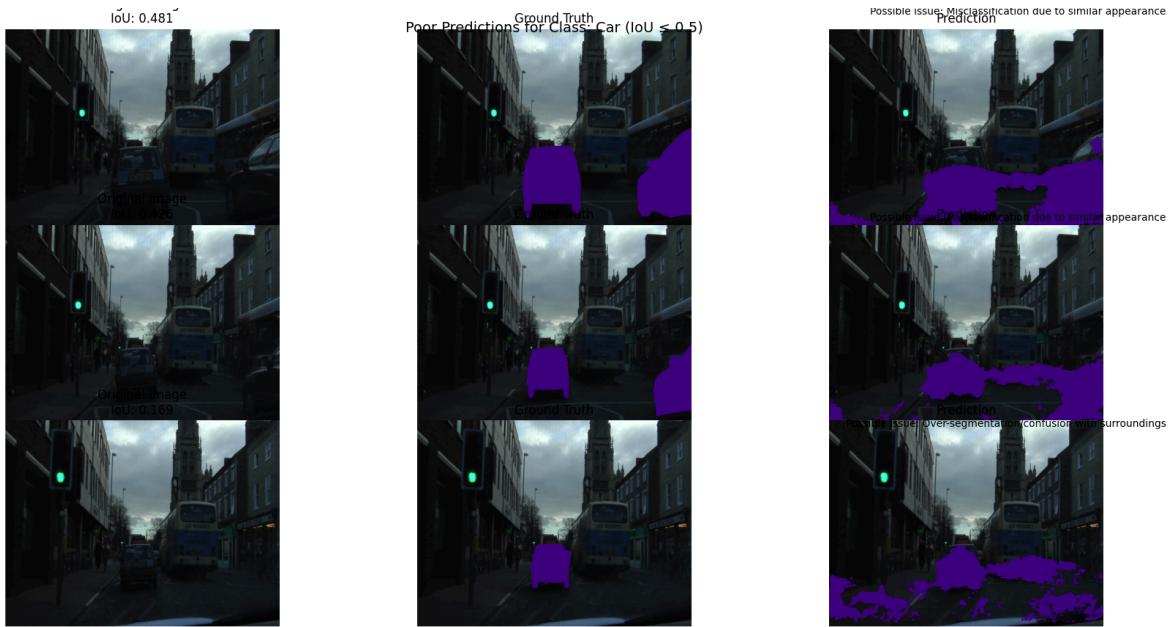
VegetationMisc	0.0	0.0953	0.0008
	0.1	0.0000	0.0000
	0.2	0.0000	0.0000
	0.3	0.0000	0.0000
	0.4	0.0000	0.0000
	0.5	0.0000	0.0000
	0.6	0.0000	0.0000
	0.7	0.0000	0.0000
	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000
Void	0.0	0.1984	0.0954
	0.1	0.0206	0.0105
	0.2	0.0000	0.0000
	0.3	0.0000	0.0000
	0.4	0.0000	0.0000
	0.5	0.0000	0.0000
	0.6	0.0000	0.0000
	0.7	0.0000	0.0000
	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000
Wall	0.0	0.4121	0.2211
	0.1	0.3784	0.2140
	0.2	0.2938	0.1802
	0.3	0.1262	0.0918
	0.4	0.0808	0.0609
	0.5	0.0530	0.0442
	0.6	0.0000	0.0000
	0.7	0.0000	0.0000
	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000

wandb:

C.

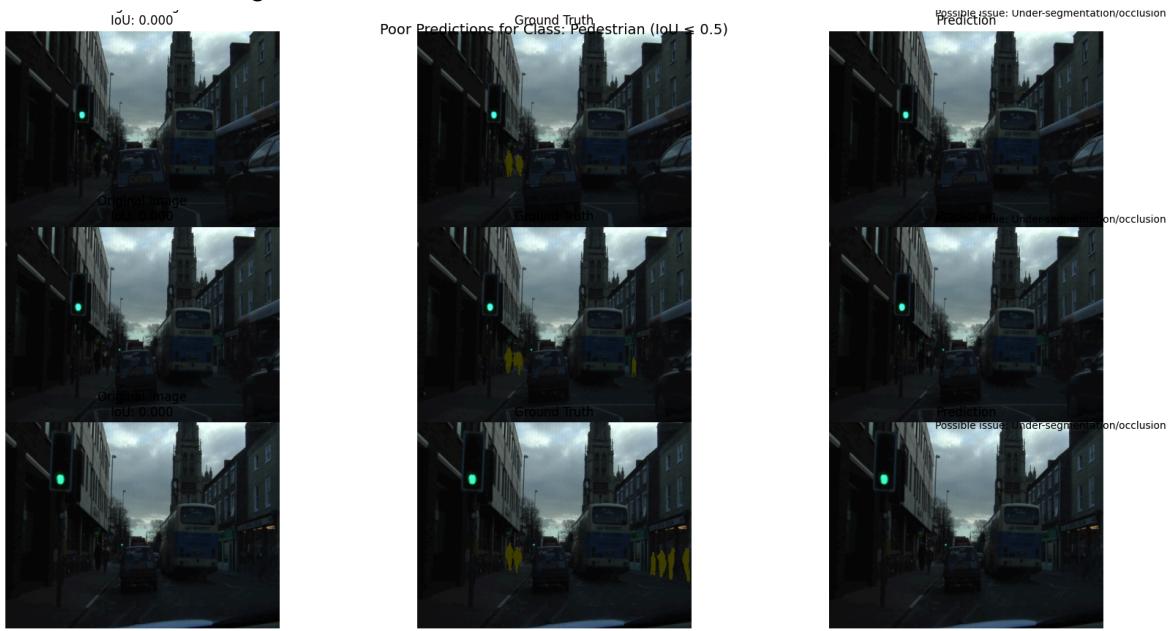
Car:

Possible issue: Misclassification due to similar appearance



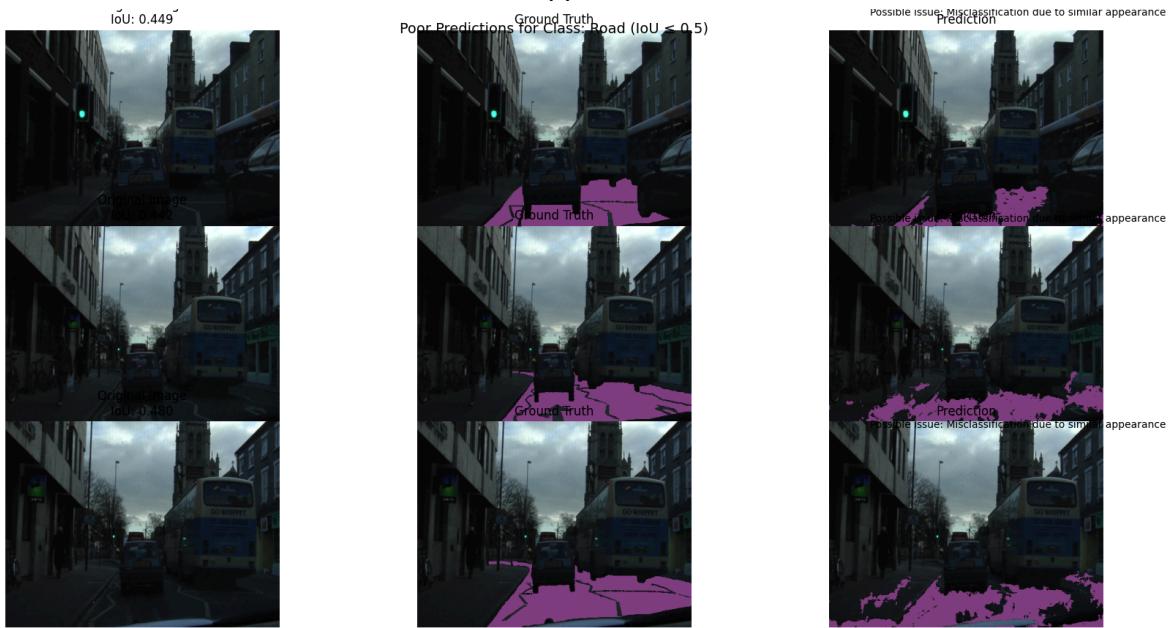
Pedestrian:

Possible issue: Under-segmentation

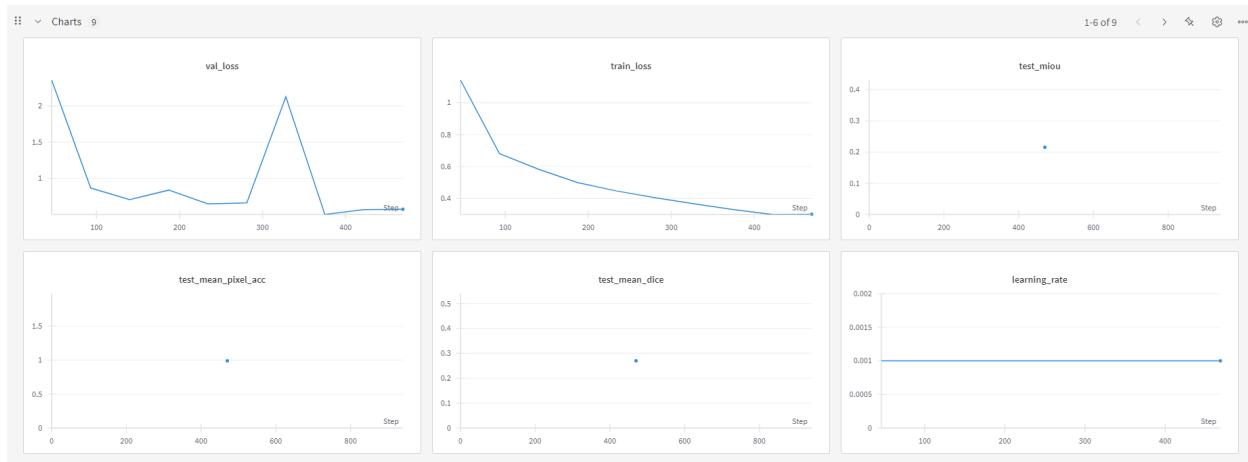


Road:

## Possible issue: Misclassification due to similar appearance



### 3. a.



b.

**Test Set Evaluation Results:**

=====  
Mean IoU (mIoU): 0.2155

**Class-wise Results:**

Class	Pixel Acc	Dice	IoU
Animal	0.9999	0.0000	0.0000
Archway	0.9998	0.0000	0.0000
Bicyclist	0.9940	0.2946	0.2090
Bridge	0.9994	0.0000	0.0000
Building	0.9444	0.8254	0.7210
Car	0.9780	0.7048	0.5789
CartLuggagePram	0.9997	0.0000	0.0000
Child	0.9997	0.0000	0.0000
Column_Pole	0.9892	0.2480	0.1477
Fence	0.9898	0.2166	0.1490
LaneMkgsDriv	0.9881	0.5131	0.3559
LaneMkgsNonDriv	1.0000	0.0000	0.0000
Misc_Text	0.9957	0.2612	0.1675
MotorcycleScooter	0.9998	0.0000	0.0000
OtherMoving	0.9968	0.1783	0.1221
ParkingBlock	0.9948	0.1229	0.0814
Pedestrian	0.9899	0.3495	0.2252
Road	0.9681	0.9270	0.8667
RoadShoulder	0.9959	0.0803	0.0663
Sidewalk	0.9814	0.7692	0.6465
SignSymbol	0.9991	0.0399	0.0267
Sky	0.9837	0.9513	0.9081
SUVPickupTruck	0.9781	0.0784	0.0463
TrafficCone	1.0000	0.0000	0.0000
TrafficLight	0.9975	0.3325	0.2485
Train	1.0000	0.0000	0.0000
Tree	0.9532	0.7334	0.6137
Truck_Bus	0.9891	0.0803	0.0530
Tunnel	1.0000	0.0000	0.0000
VegetationMisc	0.9876	0.1762	0.1208
Void	0.9735	0.3993	0.2860
Wall	0.9883	0.3550	0.2554

Precision and Recall at different IoU thresholds:

Class	Threshold	Precision	Recall
Animal	0.0	0.0000	0.0000
	0.1	0.0000	0.0000
	0.2	0.0000	0.0000
	0.3	0.0000	0.0000
	0.4	0.0000	0.0000
	0.5	0.0000	0.0000
	0.6	0.0000	0.0000
	0.7	0.0000	0.0000
	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000
Archway	0.0	0.0000	0.0000
	0.1	0.0000	0.0000
	0.2	0.0000	0.0000
	0.3	0.0000	0.0000
	0.4	0.0000	0.0000
	0.5	0.0000	0.0000
	0.6	0.0000	0.0000
	0.7	0.0000	0.0000
	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000
Bicyclist	0.0	0.3422	0.3600
	0.1	0.3043	0.3280
	0.2	0.2748	0.3227
	0.3	0.2010	0.2529
	0.4	0.1124	0.1501
	0.5	0.0926	0.1278
	0.6	0.0491	0.0682
	0.7	0.0269	0.0340
	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000
Bridge	0.0	0.0000	0.0000
	0.1	0.0000	0.0000
	0.2	0.0000	0.0000
	0.3	0.0000	0.0000
	0.4	0.0000	0.0000
	0.5	0.0000	0.0000
	0.6	0.0000	0.0000
	0.7	0.0000	0.0000
	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000

<b>Fence</b>	0 . 0	0 . 3493	0 . 1666
	0 . 1	0 . 3116	0 . 1595
	0 . 2	0 . 2730	0 . 1462
	0 . 3	0 . 1742	0 . 1102
	0 . 4	0 . 1742	0 . 1102
	0 . 5	0 . 0292	0 . 0264
	0 . 6	0 . 0292	0 . 0264
	0 . 7	0 . 0000	0 . 0000
	0 . 8	0 . 0000	0 . 0000
	0 . 9	0 . 0000	0 . 0000
	1 . 0	0 . 0000	0 . 0000
<b>LaneMkgsDriv</b>	0 . 0	0 . 5513	0 . 5019
	0 . 1	0 . 5513	0 . 5019
	0 . 2	0 . 4879	0 . 4716
	0 . 3	0 . 3870	0 . 3982
	0 . 4	0 . 2648	0 . 2689
	0 . 5	0 . 0930	0 . 1002
	0 . 6	0 . 0000	0 . 0000
	0 . 7	0 . 0000	0 . 0000
	0 . 8	0 . 0000	0 . 0000
	0 . 9	0 . 0000	0 . 0000
	1 . 0	0 . 0000	0 . 0000
<b>LaneMkgsNonDriv</b>	0 . 0	0 . 0000	0 . 0000
	0 . 1	0 . 0000	0 . 0000
	0 . 2	0 . 0000	0 . 0000
	0 . 3	0 . 0000	0 . 0000
	0 . 4	0 . 0000	0 . 0000
	0 . 5	0 . 0000	0 . 0000
	0 . 6	0 . 0000	0 . 0000
	0 . 7	0 . 0000	0 . 0000
	0 . 8	0 . 0000	0 . 0000
	0 . 9	0 . 0000	0 . 0000
	1 . 0	0 . 0000	0 . 0000
<b>Misc_Text</b>	0 . 0	0 . 5488	0 . 1909
	0 . 1	0 . 4197	0 . 1743
	0 . 2	0 . 3077	0 . 1342
	0 . 3	0 . 1605	0 . 0760
	0 . 4	0 . 0965	0 . 0509
	0 . 5	0 . 0305	0 . 0194
	0 . 6	0 . 0000	0 . 0000
	0 . 7	0 . 0000	0 . 0000
	0 . 8	0 . 0000	0 . 0000
	0 . 9	0 . 0000	0 . 0000
	1 . 0	0 . 0000	0 . 0000
<b>MotorcycleScooter</b>	0 . 0	0 . 0000	0 . 0000
	0 . 1	0 . 0000	0 . 0000
	0 . 2	0 . 0000	0 . 0000
	0 . 3	0 . 0000	0 . 0000
	0 . 4	0 . 0000	0 . 0000
	0 . 5	0 . 0000	0 . 0000
	0 . 6	0 . 0000	0 . 0000
	0 . 7	0 . 0000	0 . 0000
	0 . 8	0 . 0000	0 . 0000
	0 . 9	0 . 0000	0 . 0000
	1 . 0	0 . 0000	0 . 0000

	0.0	0.1	0.2	0.3	0.4	0.5	0.6	0.7	0.8	0.9	1.0
OtherMoving	0.2020	0.1790	0.1400	0.1079	0.0725	0.0725	0.0272	0.0000	0.0000	0.0000	0.0000
ParkingBlock	0.1499	0.0989	0.0866	0.0777	0.0371	0.0000	0.0000	0.0000	0.0000	0.0000	0.0000
Pedestrian	0.3474	0.3095	0.1952	0.0950	0.0629	0.0452	0.0259	0.0000	0.0000	0.0000	0.0000
Road	0.9242	0.9242	0.9242	0.9242	0.9242	0.9242	0.9242	0.9004	0.8394	0.3948	0.0000
RoadShoulder	0.0742	0.0691	0.0691	0.0691	0.0691	0.0514	0.0514	0.0293	0.0293	0.0000	0.0000

	1.0	0.0000	0.0000
Sidewalk	0.0	0.8338	0.7264
	0.1	0.8338	0.7264
	0.2	0.8338	0.7264
	0.3	0.7981	0.7047
	0.4	0.7804	0.6894
	0.5	0.6929	0.6346
	0.6	0.6050	0.5708
	0.7	0.3189	0.3061
	0.8	0.1942	0.1896
	0.9	0.0663	0.0652
	1.0	0.0000	0.0000
SignSymbol	0.0	0.1222	0.0294
	0.1	0.0547	0.0277
	0.2	0.0547	0.0277
	0.3	0.0547	0.0277
	0.4	0.0000	0.0000
	0.5	0.0000	0.0000
	0.6	0.0000	0.0000
	0.7	0.0000	0.0000
	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000
Sky	0.0	0.9632	0.9404
	0.1	0.9632	0.9404
	0.2	0.9632	0.9404
	0.3	0.9632	0.9404
	0.4	0.9632	0.9404
	0.5	0.9632	0.9404
	0.6	0.9632	0.9404
	0.7	0.9632	0.9404
	0.8	0.9632	0.9404
	0.9	0.6340	0.6314
	1.0	0.0000	0.0000
SUVPickupTruck	0.0	0.0766	0.1017
	0.1	0.0559	0.0794
	0.2	0.0411	0.0434
	0.3	0.0198	0.0159
	0.4	0.0000	0.0000
	0.5	0.0000	0.0000
	0.6	0.0000	0.0000
	0.7	0.0000	0.0000
	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000
TrafficCone	0.0	0.0000	0.0000
	0.1	0.0000	0.0000
	0.2	0.0000	0.0000
	0.3	0.0000	0.0000
	0.4	0.0000	0.0000
	0.5	0.0000	0.0000
	0.6	0.0000	0.0000
	0.7	0.0000	0.0000
	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000

TrafficLight	0.0	0.3638	0.3349
	0.1	0.3638	0.3349
	0.2	0.3568	0.3245
	0.3	0.2819	0.2942
	0.4	0.2430	0.2612
	0.5	0.1705	0.1957
	0.6	0.1059	0.1117
	0.7	0.0000	0.0000
	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000
Train	0.0	0.0000	0.0000
	0.1	0.0000	0.0000
	0.2	0.0000	0.0000
	0.3	0.0000	0.0000
	0.4	0.0000	0.0000
	0.5	0.0000	0.0000
	0.6	0.0000	0.0000
	0.7	0.0000	0.0000
	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000
Tree	0.0	0.6530	0.8640
	0.1	0.6530	0.8640
	0.2	0.6530	0.8640
	0.3	0.6454	0.8328
	0.4	0.5896	0.7352
	0.5	0.5744	0.7067
	0.6	0.4924	0.5817
	0.7	0.3702	0.4251
	0.8	0.0931	0.1016
	0.9	0.0315	0.0339
	1.0	0.0000	0.0000
Truck_Bus	0.0	0.1897	0.0604
	0.1	0.1161	0.0465
	0.2	0.0851	0.0420
	0.3	0.0657	0.0329
	0.4	0.0319	0.0209
	0.5	0.0319	0.0209
	0.6	0.0000	0.0000
	0.7	0.0000	0.0000
	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000
Tunnel	0.0	0.0000	0.0000
	0.1	0.0000	0.0000
	0.2	0.0000	0.0000
	0.3	0.0000	0.0000
	0.4	0.0000	0.0000
	0.5	0.0000	0.0000
	0.6	0.0000	0.0000
	0.7	0.0000	0.0000
	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000

VegetationMisc	0.0	0.1349	0.3528
	0.1	0.1261	0.2668
	0.2	0.1110	0.1856
	0.3	0.0760	0.1248
	0.4	0.0643	0.0914
	0.5	0.0643	0.0914
	0.6	0.0255	0.0297
	0.7	0.0000	0.0000
	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000
Void	0.0	0.6358	0.3205
	0.1	0.4904	0.2976
	0.2	0.4286	0.2755
	0.3	0.4048	0.2648
	0.4	0.2909	0.2108
	0.5	0.2067	0.1571
	0.6	0.0615	0.0517
	0.7	0.0324	0.0266
	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000
Wall	0.0	0.5533	0.3151
	0.1	0.4445	0.2774
	0.2	0.3455	0.2455
	0.3	0.3056	0.2275
	0.4	0.2012	0.1722
	0.5	0.2012	0.1722
	0.6	0.1478	0.1266
	0.7	0.0000	0.0000
	0.8	0.0000	0.0000
	0.9	0.0000	0.0000
	1.0	0.0000	0.0000

wandb:

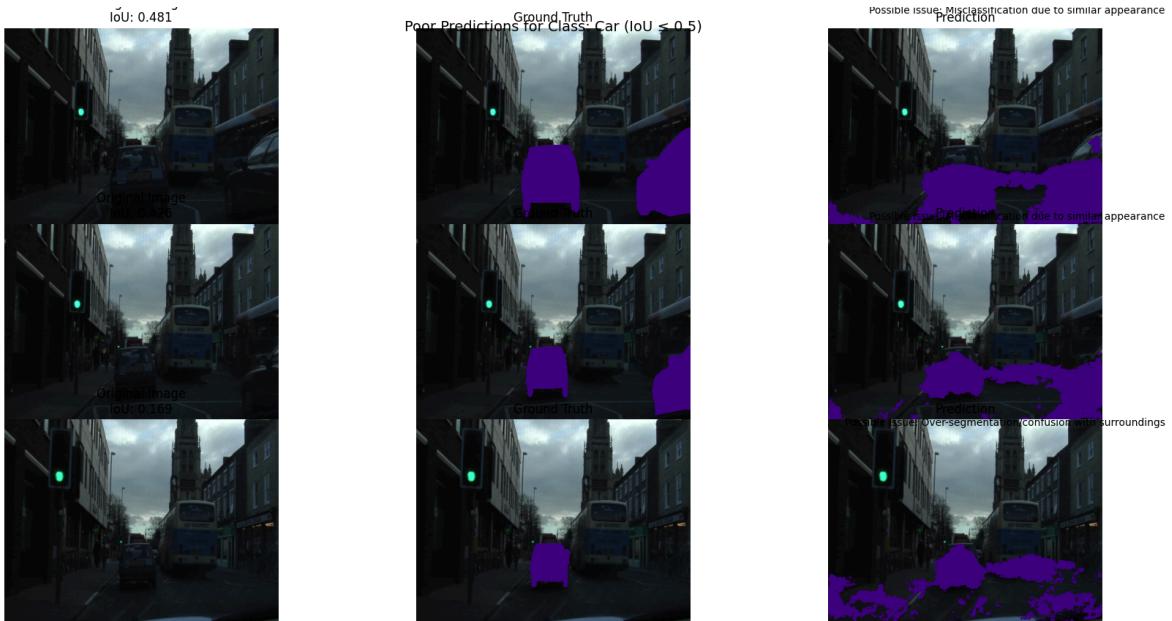
```

wandb:
wandb: Run history:
wandb:   batch [REDACTED]
wandb:   batch_loss [REDACTED]
wandb:   epoch [REDACTED]
wandb:   learning_rate [REDACTED]
wandb:   test_mean_dice [REDACTED]
wandb: test_mean_pixel_acc [REDACTED]
wandb:   test_miou [REDACTED]
wandb:   train_loss [REDACTED]
wandb:   val_loss [REDACTED]
wandb:
wandb: Run summary:
wandb:   batch 459
wandb:   batch_loss 0.31513
wandb:   epoch 10
wandb:   learning_rate 0.001
wandb:   test_mean_dice 0.26991
wandb: test_mean_pixel_acc 0.9892
wandb:   test_miou 0.21549
wandb:   train_loss 0.30113
wandb:   val_loss 0.57311
wandb:
wandb: View run trim-cherry-3 at: https://wandb.ai/vikranth2764-na/segmentation-deeplabv3/runs/e10v66ic
wandb: View project at: https://wandb.ai/vikranth2764-na/segmentation-deeplabv3
wandb: Synced 5 W&B file(s), 0 media file(s), 0 artifact file(s) and 0 other file(s)
wandb: Find logs at: ./wandb/run-20250223_172020-e10v66ic/logs

```

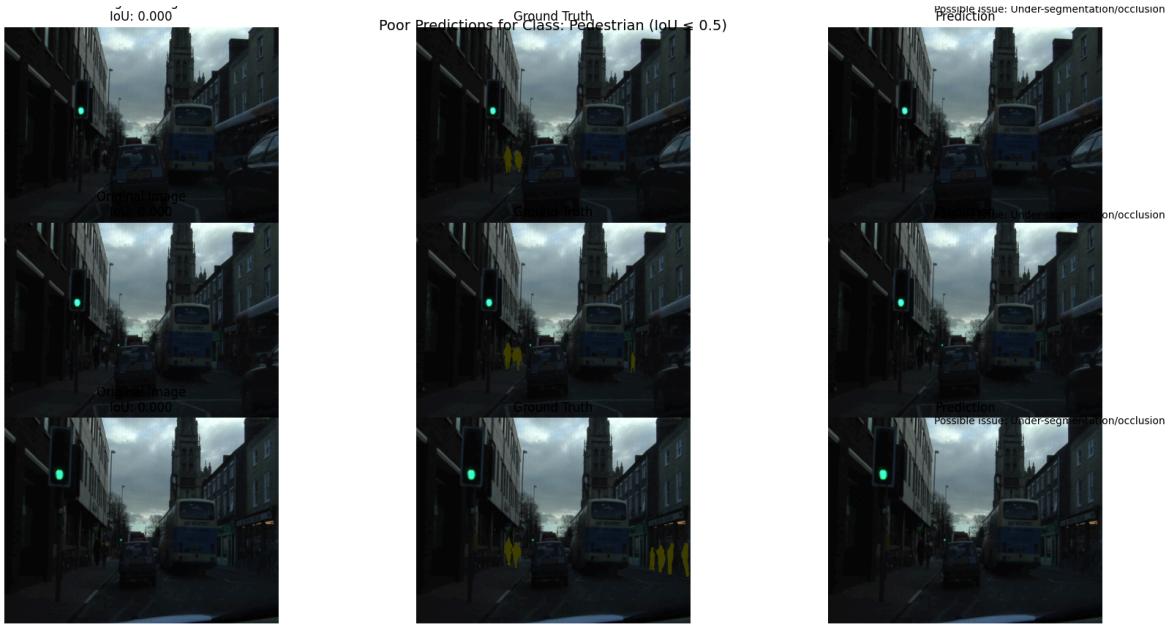
### c. Car:

Possible issue: Misclassification due to similar appearance



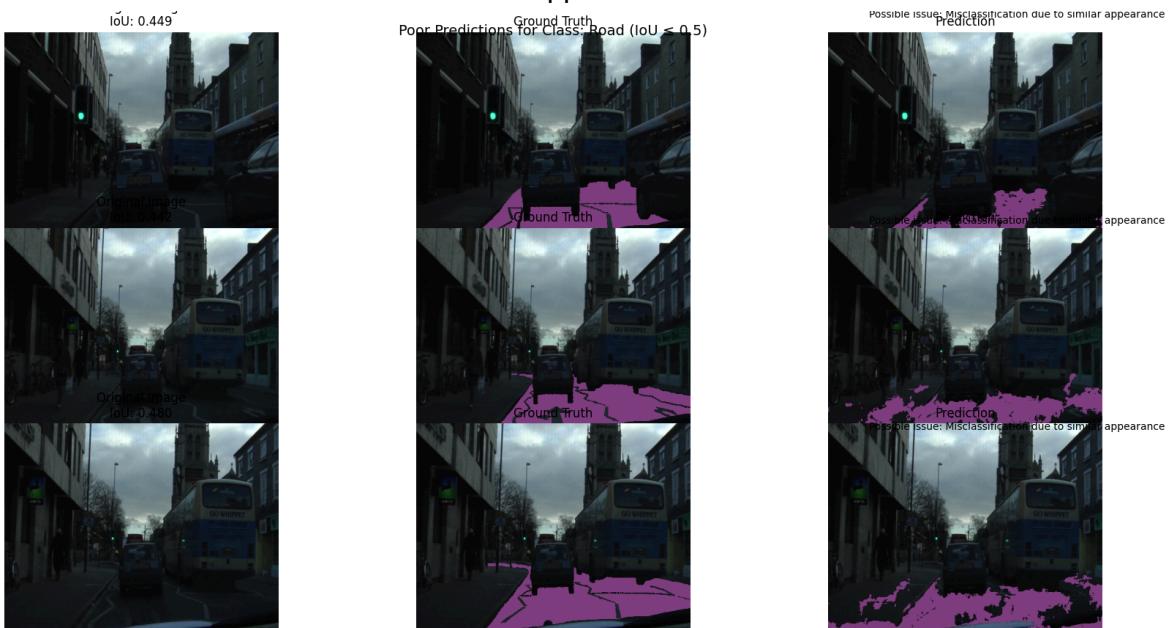
### Pedestrian:

Possible issue: Under-segmentation



Road:

Possible issue: Misclassification due to similar appearance



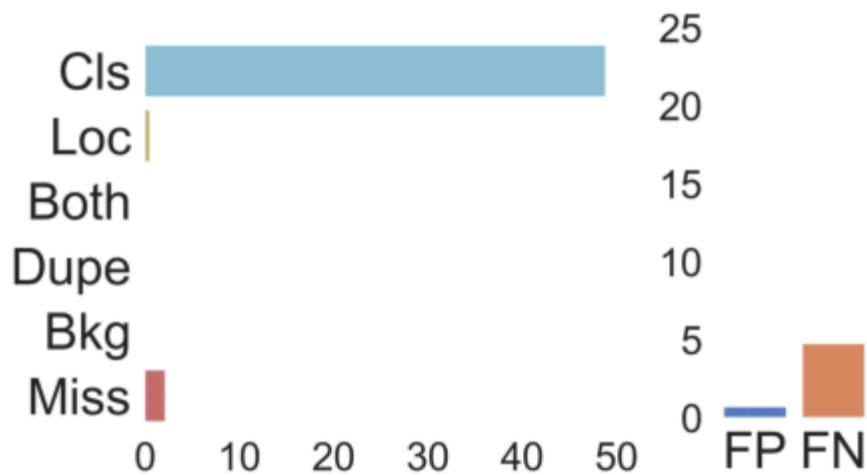
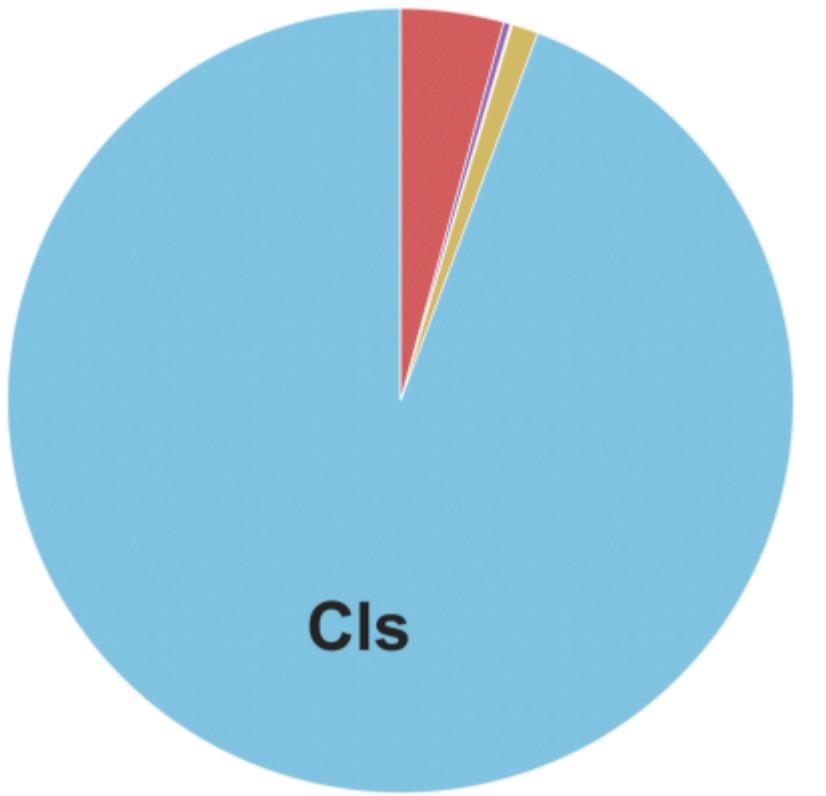
#### 4. 1. a. No deliverables

```
# Download the Dataset by using Ultralytics API

from ultralytics.data.utils import download
download(url='https://github.com/ultralytics/yolov5/releases/download/v1.0/coco2017val.zip')

with zipfile.ZipFile('coco2017val.zip', 'r') as zip_ref:
    zip_ref.extractall()
```

# coco\_predictions



b. c.

```
-- coco_predictions --
bbox AP @ 50: 8.61

      Main Errors
=====
Type    Cls     Loc     Both     Dupe     Bkg     Miss
=====
dAP    48.84   0.56   0.03   0.02   0.15   2.19
=====

      Special Error
=====
Type  FalsePos  FalseNeg
=====
dAP      0.73      4.78
=====
```

**mAP @ 50: 8.61%**

- This is the **Mean Average Precision (mAP)** at an **IoU threshold of 0.5**, which measures how well the model's predictions align with the ground truth.

**TIDE Statistics:**

**Cls (Classification Error): 48.84%**

- The **dominant error type**.
- Model is **misclassifying objects** (wrong labels).

**Loc (Localization Error): 0.56%**

- Very low, meaning **bounding boxes are mostly accurate**.

**Both (Cls + Loc Errors Combined): 0.03%**

- Model rarely makes both classification and localization errors together.

**Dupe (Duplicate Detections): 0.02%**

- Model almost never predicts the **same object multiple times**.

**Bkg (Background Error): 0.15%**

- Minimal **false positives from the background**.

**Miss (Missed Detections): 2.19%**

- The model misses a small percentage of actual objects.

**also Positives (FP): 0.73%**

- Incorrect detections where nothing should be detected.

#### **False Negatives (FN): 4.78%**

- Ground truth objects were not detected.

d. ECE: 0.2308927360370347.

This indicates significant **overconfidence**, meaning the model assigns high confidence to incorrect predictions. This aligns with the **48.84% classification error** from TIDE, showing that misclassification is the biggest issue. To improve, apply **temperature scaling** or **lower the confidence threshold** to better align confidence scores with accuracy.

```

loading annotations into memory...
Done (t=0.12s)
creating index...
index created!

Analyzing small objects...
Loading and preparing results...
DONE (t=0.04s)
creating index...
index created!
Processing images: 100%|██████████| 5000/5000 [00:00<00:00, 42112.75it/s]
ECE for small objects: 0.12887502587339783

Analyzing medium objects...
Loading and preparing results...
DONE (t=0.14s)
creating index...
index created!
Processing images: 100%|██████████| 5000/5000 [00:00<00:00, 48291.88it/s]
ECE for medium objects: 0.15358833936401287

Analyzing large objects...
Loading and preparing results...
DONE (t=0.05s)
creating index...
index created!
Processing images: 100%|██████████| 5000/5000 [00:00<00:00, 26996.56it/s]
ECE for large objects: 0.16464595028723078

```

e.

```

Analyzing small objects...

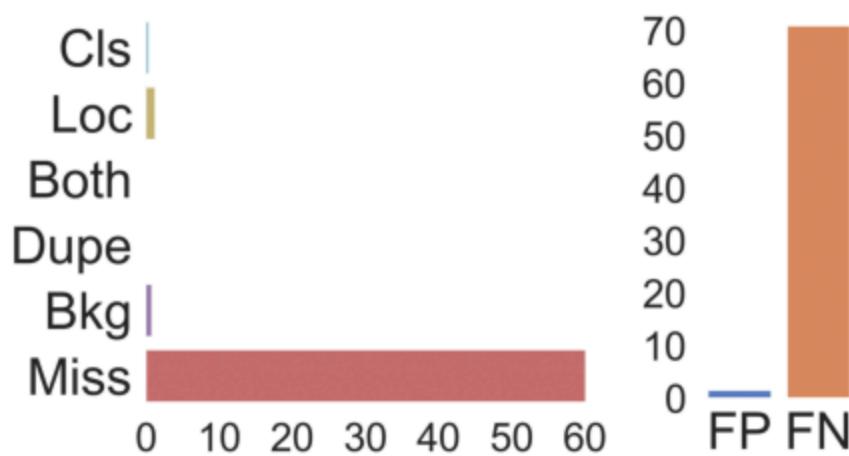
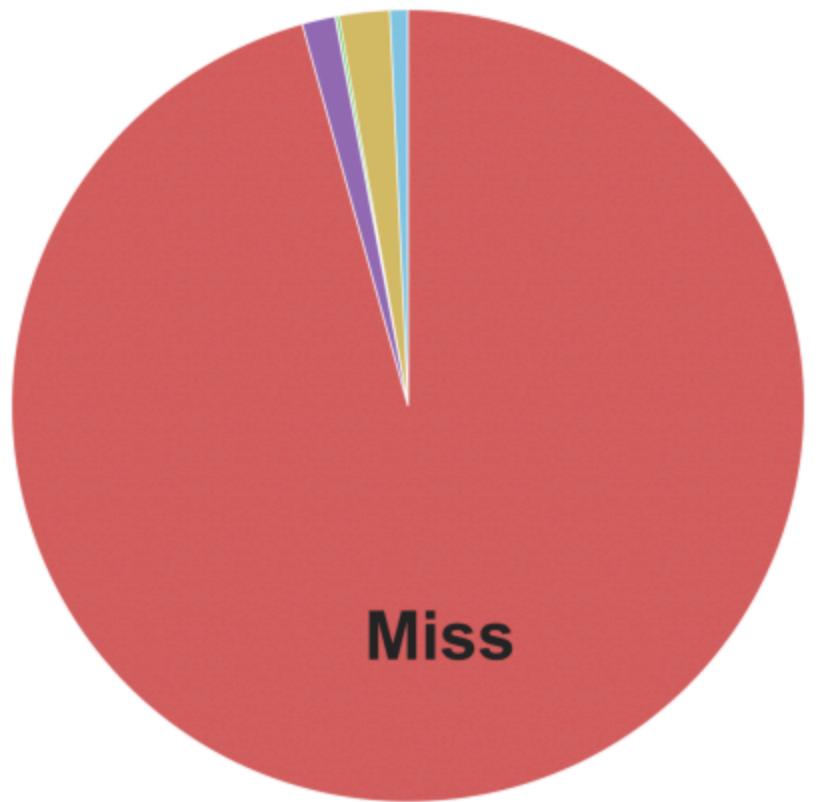
TIDE Statistics for small objects:
-- coco_predictions_small --

bbox AP @ 50: 9.38

Main Errors
=====
Type    Cls      Loc     Both     Dupe     Bkg     Miss
=====
dAP     0.46     1.29    0.08    0.02     0.85    60.12
=====

Special Error
=====
Type  FalsePos  FalseNeg
-----
dAP      1.54     70.94
=====
```

# coco\_predictions\_small



```
Analyzing medium objects...
```

```
TIDE Statistics for medium objects:  
-- coco_predictions_medium --
```

```
bbox AP @ 50: 22.19
```

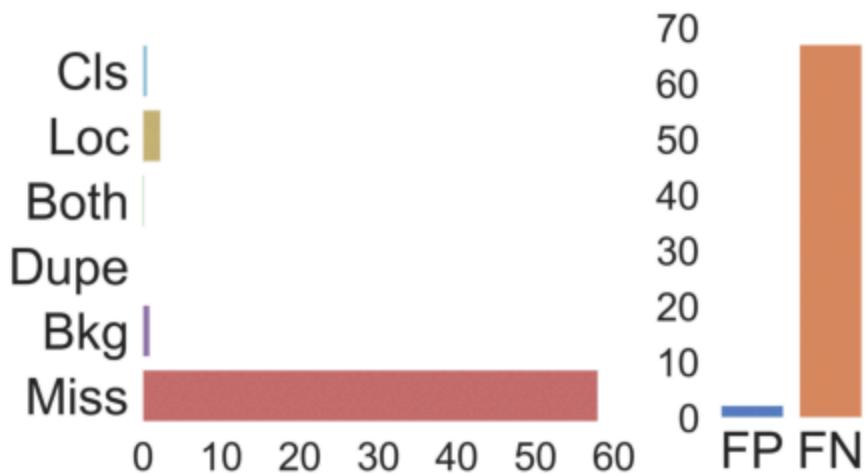
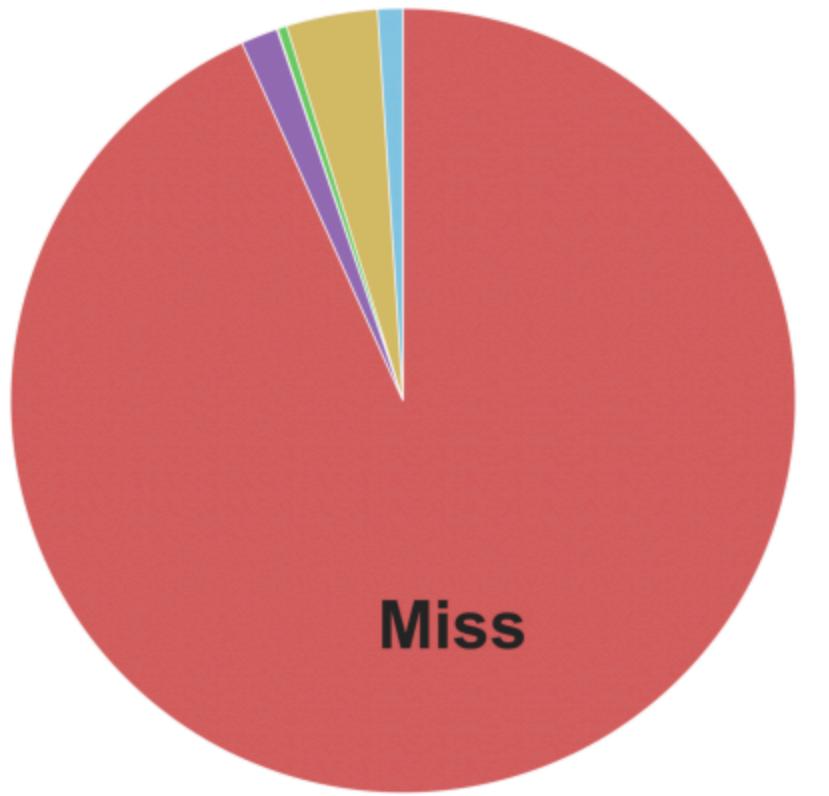
```
    Main Errors
```

Type	Cls	Loc	Both	Dupe	Bkg	Miss
dAP	0.64	2.34	0.24	0.02	0.94	58.00

```
    Special Error
```

Type	FalsePos	FalseNeg
dAP	2.27	67.04

# coco\_predictions\_medium



```
Analyzing large objects...
```

```
TIDE Statistics for large objects:
```

```
-- coco_predictions_large --
```

```
bbox AP @ 50: 31.29
```

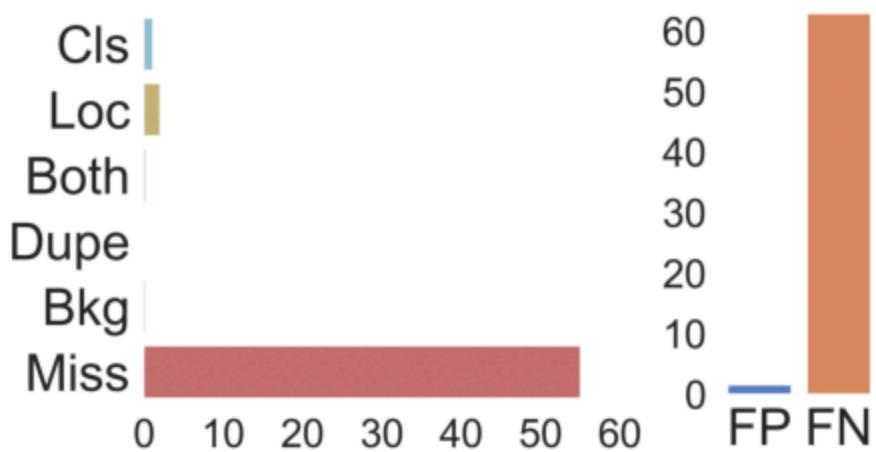
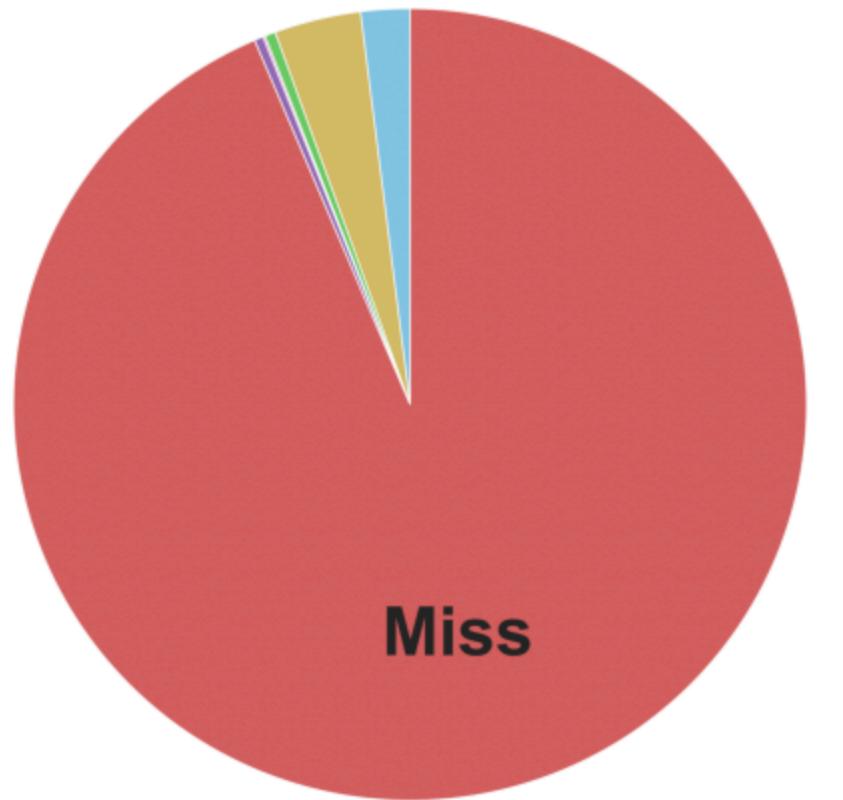
```
    Main Errors
```

Type	Cls	Loc	Both	Dupe	Bkg	Miss
dAP	1.18	2.07	0.26	0.05	0.21	55.07

```
    Special Error
```

Type	FalsePos	FalseNeg
dAP	1.56	62.91

# coco\_predictions\_large



ECE (Calibration Error)

- Small objects: 0.1289, Medium: 0.1536, Large: 0.1646 → ECE increases with object size.

- **Lower ECE for small objects** suggests the model is better calibrated for them.
- **Higher ECE for large objects** indicates overconfidence in incorrect predictions.

**AP @ 50 for Small: 9.38, Medium: 22.19, Large: 31.29 → Performance improves with size.**

- **Missed Detections are dominant (Small: 60.12%, Medium: 58.00%, Large: 55.07%)**  
→ The model struggles to detect objects, especially small ones.
- **Localization error is highest for medium & large objects (2.34%, 2.07%),** meaning bounding box accuracy decreases as object size increases.
- **False Negatives are very high across all sizes (Small: 70.94%, Medium: 67.04%, Large: 62.91%),** indicating many objects remain undetected.

**The model performs best on large objects but still struggles with missed detections.**

**False Negatives are the biggest issue across all sizes, suggesting objects are not being detected at all.**

**Calibration is better for small objects but worsens for medium and large ones.**