

CV Assignment 2

Vikranth Udandarao
2022570

Theory

C V Theory

$$N.S = N2022570$$

$$1- \mathbf{v} = \begin{bmatrix} 3 \\ -1 \\ 4 \end{bmatrix}$$

Step 1 - Rotation of $\underline{\underline{n}}$ about y-axis.

A rotation about y-axis by an angle θ

has the direction

$$R_y(\theta) = \begin{bmatrix} \cos \theta & 0 & \sin \theta \\ 0 & 1 & 0 \\ -\sin \theta & 0 & \cos \theta \end{bmatrix}$$

$$\theta = -\frac{\pi}{6}$$

$$\rightarrow \cos -\frac{\pi}{6} = \frac{\sqrt{3}}{2}, \sin -\frac{\pi}{6} = -\frac{1}{2}$$

$$R_y\left(-\frac{\pi}{6}\right) = \begin{bmatrix} \frac{\sqrt{3}}{2} & 0 & \frac{1}{2} \\ 0 & 1 & 0 \\ \frac{1}{2} & 0 & \frac{\sqrt{3}}{2} \end{bmatrix}$$

Step²
Rotation $\frac{\pi}{4}$ about x-axis -

$$R_x(\theta) = \begin{bmatrix} 1 & 0 & 0 \\ 0 & \cos \theta & -\sin \theta \\ 0 & \sin \theta & \cos \theta \end{bmatrix}$$

$$R_x(\frac{\pi}{4}) = \begin{bmatrix} 1 & 0 & 0 \\ 0 & \frac{1}{\sqrt{2}} & -\frac{1}{\sqrt{2}} \\ 0 & \frac{1}{\sqrt{2}} & \frac{1}{\sqrt{2}} \end{bmatrix}$$

Step³ - Reflection across XZ plane -

$$n \rightarrow n$$

$$y \rightarrow -y$$

$$z \rightarrow z$$

~~S_{xz}~~ $S_{xz} = \begin{bmatrix} 1 & 0 & 0 \\ 0 & -1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$

$$\Rightarrow V' = S_{xz} - R_x(\frac{\pi}{4}) - R_y(-\frac{\pi}{6}) \cdot V$$

$$R_x(\frac{\pi}{4}) \cdot R_y(-\frac{\pi}{6}) = \begin{bmatrix} 1 & 0 & 0 \\ 0 & \frac{1}{\sqrt{2}} & -\frac{1}{\sqrt{2}} \\ 0 & \frac{1}{\sqrt{2}} & \frac{1}{\sqrt{2}} \end{bmatrix} \begin{bmatrix} \frac{\sqrt{3}}{2} & 0 & \frac{-1}{2} \\ 0 & 1 & 0 \\ \frac{1}{2} & 0 & \frac{\sqrt{3}}{2} \end{bmatrix}$$

$$= \begin{pmatrix} \frac{\sqrt{3}}{2} & 0 & -\frac{1}{2} \\ -\frac{1}{2\sqrt{2}} & \frac{1}{2\sqrt{2}} & -\frac{\sqrt{3}}{2\sqrt{2}} \\ \frac{1}{2\sqrt{2}} & \frac{1}{2\sqrt{2}} & \frac{\sqrt{3}}{2\sqrt{2}} \end{pmatrix}$$

$$M = S_{n_2} \cdot \left(R_n \left(\frac{\pi}{3} \right) \cdot R_y \left(-\frac{\pi}{6} \right) \right)$$

$$M = \begin{bmatrix} 1 & 0 & 0 \\ 0 & -1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \cdot \begin{bmatrix} \frac{\sqrt{3}}{2} & 0 & -\frac{1}{2} \\ -\frac{1}{2\sqrt{2}} & \frac{1}{2\sqrt{2}} & -\frac{\sqrt{3}}{2\sqrt{2}} \\ \frac{1}{2\sqrt{2}} & \frac{1}{2\sqrt{2}} & \frac{\sqrt{3}}{2\sqrt{2}} \end{bmatrix}$$

-1- $M = \begin{bmatrix} \frac{\sqrt{3}}{2} & 0 & -\frac{1}{2} \\ \frac{1}{2\sqrt{2}} & -\frac{1}{2\sqrt{2}} & \frac{\sqrt{3}}{2\sqrt{2}} \\ \frac{1}{2\sqrt{2}} & \frac{1}{2\sqrt{2}} & -\frac{\sqrt{3}}{2\sqrt{2}} \end{bmatrix}$

1-2- $V = M \cdot V + T A^T \quad T = \begin{bmatrix} 1 \\ 0 \\ -2 \end{bmatrix}$

$$M - V = \begin{pmatrix} \frac{\sqrt{3}}{2} & 0 & -\frac{1}{2} & 1 \\ \frac{1}{2\sqrt{2}} & -\frac{1}{\sqrt{2}} & \frac{\sqrt{3}}{2\sqrt{2}} & 0 \\ \frac{1}{2\sqrt{2}} & \frac{1}{\sqrt{2}} & \frac{\sqrt{3}}{2\sqrt{2}} & -2 \\ 0 & 0 & 0 & 1 \end{pmatrix} \cdot \begin{pmatrix} 3 \\ -1 \\ -4 \\ 1 \end{pmatrix}$$

$$M \cdot V = \begin{pmatrix} \frac{3\sqrt{3} - 4}{2} \\ \frac{5\sqrt{3} + 4\sqrt{6}}{4} \\ \frac{\sqrt{2} + 4\sqrt{6}}{4} \\ 1 \end{pmatrix}$$

$$\rightarrow V^1 = \begin{pmatrix} \frac{3\sqrt{3} - 4}{2} \\ \frac{5\sqrt{2} + 4\sqrt{6}}{4} \\ \frac{\sqrt{2} + 4\sqrt{6} - 8}{4} \\ 0 \end{pmatrix}$$

$$\rightarrow V^1 = \begin{pmatrix} \frac{3\sqrt{3} - 4}{2} \\ \frac{5\sqrt{2} + 4\sqrt{6}}{4} \\ \frac{\sqrt{2} + 4\sqrt{6} - 8}{4} \\ 1 \end{pmatrix}$$

$$0 = \begin{bmatrix} 0 \\ 0 \\ 0 \end{bmatrix}$$

$$v = m \cdot 0 + t$$

$$v = \begin{bmatrix} -1 \\ 0 \\ -2 \end{bmatrix}$$

1-3- $R = R_x\left(\frac{\pi}{4}\right) \cdot R_y\left(-\frac{\pi}{2}\right)$, $\text{tr}(R) = 1 + 2 \cos \theta$

$$= \frac{\sqrt{3}}{2} + \frac{\sqrt{2}}{2} + \frac{\sqrt{6}}{4}$$

$$\Rightarrow 1 + 2 \cos \theta = \frac{\sqrt{3}}{2} + \frac{\sqrt{2}}{2} + \frac{\sqrt{6}}{4}$$

$$\Rightarrow \cos \theta \approx 0.5925$$

$$\Rightarrow \theta = \arcsin(0.5925)$$

$$= 53.7^\circ$$

Axes direction of current rotation in original frame
of reference.

Rotation Matrix R ,

$$R_n = h \rightarrow (R - I) h = 0$$

$$\rightarrow n = \cancel{0} \left(\frac{1}{\sqrt{2}} + \frac{\sqrt{3}}{2\sqrt{2}}, -\frac{1}{2} - \frac{1}{2\sqrt{2}}, -\frac{1}{2\sqrt{2}} \right)$$

1. 4- Rodriguez formula

$$R = I + \sin \theta K + (1 - \cos \theta) K^2$$

K is skew symmetric,

$$K = \begin{bmatrix} 0 & -n_2 & n_1 \\ n_2 & 0 & -n_3 \\ -n_1 & n_3 & 0 \end{bmatrix}$$

$$\rightarrow n' = (1.32, -0.85, -0.35)$$

Normalise -)

$$\|n'\| = \sqrt{(1.32)^2 + (-0.85)^2 + (-0.35)^2}$$

$$\therefore = \sqrt{2.6} = 1.6$$

$$\rightarrow h = (0.82, -0.53, -0.22) \text{ a unit vector}$$

and

$$n \cos = 53.7^\circ = 0.93 \text{ rad,}$$

$$\rightarrow \sin \theta = \sin(0.93) \approx 0.8.$$

$$\cos \theta \approx 0.6.$$

$$1 - \cos \theta = 0.4$$

$$K = \begin{pmatrix} 0, & 0.22 & -0.53 \\ -0.22 & 0 & -0.82 \\ 0.53 & 0.82 & 0 \end{pmatrix}$$

$$K^2 = \begin{pmatrix} -0.33 & -0.18 & -0.43 \\ -0.18 & -0.48 & 0.28 \\ -0.43 & 0.28 & -0.19 \end{pmatrix}$$

$$\rightarrow R = \begin{pmatrix} 0.86 & 0 & -0.15 \\ -0.25 & 0.7 & -0.61 \\ 0.15 & 0.7 & 0.61 \end{pmatrix}$$

This matches our previous R .

2- Image formation equations: $\mathbf{u} = \mathbf{K}[\mathbf{R}|\mathbf{t}] \mathbf{x}$

$$= \mathbf{u} = \mathbf{K}[\mathbf{R}|\mathbf{t}] \mathbf{x}$$

\mathbf{K} → intrinsic parameters

$[\mathbf{R}|\mathbf{t}]$ → extrinsic parameters

\mathbf{x} → 3D point

\mathbf{u} → image point in homogenous system

→ The extrinsic matrix $[\mathbf{R}|\mathbf{t}]$ transforms a world point into camera's coordinate.

$$\mathbf{x}_{\text{cam}} = [\mathbf{R}|\mathbf{t}] \mathbf{x} = \mathbf{R} \mathbf{x}_{3D} + \mathbf{t}$$

Step 1 - Defining camera coordinate frame (G)

→ The world is aligned with G's frame.

$$\rightarrow \mathbf{u}_1 = \mathbf{K}_1 [\mathbf{I}|0] = \mathbf{K}_1 \begin{bmatrix} \mathbf{I} & 0 \\ 0^T & 1 \end{bmatrix} \begin{bmatrix} \mathbf{x} \\ 1 \end{bmatrix}$$

$$\mathbf{u}_1 = \mathbf{K}_1 \begin{bmatrix} \mathbf{x} \\ 1 \end{bmatrix}$$

$$\mathbf{u}_1 = \mathbf{K}_1 \begin{bmatrix} \mathbf{x} \\ 1 \end{bmatrix}$$

$$\rightarrow \text{so, } u_1 = K_1 X_{3D},$$

$$X_{3D} = \begin{bmatrix} X \\ Y \\ Z \end{bmatrix}^T$$

$Step^2 - Camera 2$

C_2 's orientation is obtained by a pure 3D rotation R applied to C_1 's orientation and then is no translation between the cameras.

for the same world point X -

$$u_2 = K_2 [R^{(0)}] X$$

$$= K_2 \begin{bmatrix} R^0 \\ 0^T 1 \end{bmatrix} \begin{bmatrix} X \\ Y \\ Z \end{bmatrix}$$

$$\rightarrow u_2 = K_2 R X_{3D}$$

Now, we need $u_1 = K u_2$

$$\rightarrow u_1 = K_1 X_{3D}$$

$$u_2 = K_2 R X_{3D}$$

$$x_{3D} = R^{-1} K_2^{-1} u_2$$

$$\Rightarrow u_1 = K_1 x_{3D} \Rightarrow K_1 (R^{-1} K_2^{-1}, u_2)$$

$$= K_1 R T \cdot R_2^{-1} u_2$$

$$u_1 = H u_2$$

$$H = K_1 R T \cdot K_2^{-1}$$

H is ~~93×3~~ invertible matrix expressed

in terms of K_1 (internal unitary of C_1)

& K_2 (internal unitary of C_2) and

R (rotation from C_1 to C_2)

$$\begin{pmatrix} R & 0 \\ 0 & I_{70} \end{pmatrix} \cdot \begin{pmatrix} K_1 & 0 \\ 0 & K_2 \end{pmatrix} \cdot \begin{pmatrix} T & 0 \\ 0 & I_{70} \end{pmatrix} = H$$

Coding

3. Given Images:

1. **Estimated Intrinsic Camera Parameters:**

- **Focal Lengths:**

- $f_x = 956.64$ pixels
- $f_y = 957.55$ pixels

- **Principal Point:**

- $c_x = 369.05$ pixels
- $c_y = 651.41$ pixels

- **Skew Parameter:**

- 0.00 (indicating no skew between image axes)

- **Intrinsic Matrix (K):**

$$\begin{bmatrix} [956.63717909 & 0. & 369.05487997] \\ [0. & 957.55139547 & 651.41419153] \\ [0. & 0. & 1.] \end{bmatrix}$$

- **Reprojection Error:**

- 0.5091 pixels (indicating the average discrepancy between observed and projected points)

```
code.py 2, M  intrinsic_params_provided.txt X
A2 > Q3 >  intrinsic_params_provided.txt
1   Intrinsic Parameters (Provided Dataset):
2   Intrinsic Matrix (K):
3   [[956.63717909  0.          369.05487997]
4   [ 0.          957.55139547 651.41419153]
5   [ 0.          0.          1.        ]]
6
7   Focal Lengths: fx = 956.64, fy = 957.55 pixels
8   Principal Point: cx = 369.05, cy = 651.41 pixels
9   Skew Parameter: 0.00
10  Reprojection Error: 0.5091 pixels
11
```

2. Estimated Extrinsic Camera Parameters

Image 1

- **Rotation Matrix (R):**

```
[[ 0.99699847  0.01387206 -0.07616832]
```

```
[-0.02511535  0.9885627 -0.14870444]
```

```
[ 0.07323433  0.1501711  0.9859439 ]]
```

- **Translation Vector (t):**

```
[-118.14885488]
```

```
[-108.51488489]
```

```
[ 439.77334073]]
```

Image 2

- **Rotation Matrix (R):**

```
[[ 0.99828333  0.05391999 -0.02286982]
```

```
[-0.05069018  0.99100787  0.12383014]
```

```
[ 0.0293411 -0.12245829  0.99203985]]
```

- **Translation Vector (t):**

```
[[ -131.13909032]
```

```
[ -43.7887548 ]
```

```
[ 474.41320287]]
```

```
A2 > Q3 > extrinsic_params_provided.txt
1   Extrinsic Parameters (Provided Dataset):
2
3   Image 1:
4     Rotation Matrix (R):
5     [[ 0.99699847  0.01387206 -0.07616832]
6     [-0.02511535  0.9885627  -0.14870444]
7     [ 0.07323433  0.1501711   0.9859439 ]]
8   Translation Vector (t):
9   [[ -118.14885488]
10  [-108.51488489]
11  [ 439.77334073]]
12
13  Image 2:
14  Rotation Matrix (R):
15  [[ 0.99828333  0.05391999 -0.02286982]
16  [-0.05069018  0.99100787  0.12383014]
17  [ 0.0293411  -0.12245829  0.99203985]]
18  Translation Vector (t):
19  [[ -131.13909032]
20  [ -43.7887548 ]
21  [ 474.41320287]]
22
```

3. Radial Distortion Coefficients and Undistortion Results

The estimated **radial distortion coefficients** obtained from camera calibration are:

- **$k_1 = 0.1976$**
- **$k_2 = -0.7069$**
- **$k_3 = 0.0488$**

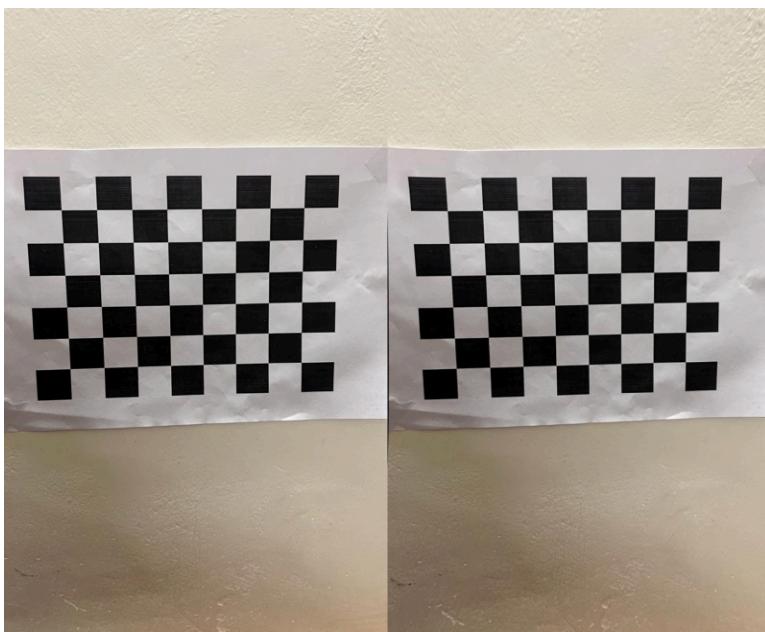
In addition to radial distortion, the tangential distortion coefficients are:

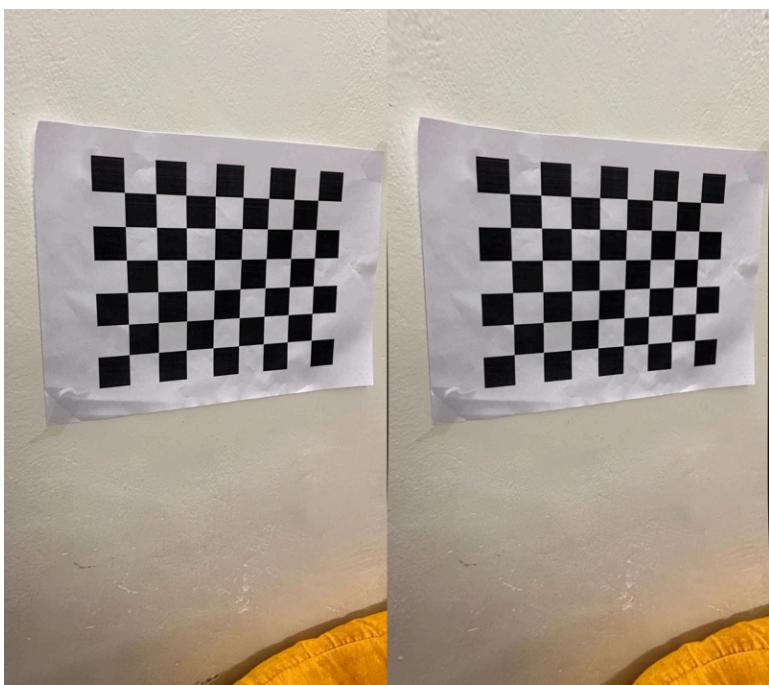
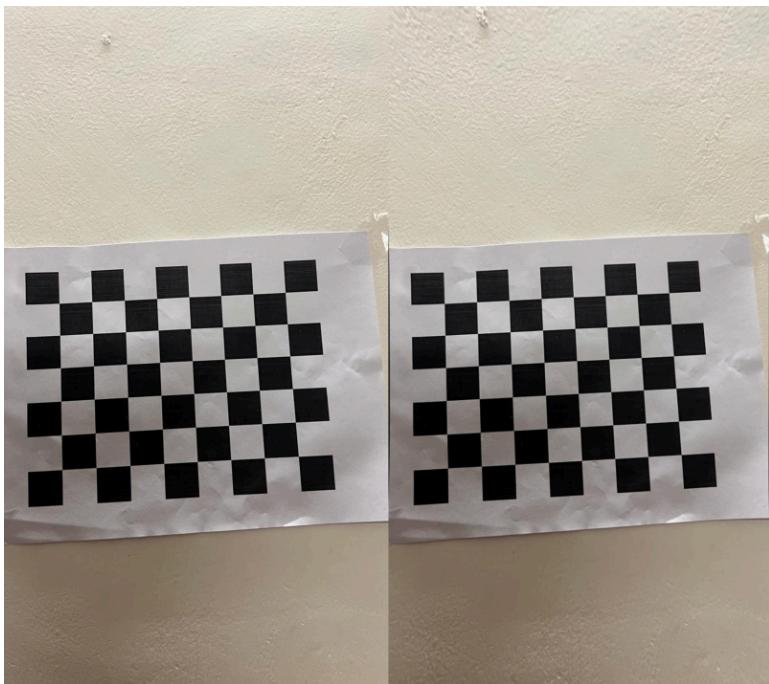
- $p_1 = 0.0034$
- $p_2 = 0.0070$

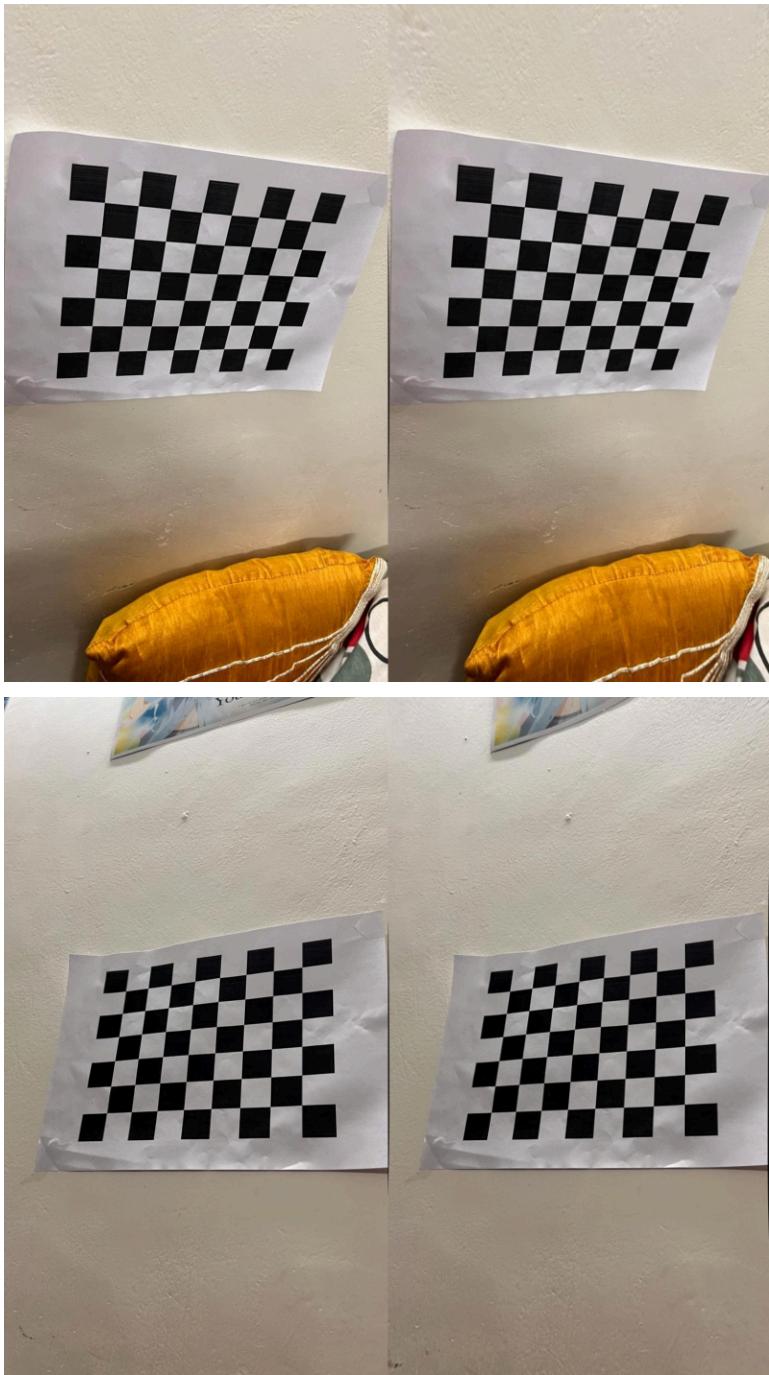
These coefficients were used to **undistort five raw images** of the calibration checkerboard. The undistortion was performed using OpenCV's `cv2.undistort()` function, applying the estimated intrinsic matrix and distortion coefficients.

Visual Results

The raw and undistorted images are shown side-by-side below:







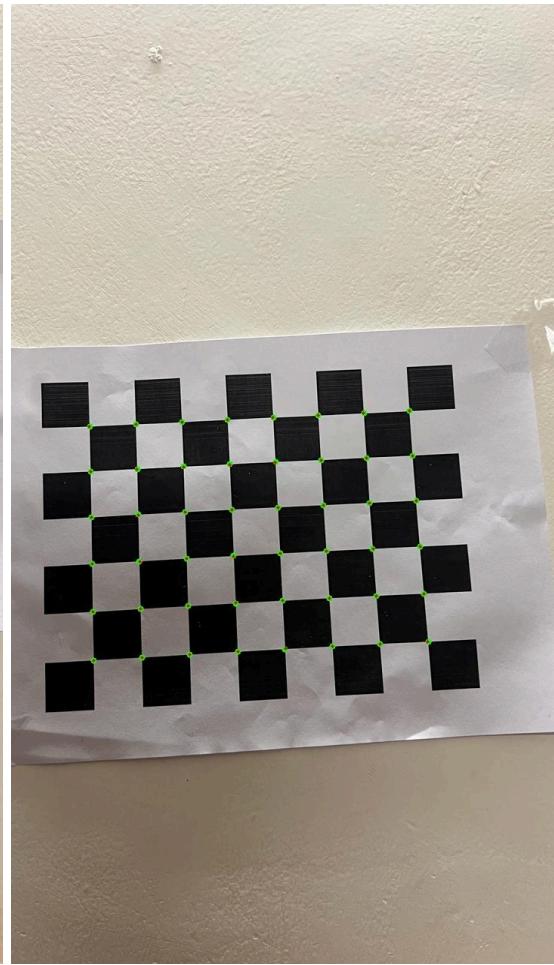
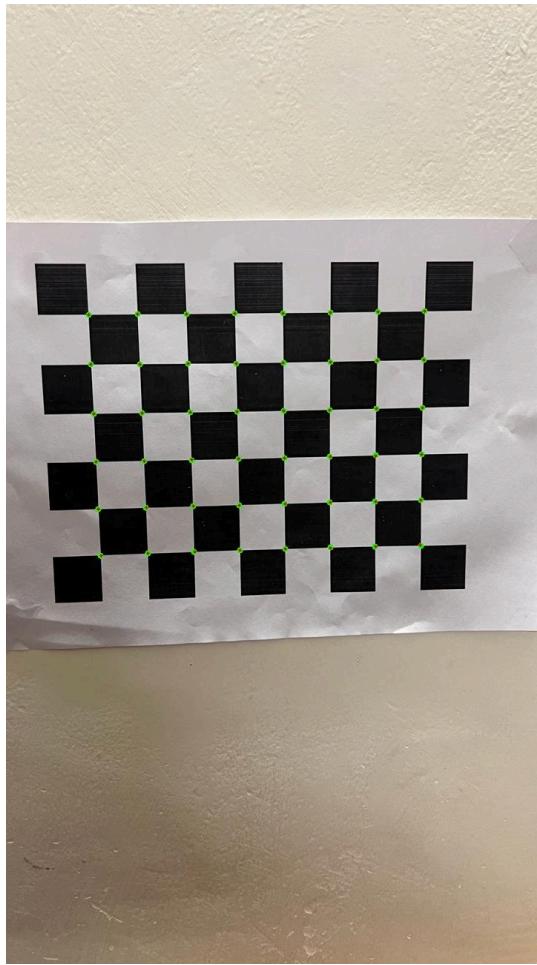
Observation

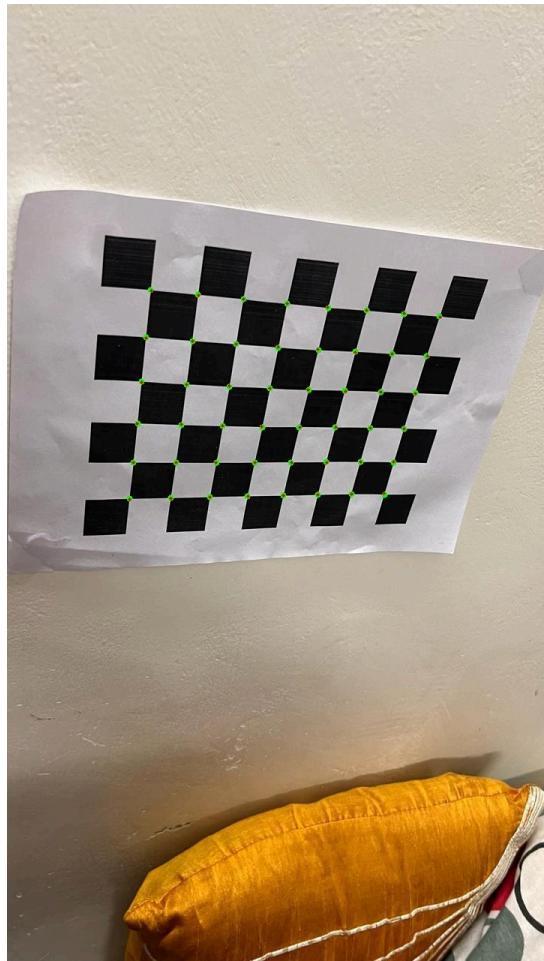
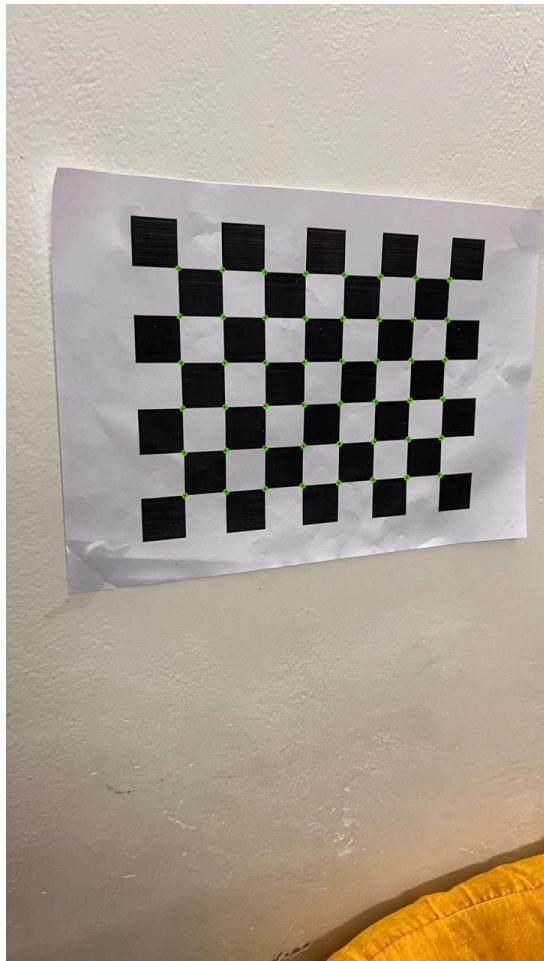
Upon examining the undistorted images:

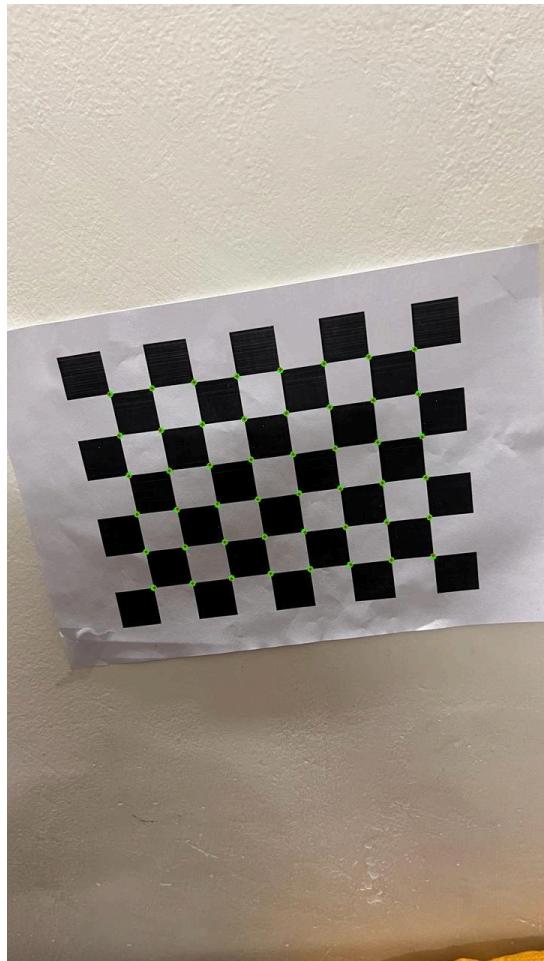
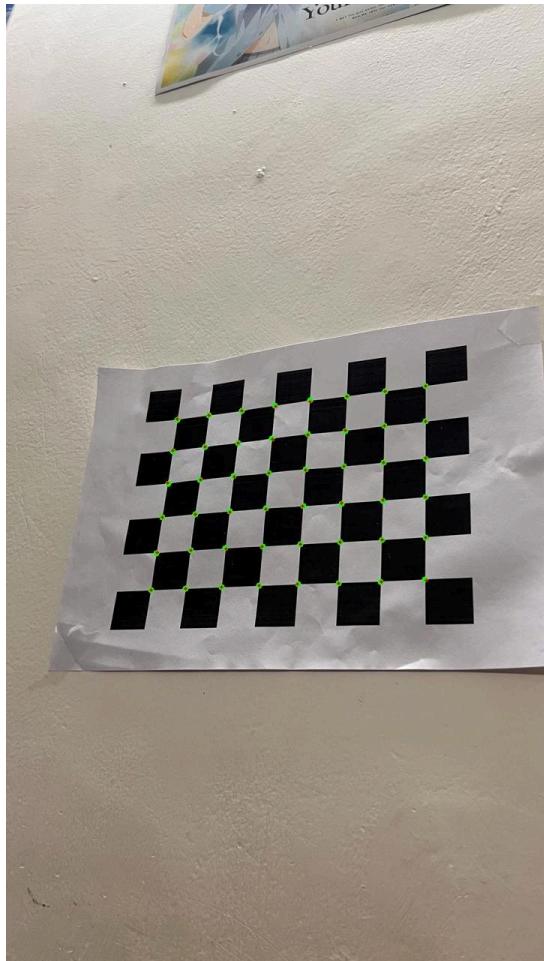
- The **curved lines near the image corners** in the raw images, caused by **barrel distortion**, were successfully straightened.
- The checkerboard squares near the edges, which appeared **warped outward**, now appear **rectangular and properly aligned**.
- This confirms that the distortion correction removed the radial distortion effectively, resulting in a more geometrically accurate image projection.

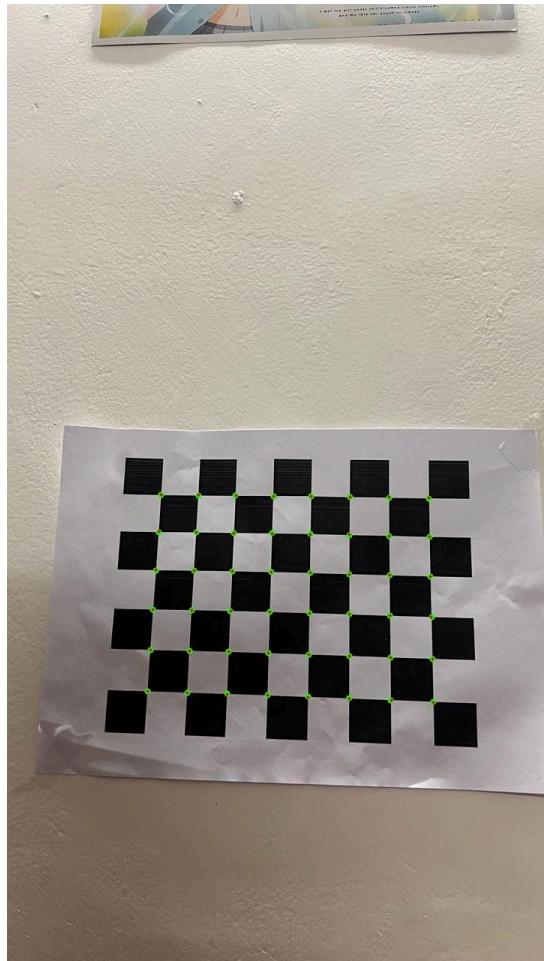
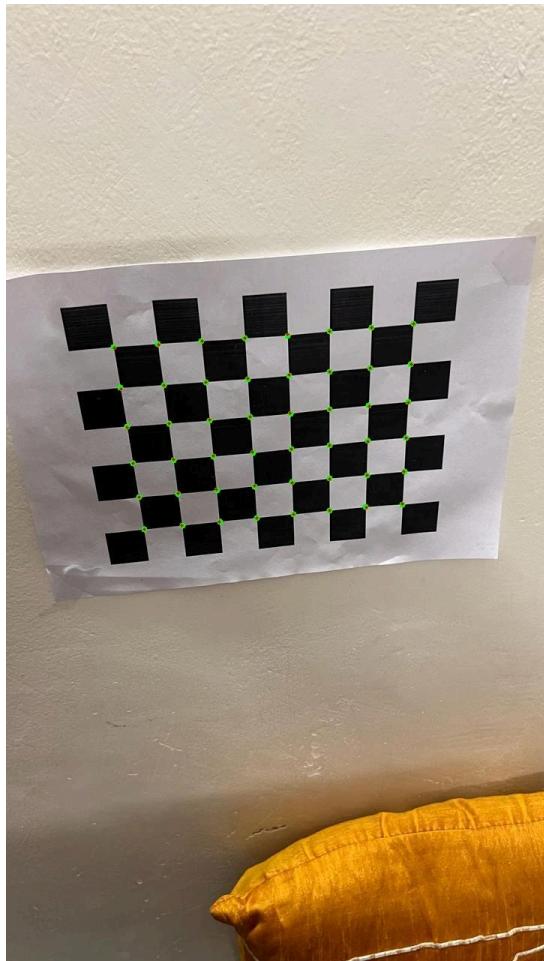
4. Reprojection Error Analysis

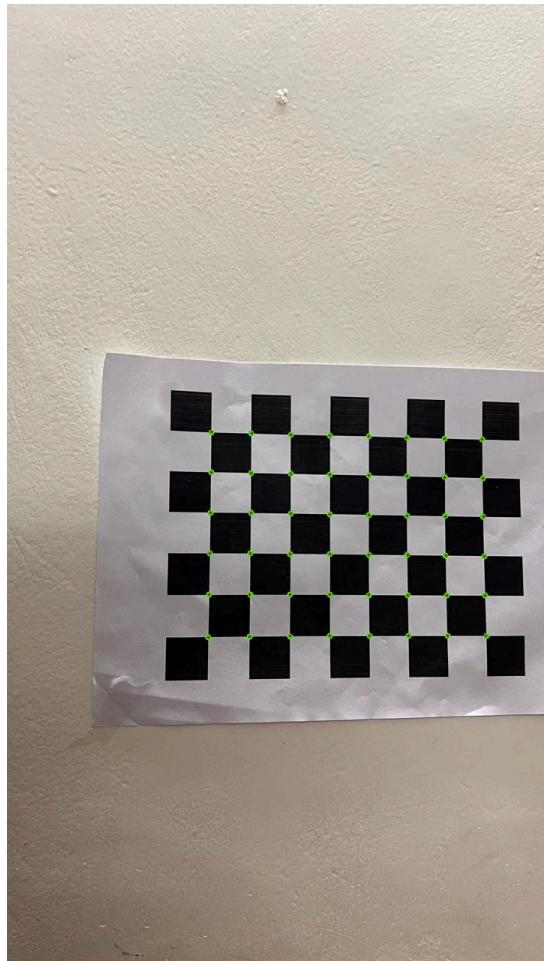
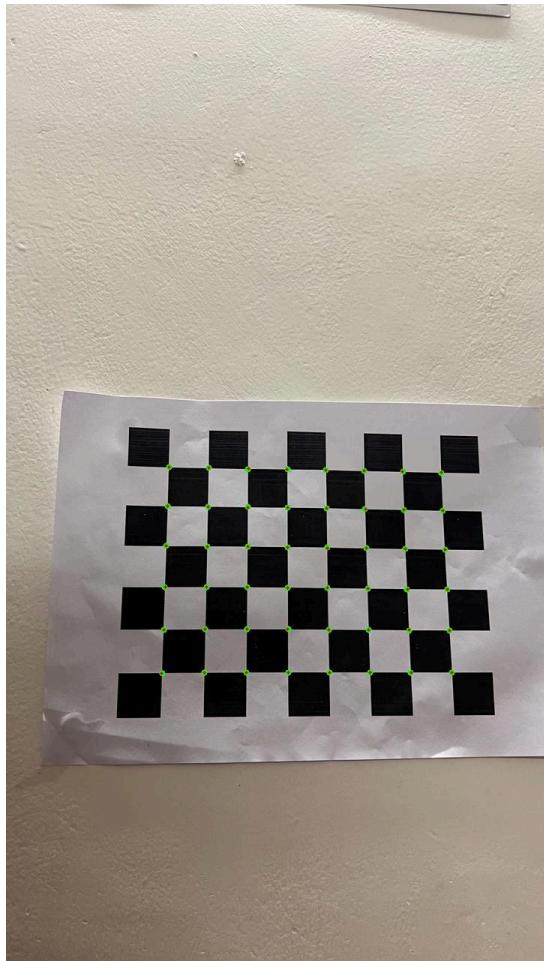
The **reprojection error** quantifies the difference between the actual detected image points and the projected image points (obtained by projecting known 3D world points using the estimated intrinsic and extrinsic parameters). It is a key metric used to evaluate the accuracy of camera calibration.

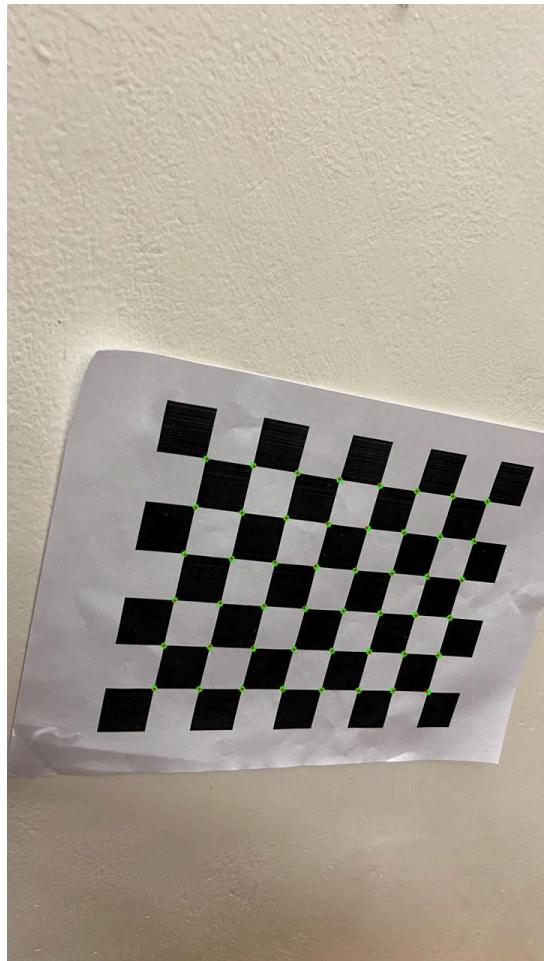
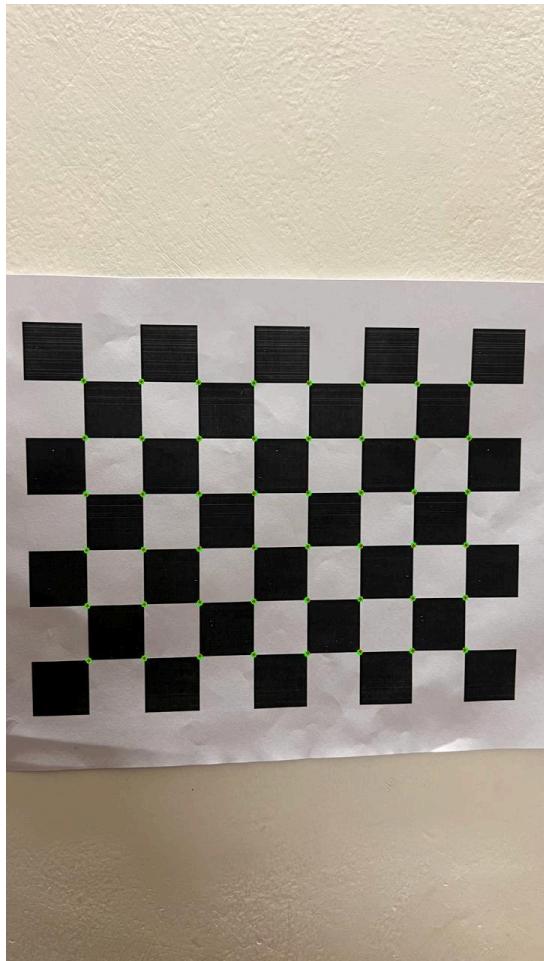


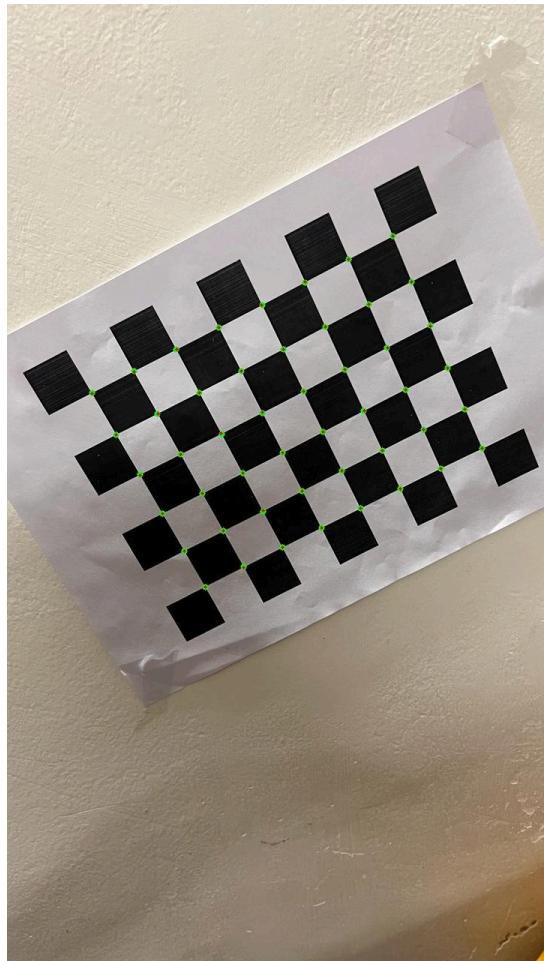
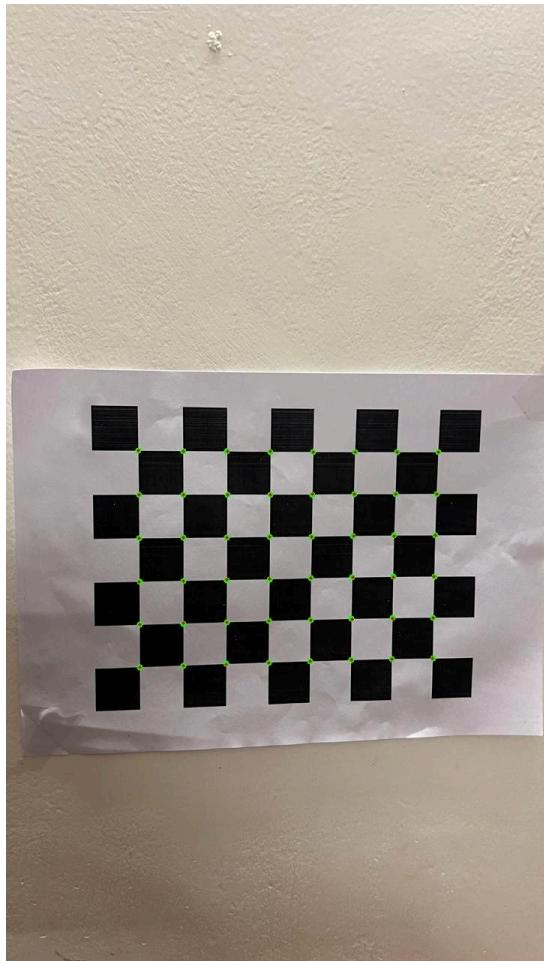


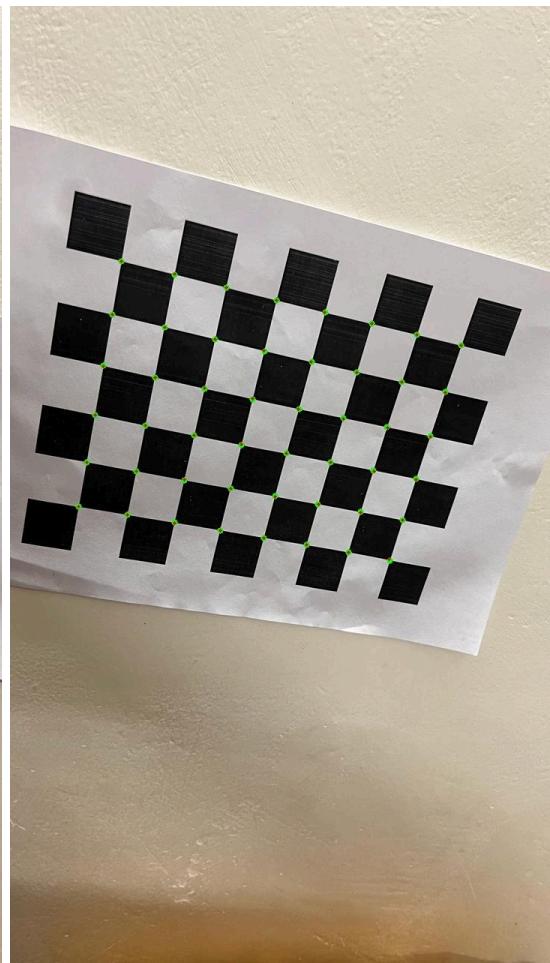
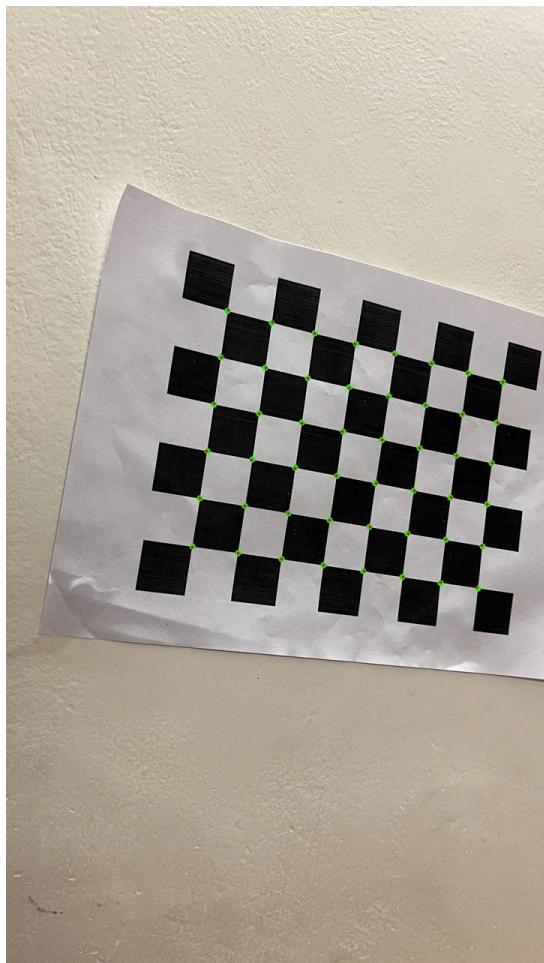


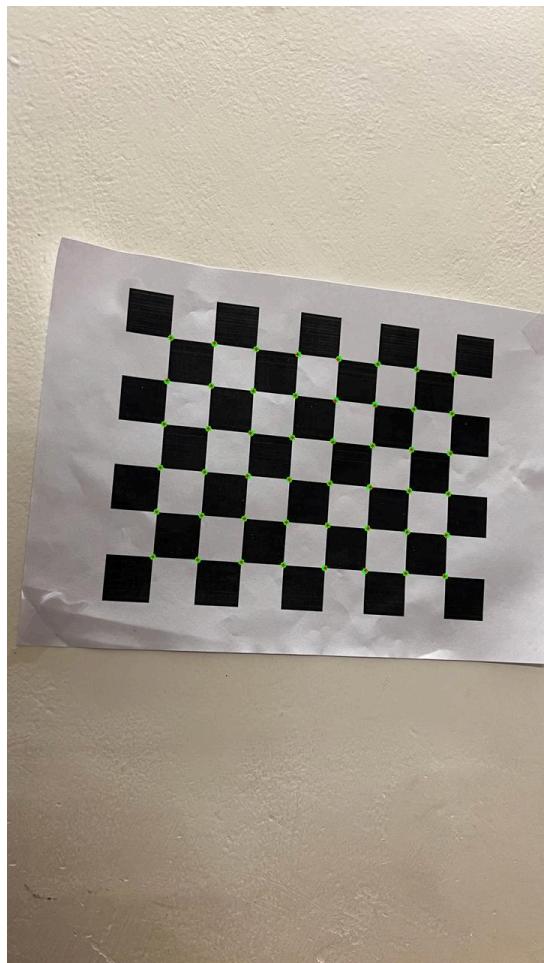
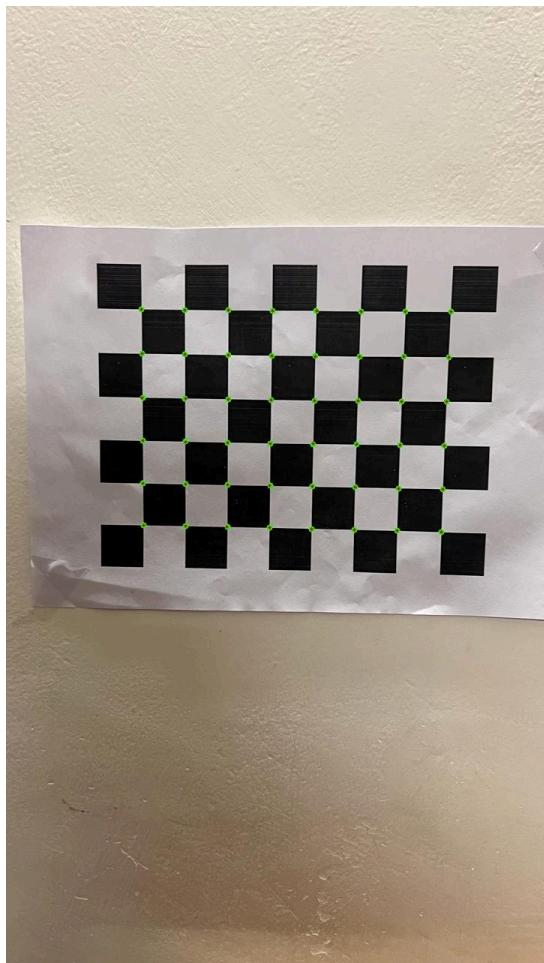


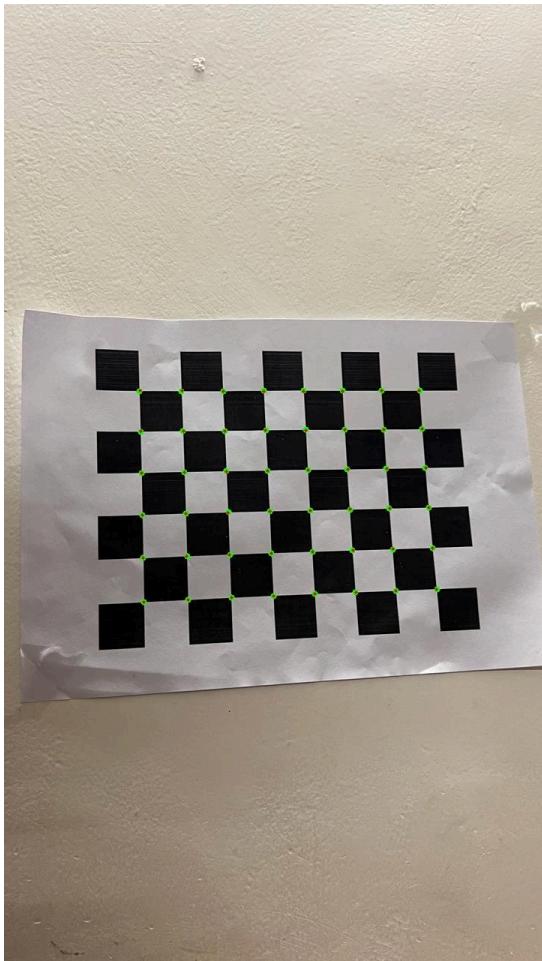
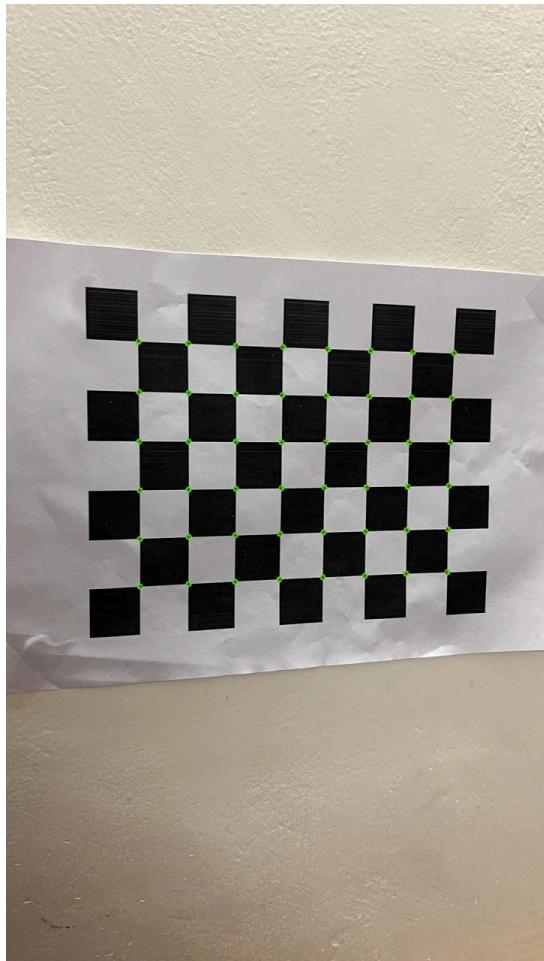


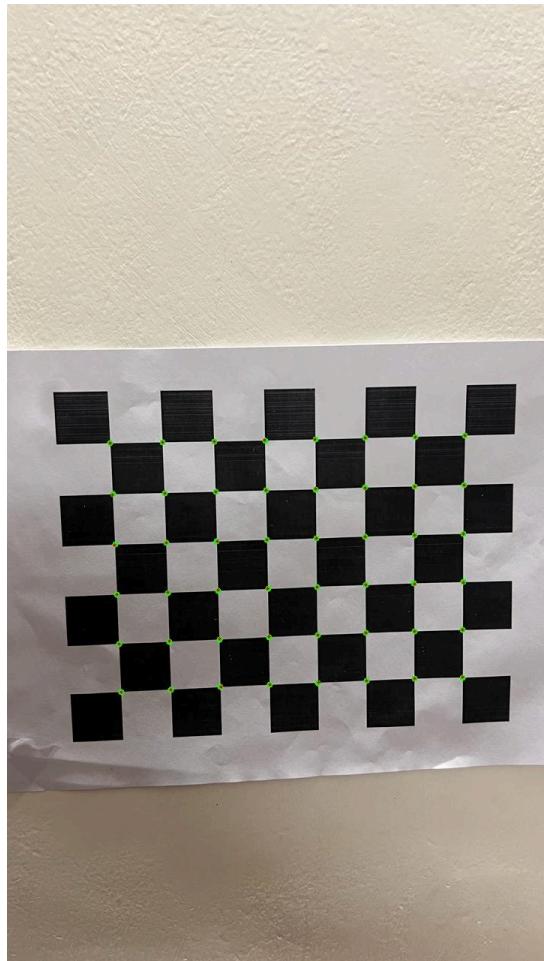
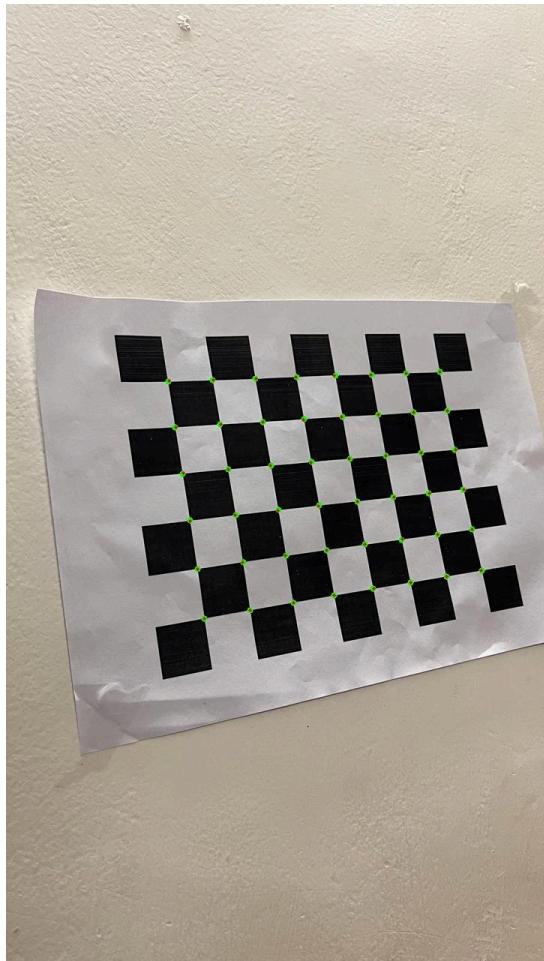


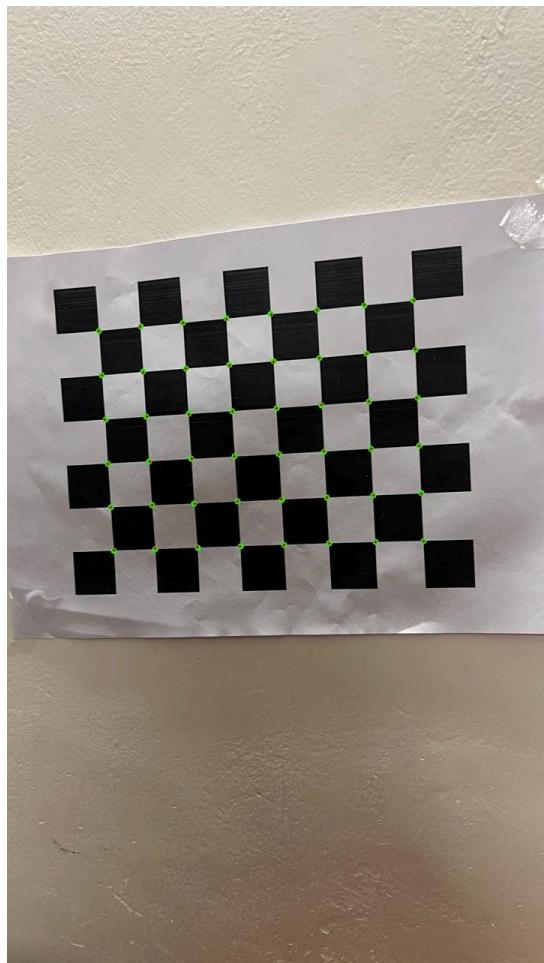
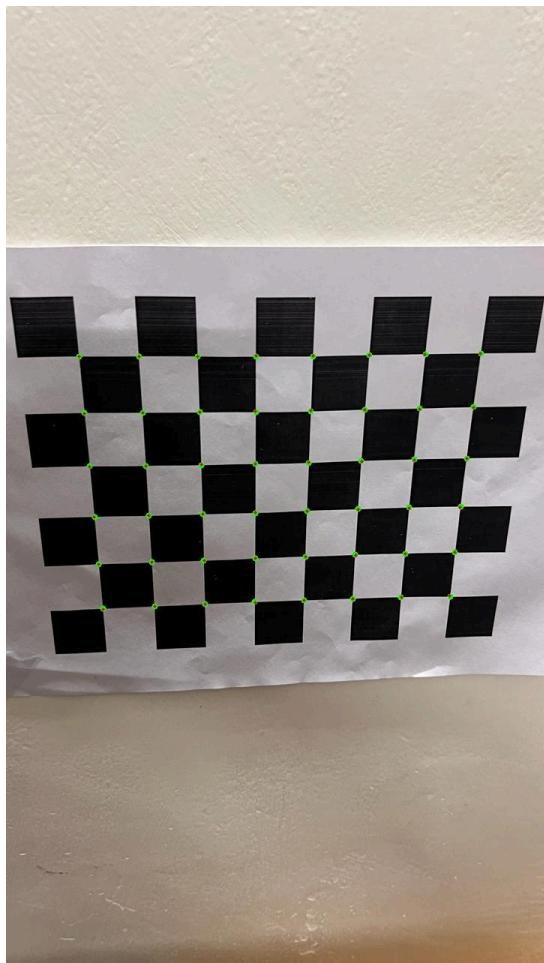


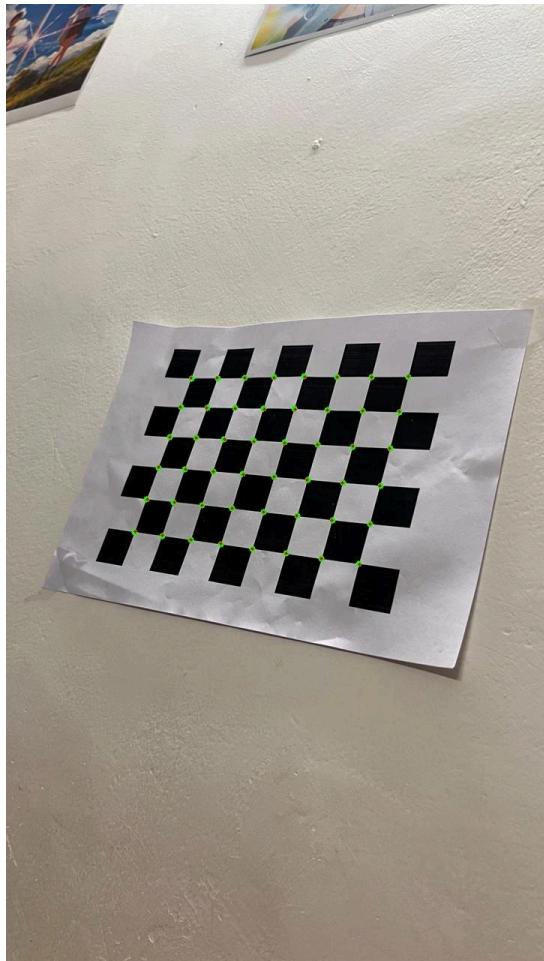






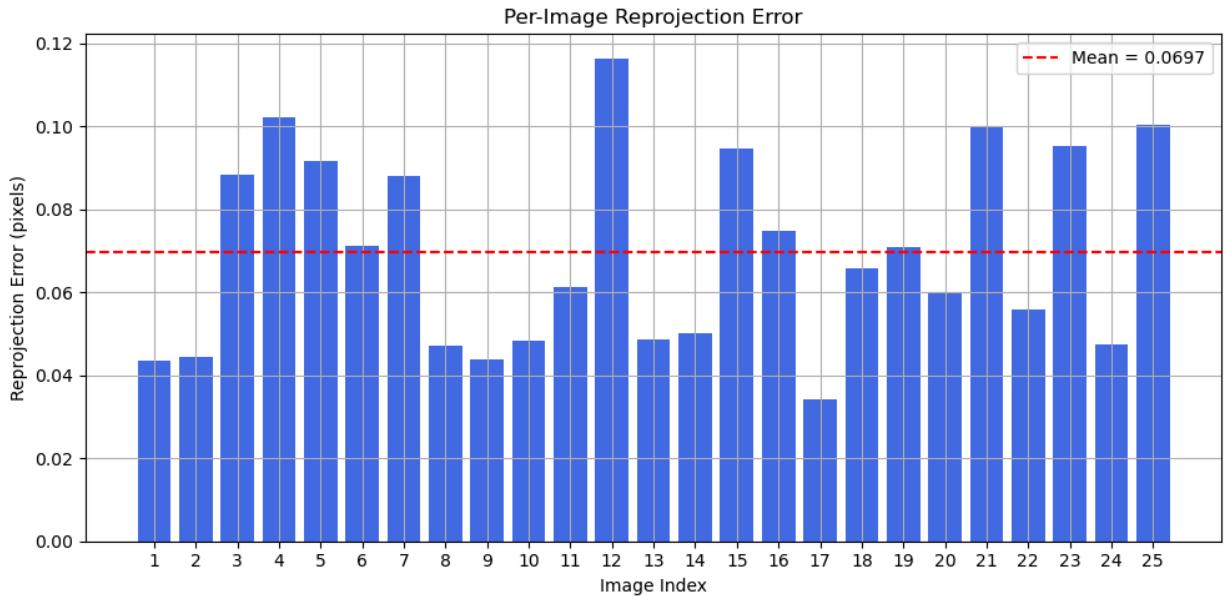






Results

- **Mean Reprojection Error:** 0.0697 pixels
- **Standard Deviation:** 0.0232 pixels



5. The **reprojection error** quantifies how accurately the estimated camera parameters (intrinsic, extrinsic, and distortion coefficients) model the actual projection of 3D points to the 2D image plane.

Computation:

For each image:

1. We **project the known 3D object points** (i.e., the 3D checkerboard corners) into the image using the camera matrix, distortion coefficients, and the pose (rotation and translation vectors).
2. We compare these **projected 2D points** to the **actual detected corner locations** in the image.
3. The reprojection error is the **average Euclidean distance** between the detected and projected points:

$$\text{Reprojection Error} = \frac{1}{N} \sum_{i=1}^N \| p_i - \hat{p}_i \|_2$$

$$\text{Reprojection Error} = \frac{1}{N} \sum_{i=1}^N \| p_i - \hat{p}_i \|_2$$
where p_i is the detected corner, and \hat{p}_i is the reprojected point.

6. Image Normal Vector

- 1 [-0.0762, -0.1487, 0.9859]
- 2 [-0.0229, 0.1238, 0.9920]
- 3 [-0.3983, -0.1872, 0.8980]

4 [-0.3464, -0.5281, 0.7753]
5 [0.3582, 0.3986, 0.8443]
6 [0.0452, -0.3252, 0.9446]
7 [-0.1278, -0.5073, 0.8523]
8 [0.0720, 0.2172, 0.9735]
9 [-0.0954, 0.1972, 0.9757]
10 [-0.1461, 0.0739, 0.9865]
11 [-0.1077, -0.0864, 0.9904]
12 [-0.3853, -0.3092, 0.8695]
13 [-0.1355, -0.1060, 0.9851]
14 [-0.0976, -0.2745, 0.9566]
15 [-0.3603, 0.0181, 0.9326]
16 [-0.1646, -0.3628, 0.9172]
17 [0.0357, -0.1105, 0.9932]
18 [-0.2233, 0.0683, 0.9724]
19 [-0.2714, -0.0962, 0.9577]
20 [-0.1658, 0.1601, 0.9731]
21 [-0.3974, 0.1908, 0.8976]
22 [-0.0833, -0.1275, 0.9883]
23 [-0.1031, -0.4017, 0.9100]
24 [0.1888, -0.0664, 0.9798]
25 [0.4546, 0.4001, 0.7958]

Custom Images:

1. **Estimated Intrinsic Camera Parameters:**
Focal Lengths:

- $f_x = 1306.50$ pixels
- $f_y = 1319.42$ pixels

Principal Point:

- $c_x = 631.95$ pixels
- $c_y = 531.41$ pixels

Skew Parameter:

- 0.00 (indicating no skew between image axes)

Intrinsic Matrix (K):

```
[[1.30650443e+03 0.00000000e+00 6.31947651e+02]
 [0.00000000e+00 1.31941647e+03 5.31411939e+02]
 [0.00000000e+00 0.00000000e+00 1.00000000e+00]]
```

Reprojection Error:

- 1.8691 pixels (indicating the average discrepancy between observed and projected points)

2. **Estimated Extrinsic Camera Parameters:**

Image 1:

Rotation Matrix (R):

```
[[ 0.07173739  0.99658906  0.04079217]
 [-0.84302223  0.038724   0.53648296]]
```

[0.53307341 -0.0728746 0.84292469]]

Translation Vector (t):

[-48.40408507]

[120.3335496]

[455.45193068]]

Image 2:

Rotation Matrix (R):

[[0.74750942 0.66145545 -0.06087978]

[-0.50668867 0.6270634 0.59165706]

[0.42953027 -0.41142214 0.80388779]]

Translation Vector (t):

[-137.16624448]

[158.24317515]

[608.50635196]]

3. Estimated Radial Distortion Coefficients:

- Distortion Coefficients:

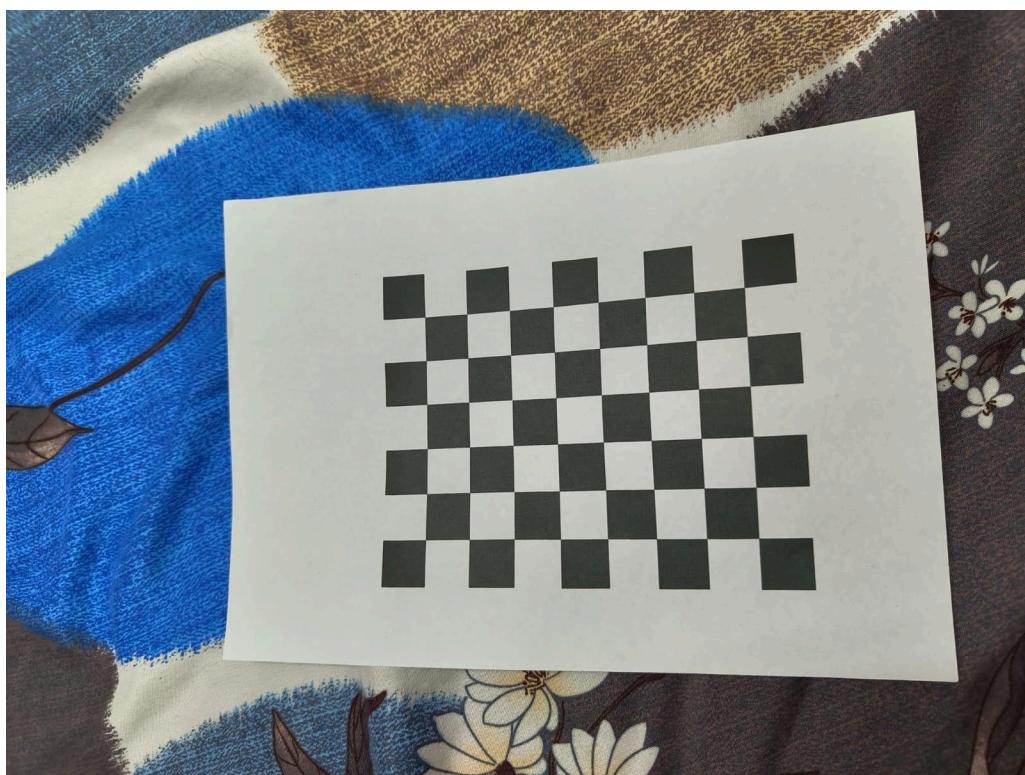
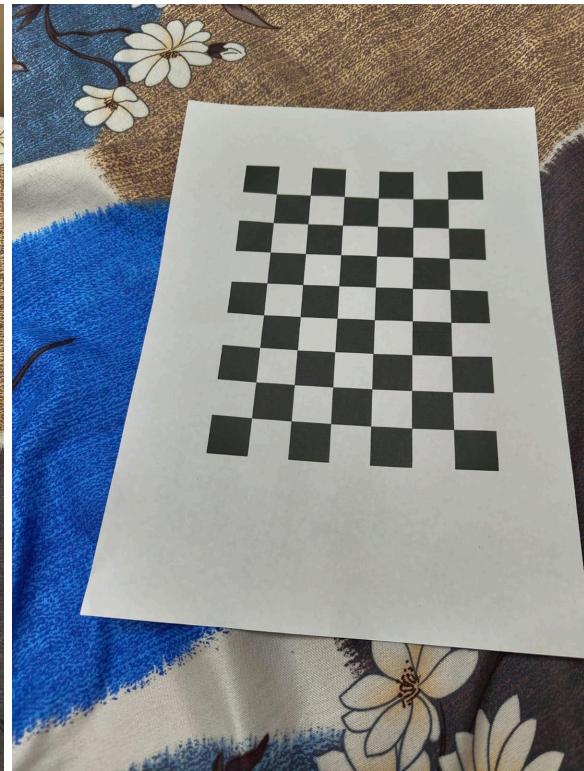
[-0.23592874 0.67182843 -0.02704358 0.00346771 -0.70792782]]

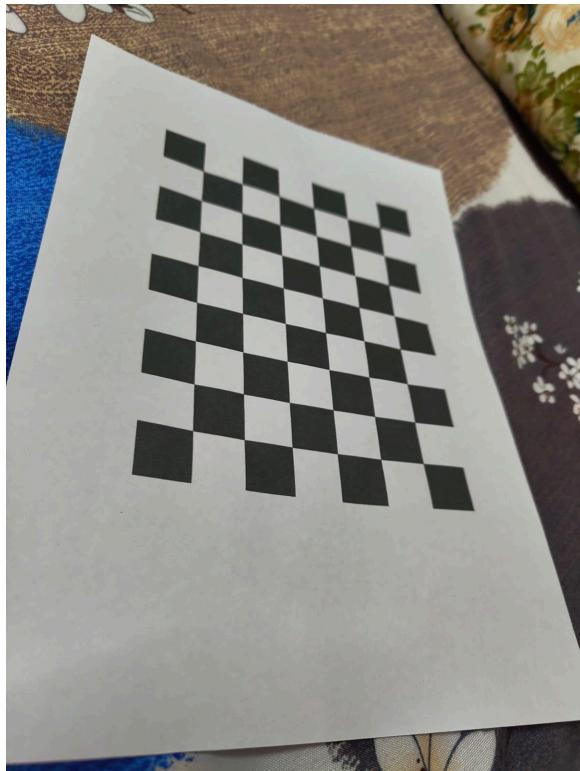
where:

- $k_1 = -0.2359$
- $k_2 = 0.6718$
- $p_1 = -0.0270$

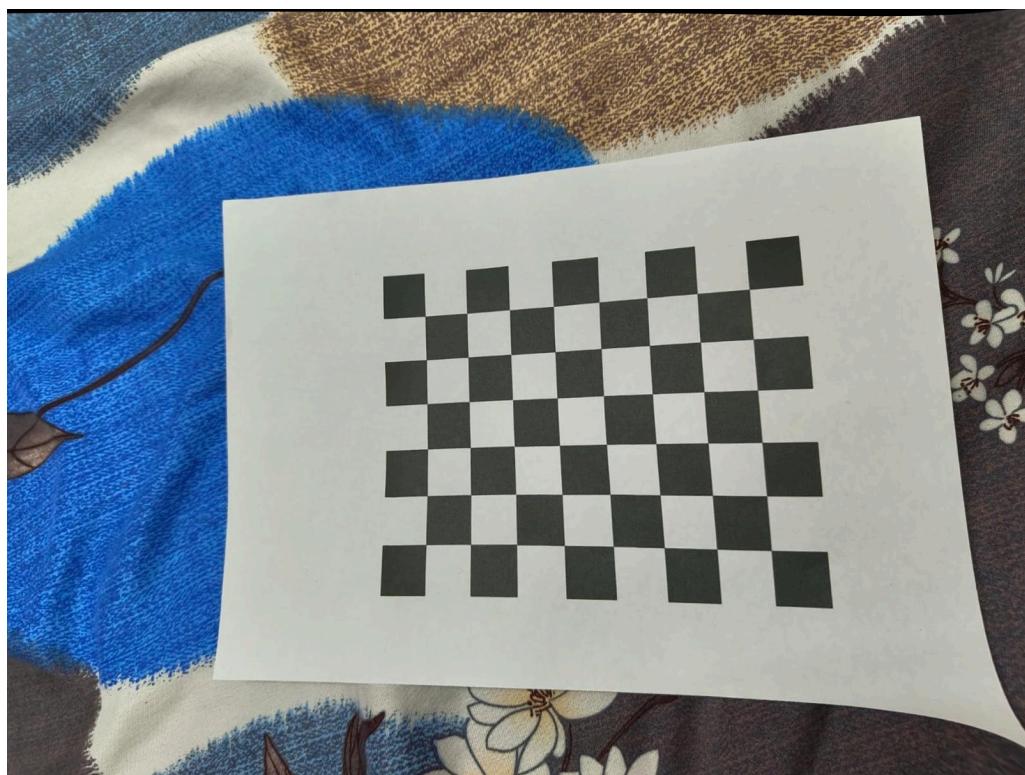
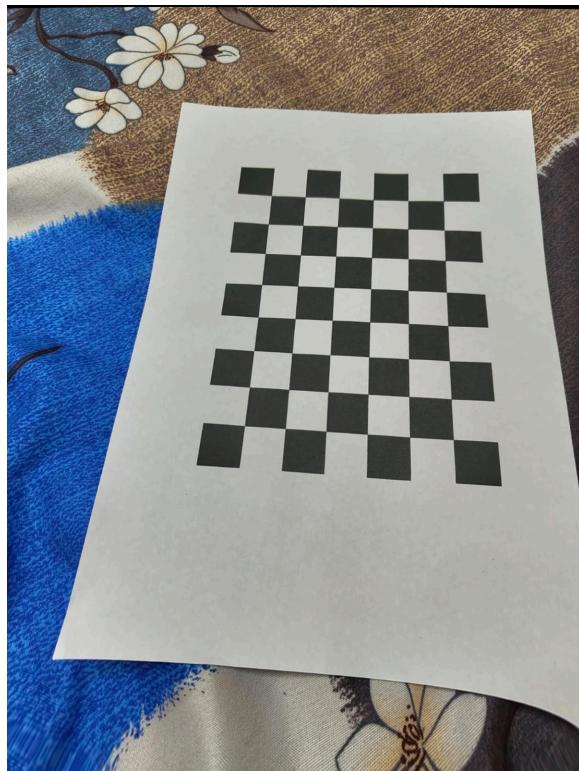
- $p_2=0.0035$
 $p_2 = 0.0035$
- $k_3=-0.7079$
 $k_3 = -0.7079$

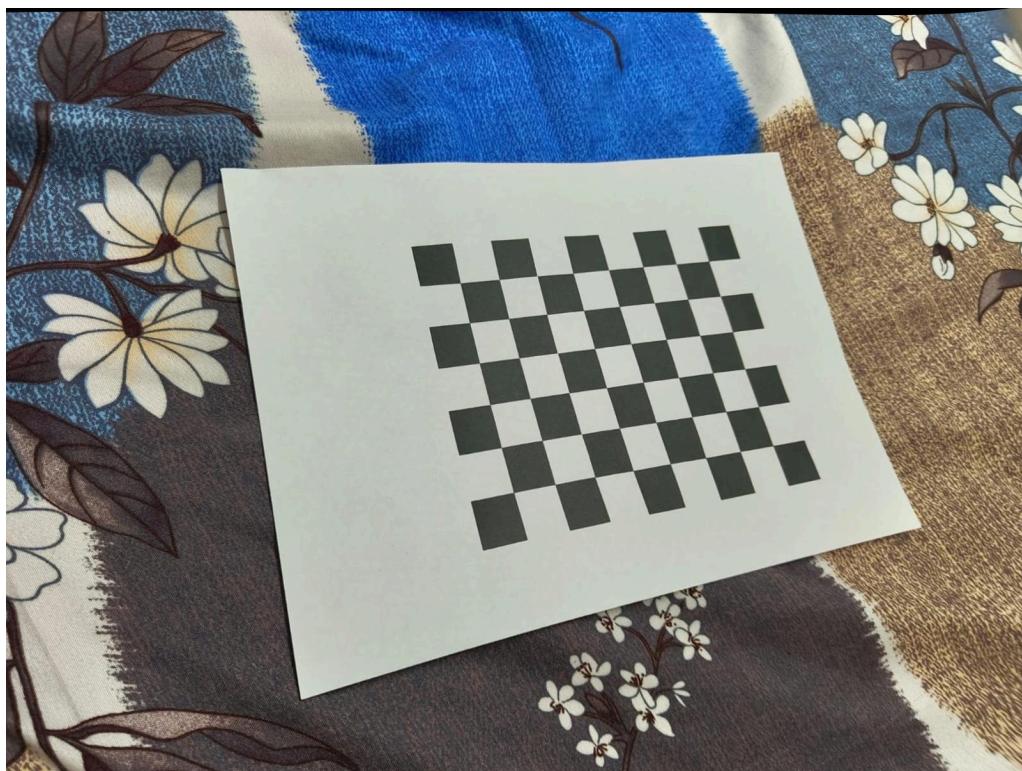
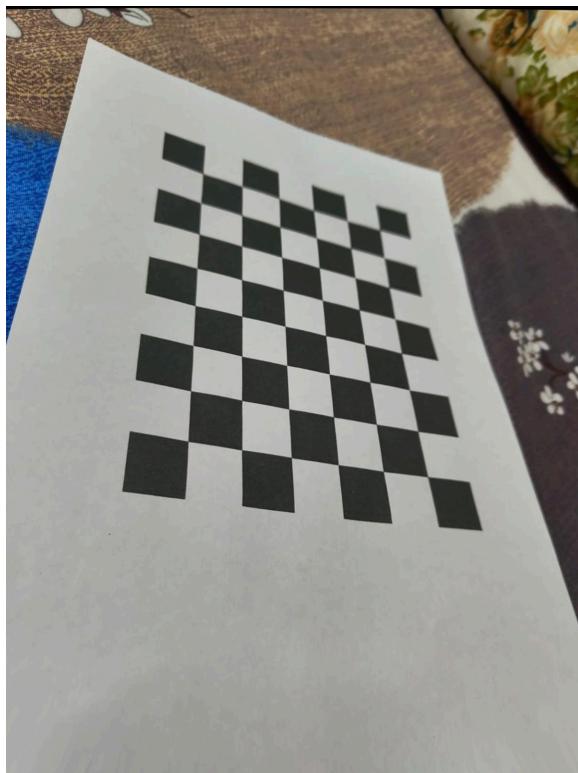
Raw Images:





Undistorted Images:





Observation:

After applying the distortion correction:

4. The **checkerboard corners** near the image edges appear **straighter** and more **uniformly spaced**.
5. In the **raw images**, slight curvature or warping was visible along the edges — especially near the corners — due to radial distortion.
6. In the **undistorted images**, this effect is **corrected**, and the edges appear more linear, indicating that the **camera calibration and distortion correction were effective**.

4. Image 01 - Reprojection Error: 0.3172 pixels

Image 02 - Reprojection Error: 0.1248 pixels

Image 03 - Reprojection Error: 0.3472 pixels

Image 04 - Reprojection Error: 0.3203 pixels

Image 05 - Reprojection Error: 0.1732 pixels

Image 06 - Reprojection Error: 0.2241 pixels

Image 07 - Reprojection Error: 0.2066 pixels

Image 08 - Reprojection Error: 0.2184 pixels

Image 09 - Reprojection Error: 0.2650 pixels

Image 10 - Reprojection Error: 0.2567 pixels

Image 11 - Reprojection Error: 0.2168 pixels

Image 12 - Reprojection Error: 0.1763 pixels

Image 13 - Reprojection Error: 0.2654 pixels

Image 14 - Reprojection Error: 0.2183 pixels

Image 15 - Reprojection Error: 0.1676 pixels

Image 16 - Reprojection Error: 0.2241 pixels

Image 17 - Reprojection Error: 0.2406 pixels

Image 18 - Reprojection Error: 0.1930 pixels

Image 19 - Reprojection Error: 0.6814 pixels

Image 20 - Reprojection Error: 0.1401 pixels

Image 21 - Reprojection Error: 0.1874 pixels

Image 22 - Reprojection Error: 0.2546 pixels

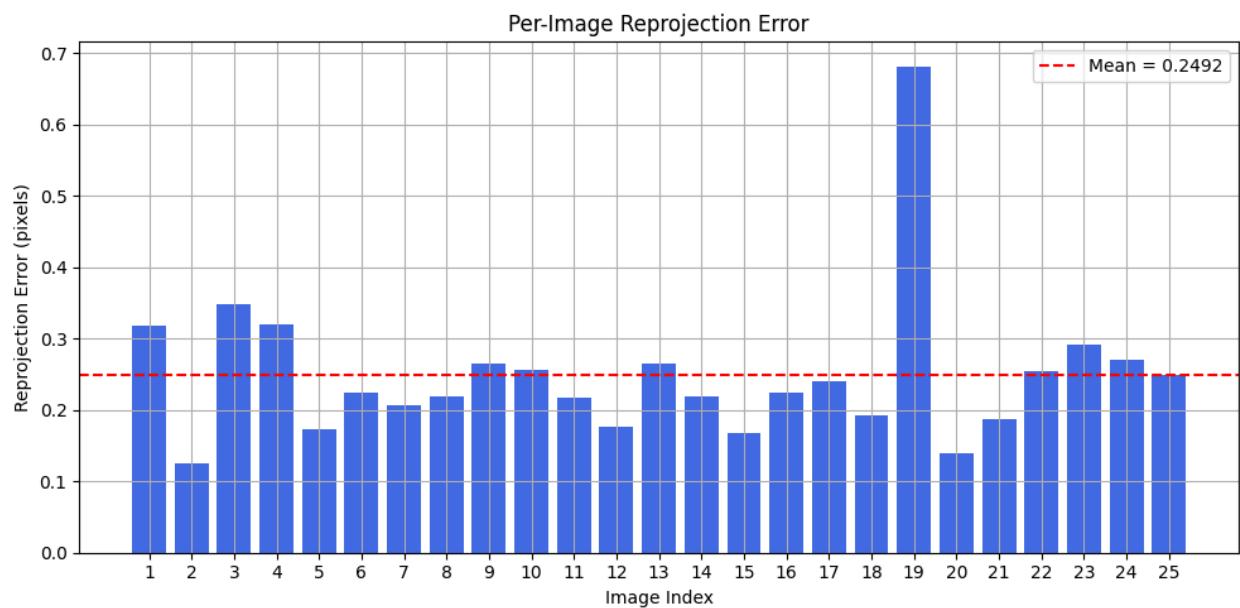
Image 23 - Reprojection Error: 0.2913 pixels

Image 24 - Reprojection Error: 0.2705 pixels

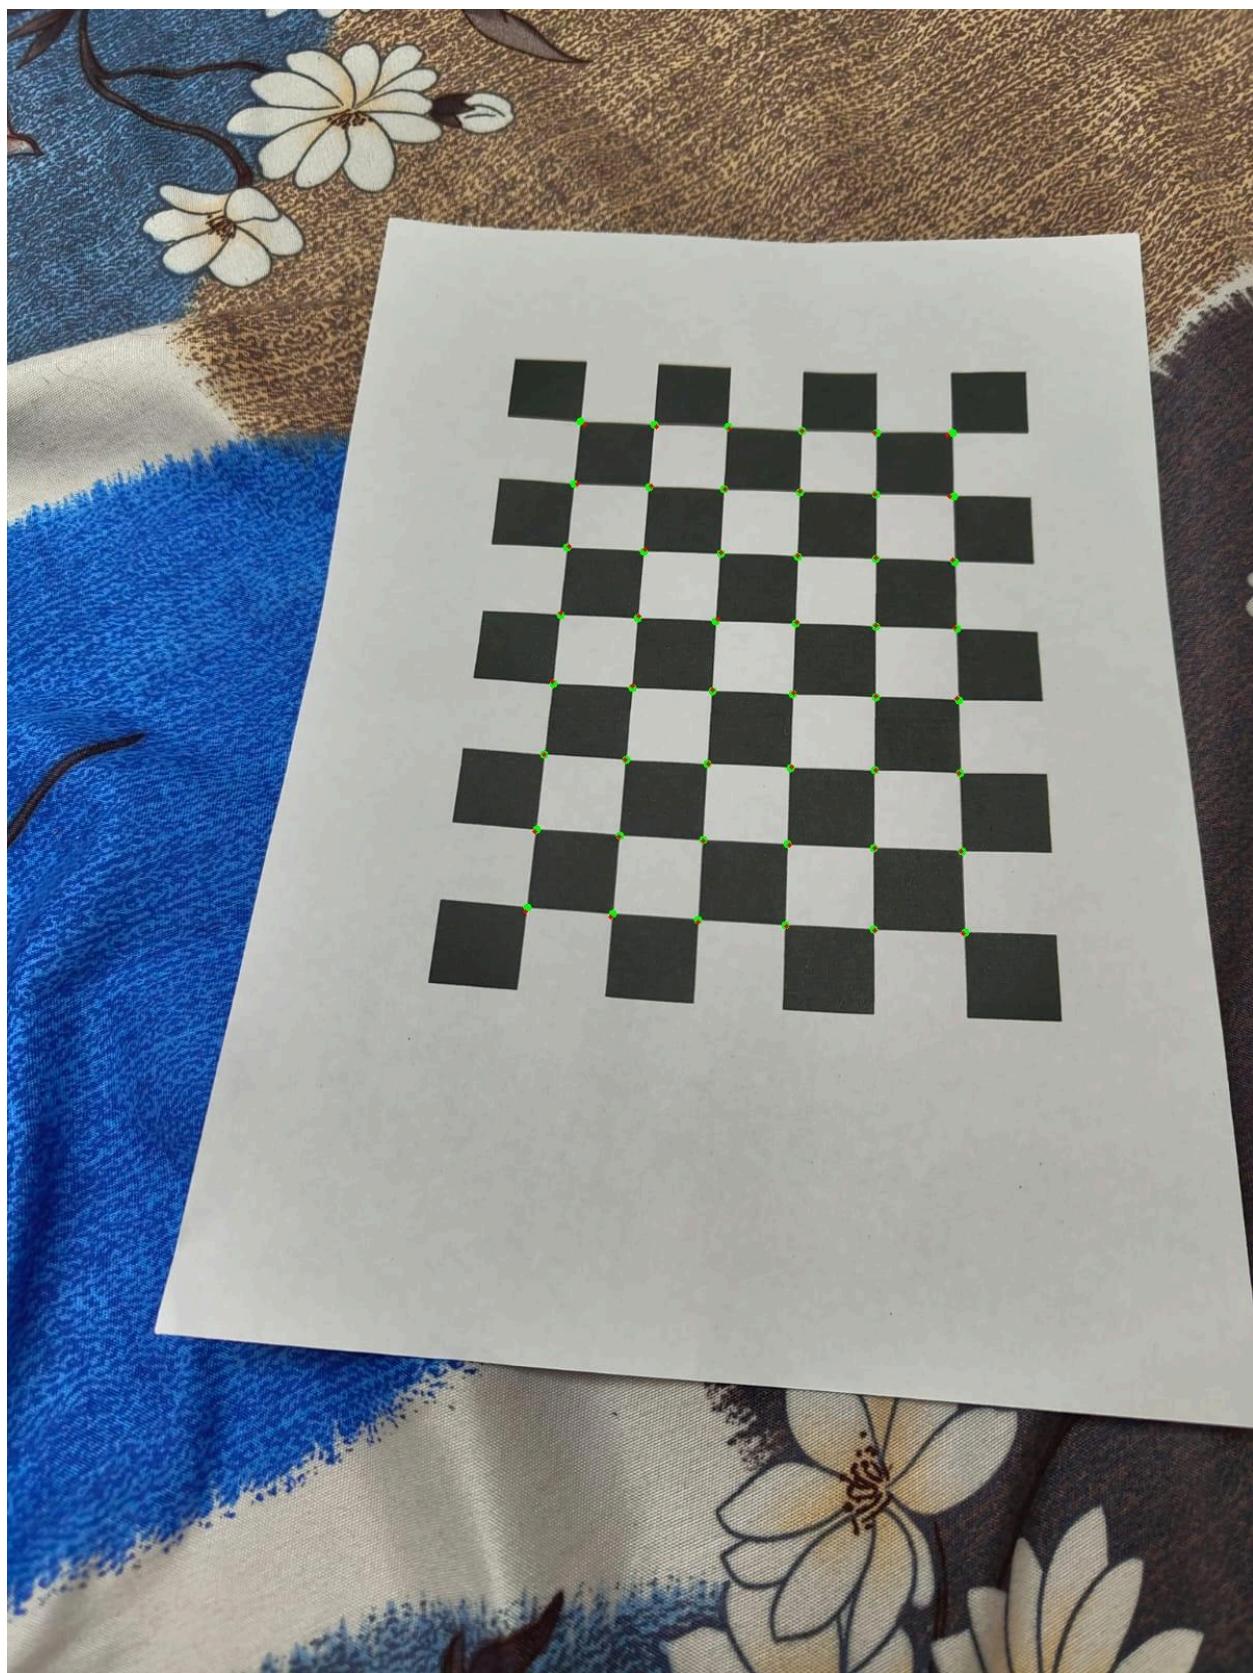
Image 25 - Reprojection Error: 0.2485 pixels

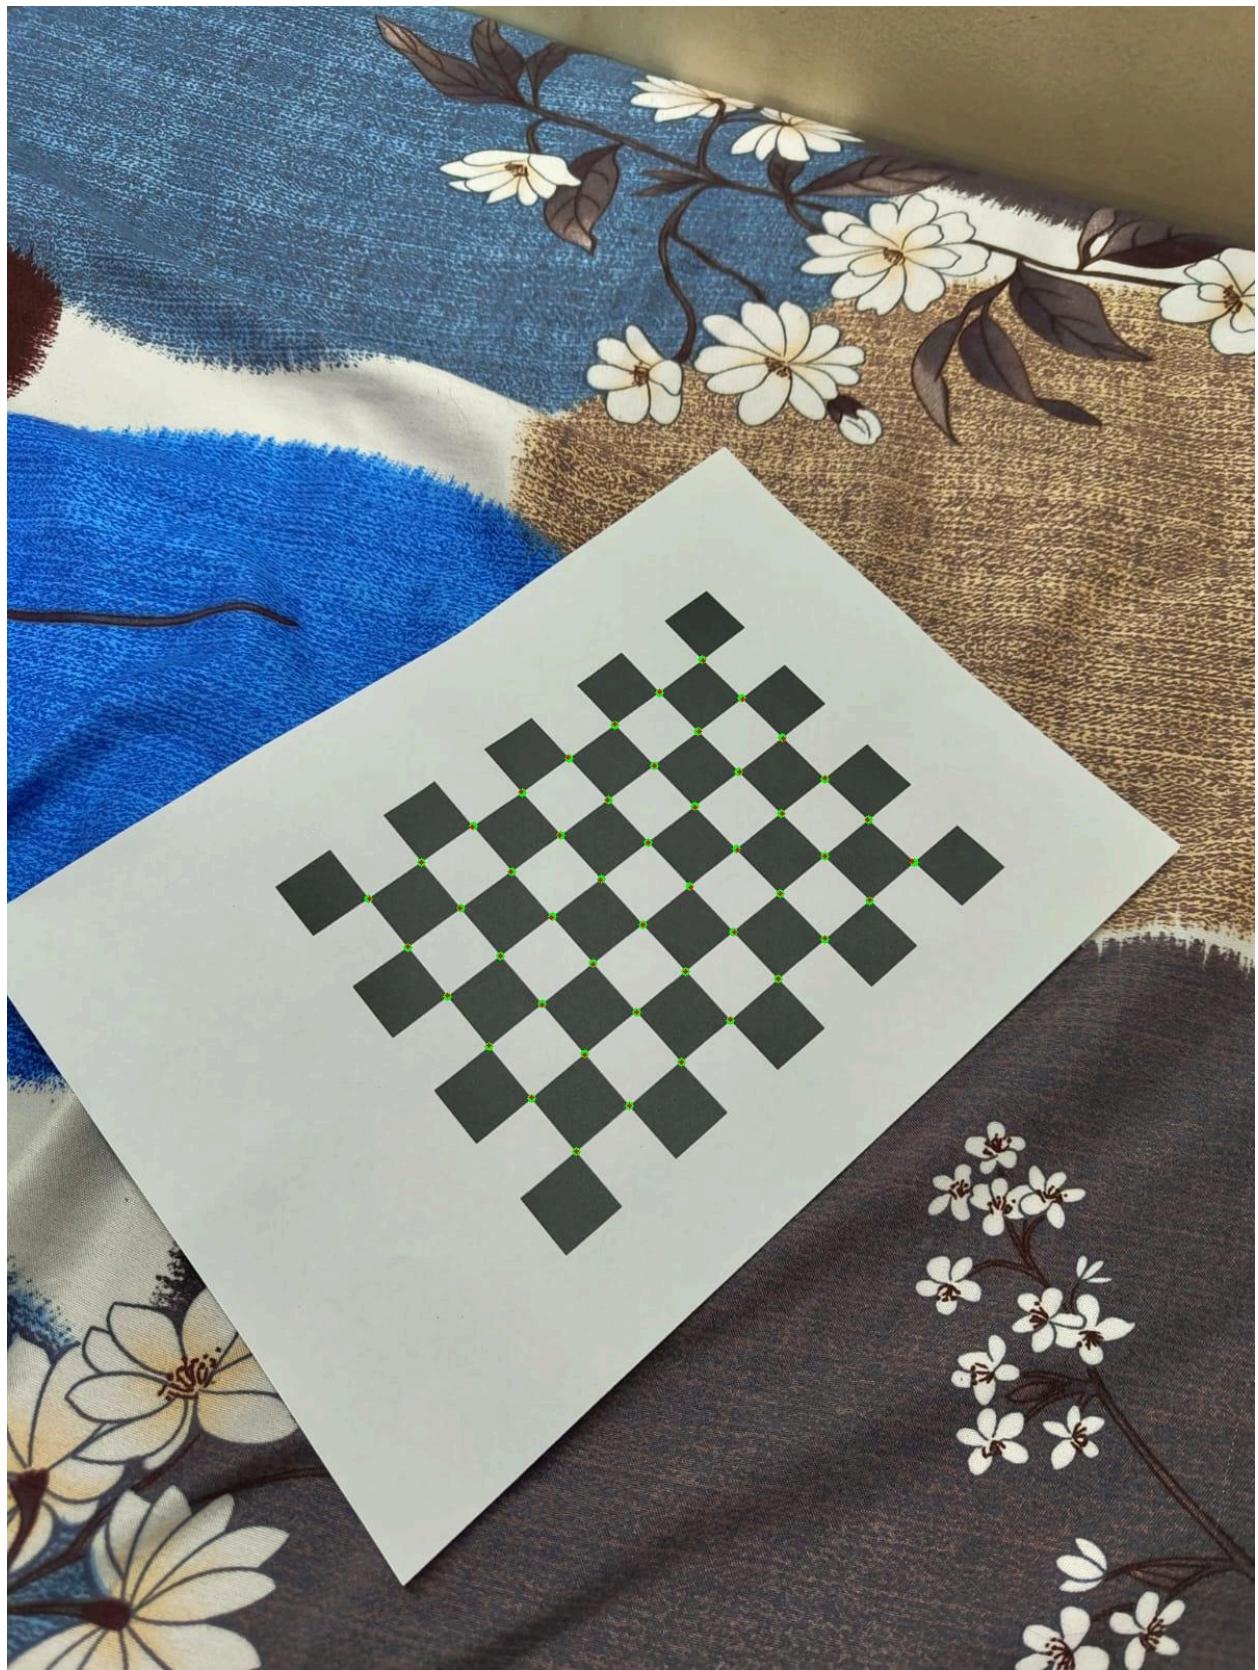
Mean Reprojection Error: 0.2492 pixels

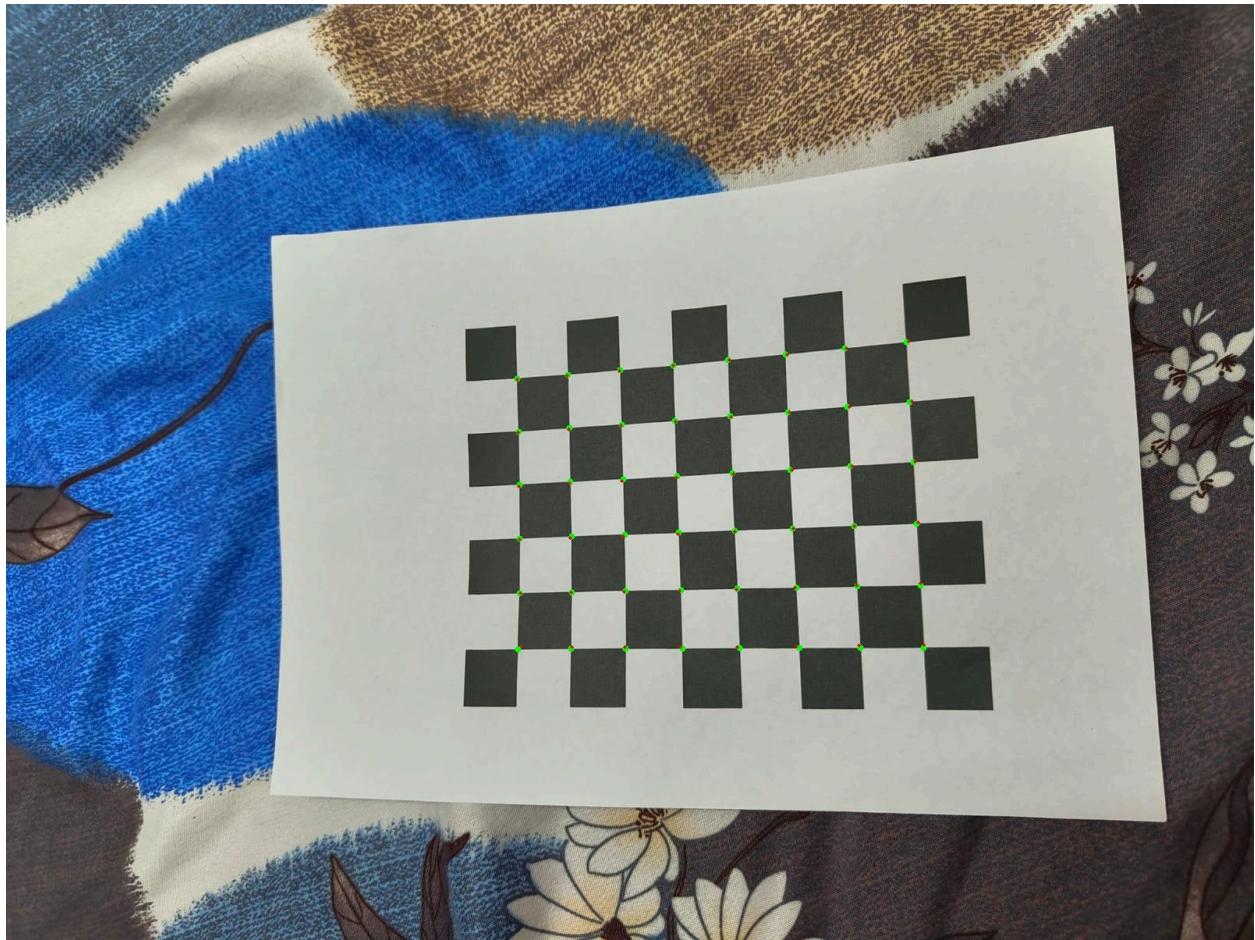
Standard Deviation: 0.1034 pixels

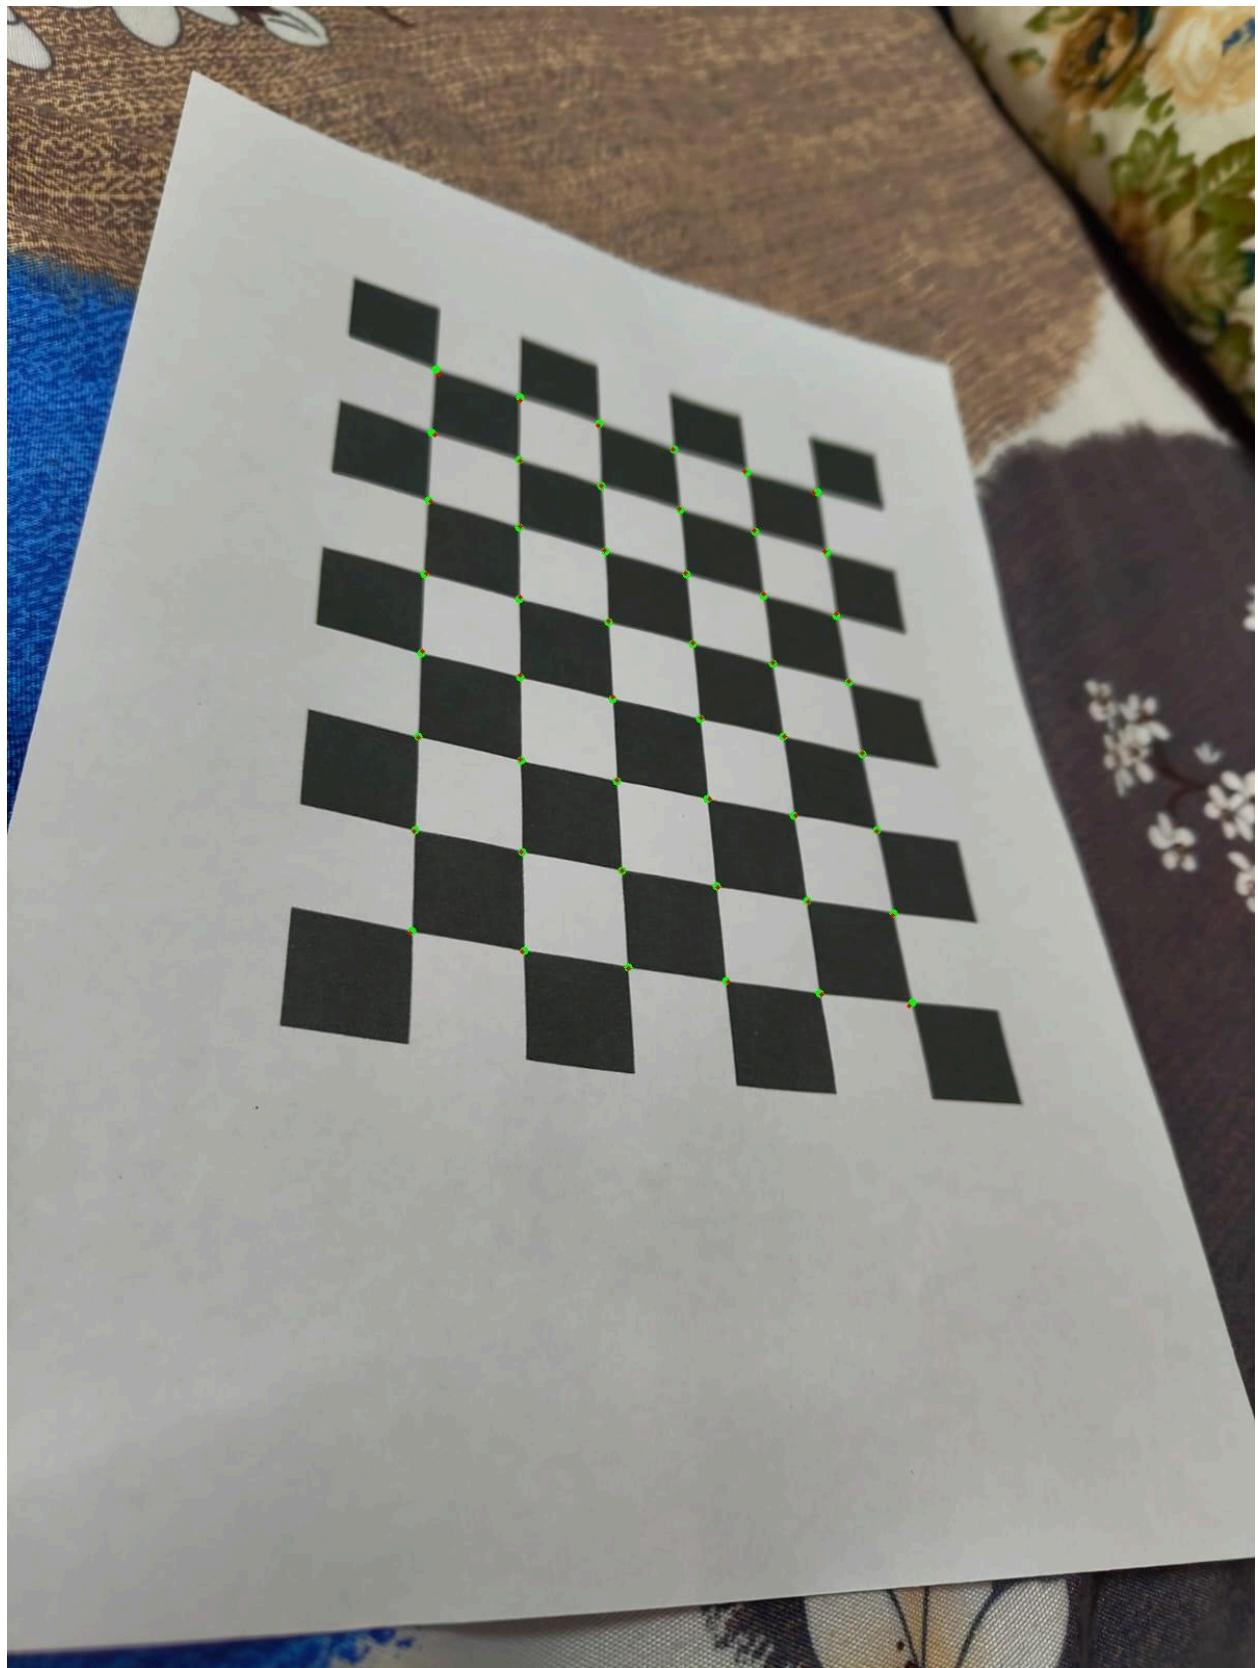


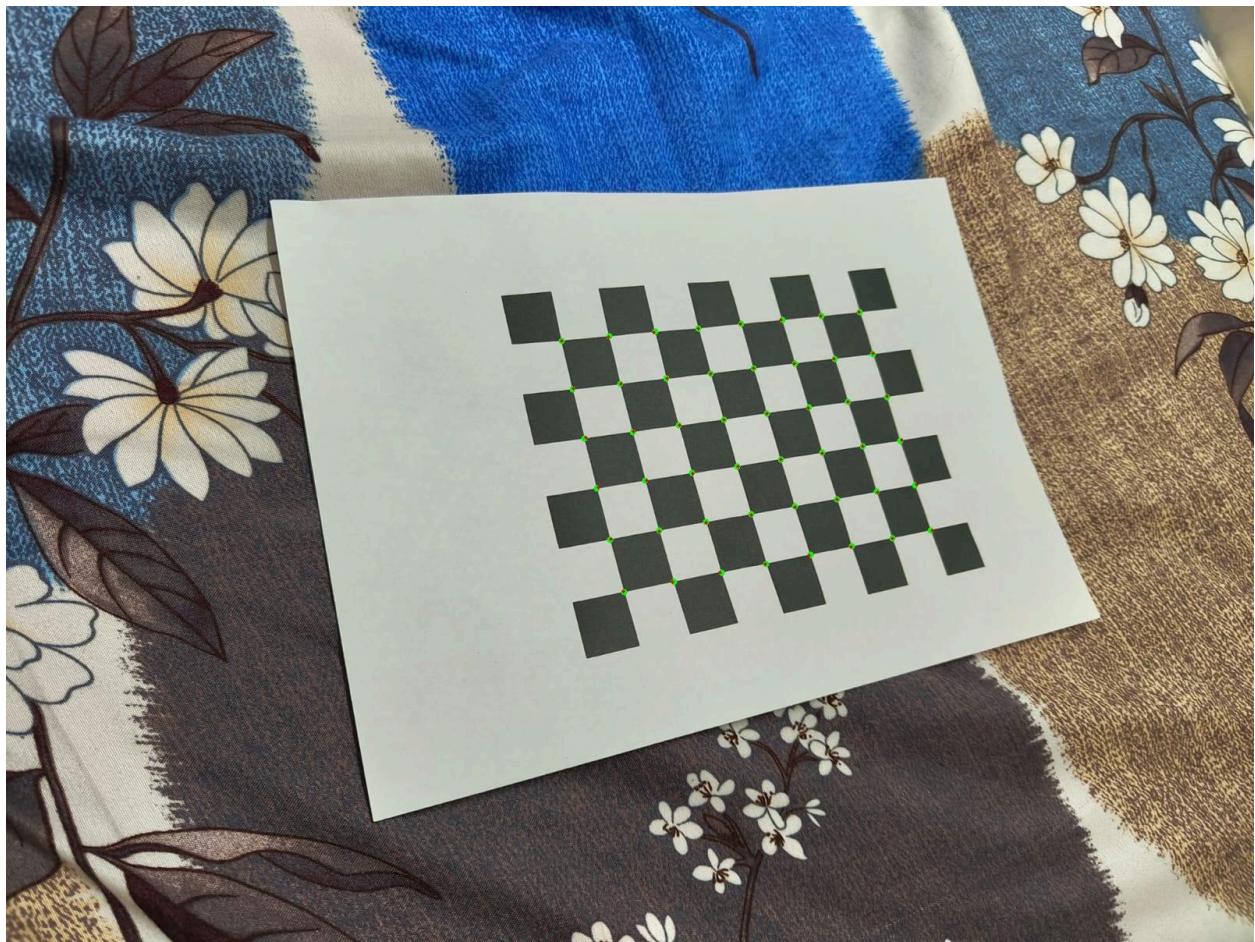
5.



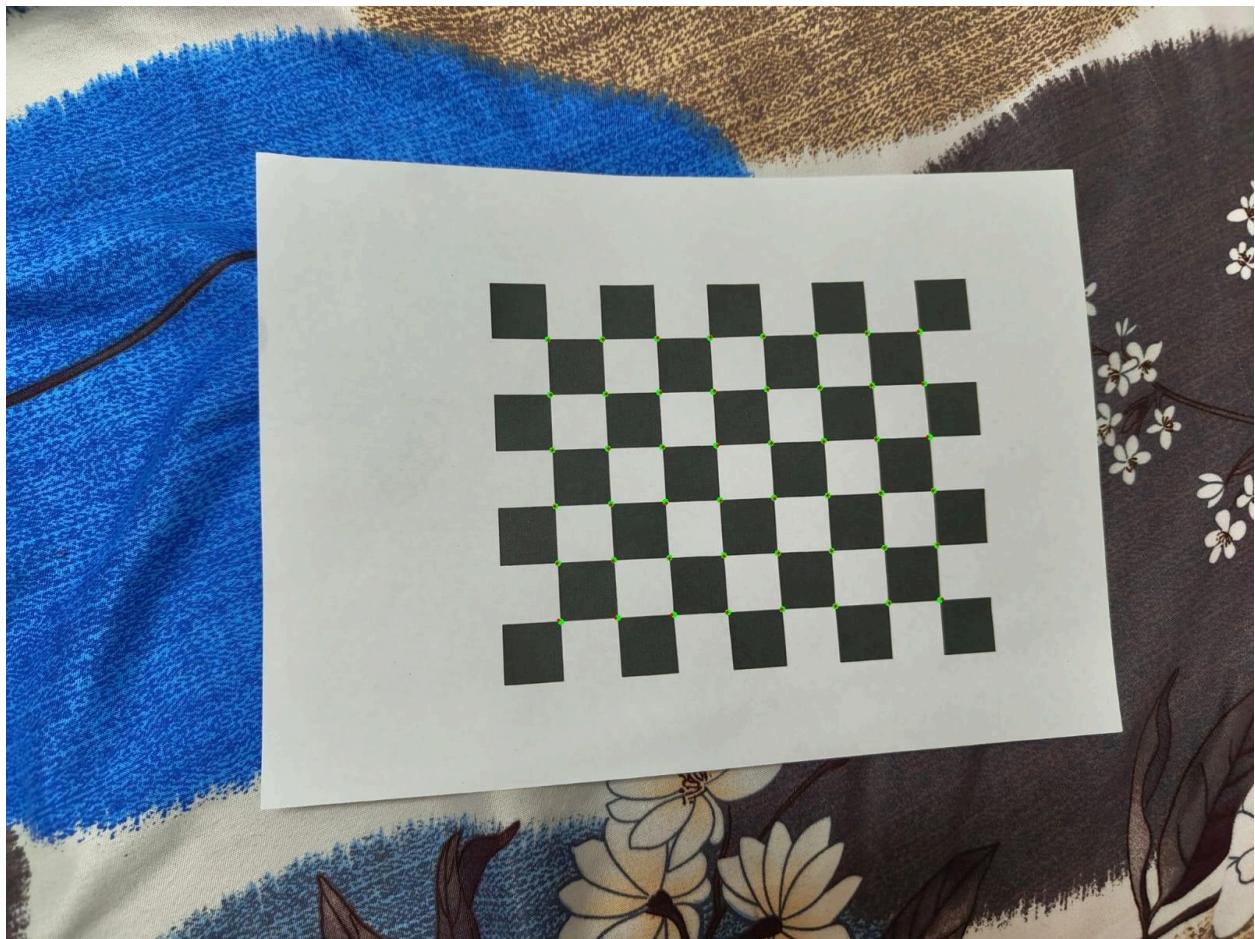






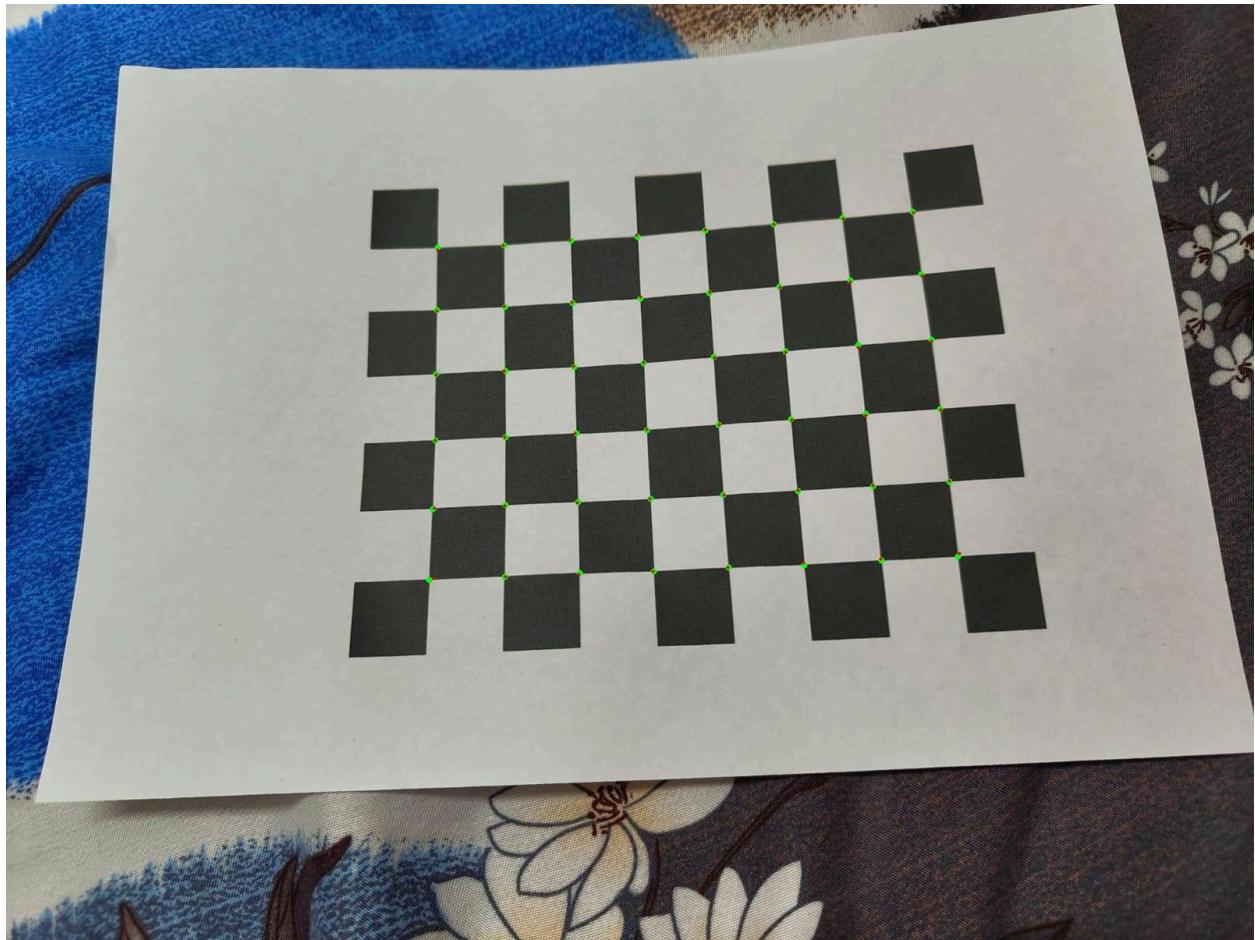




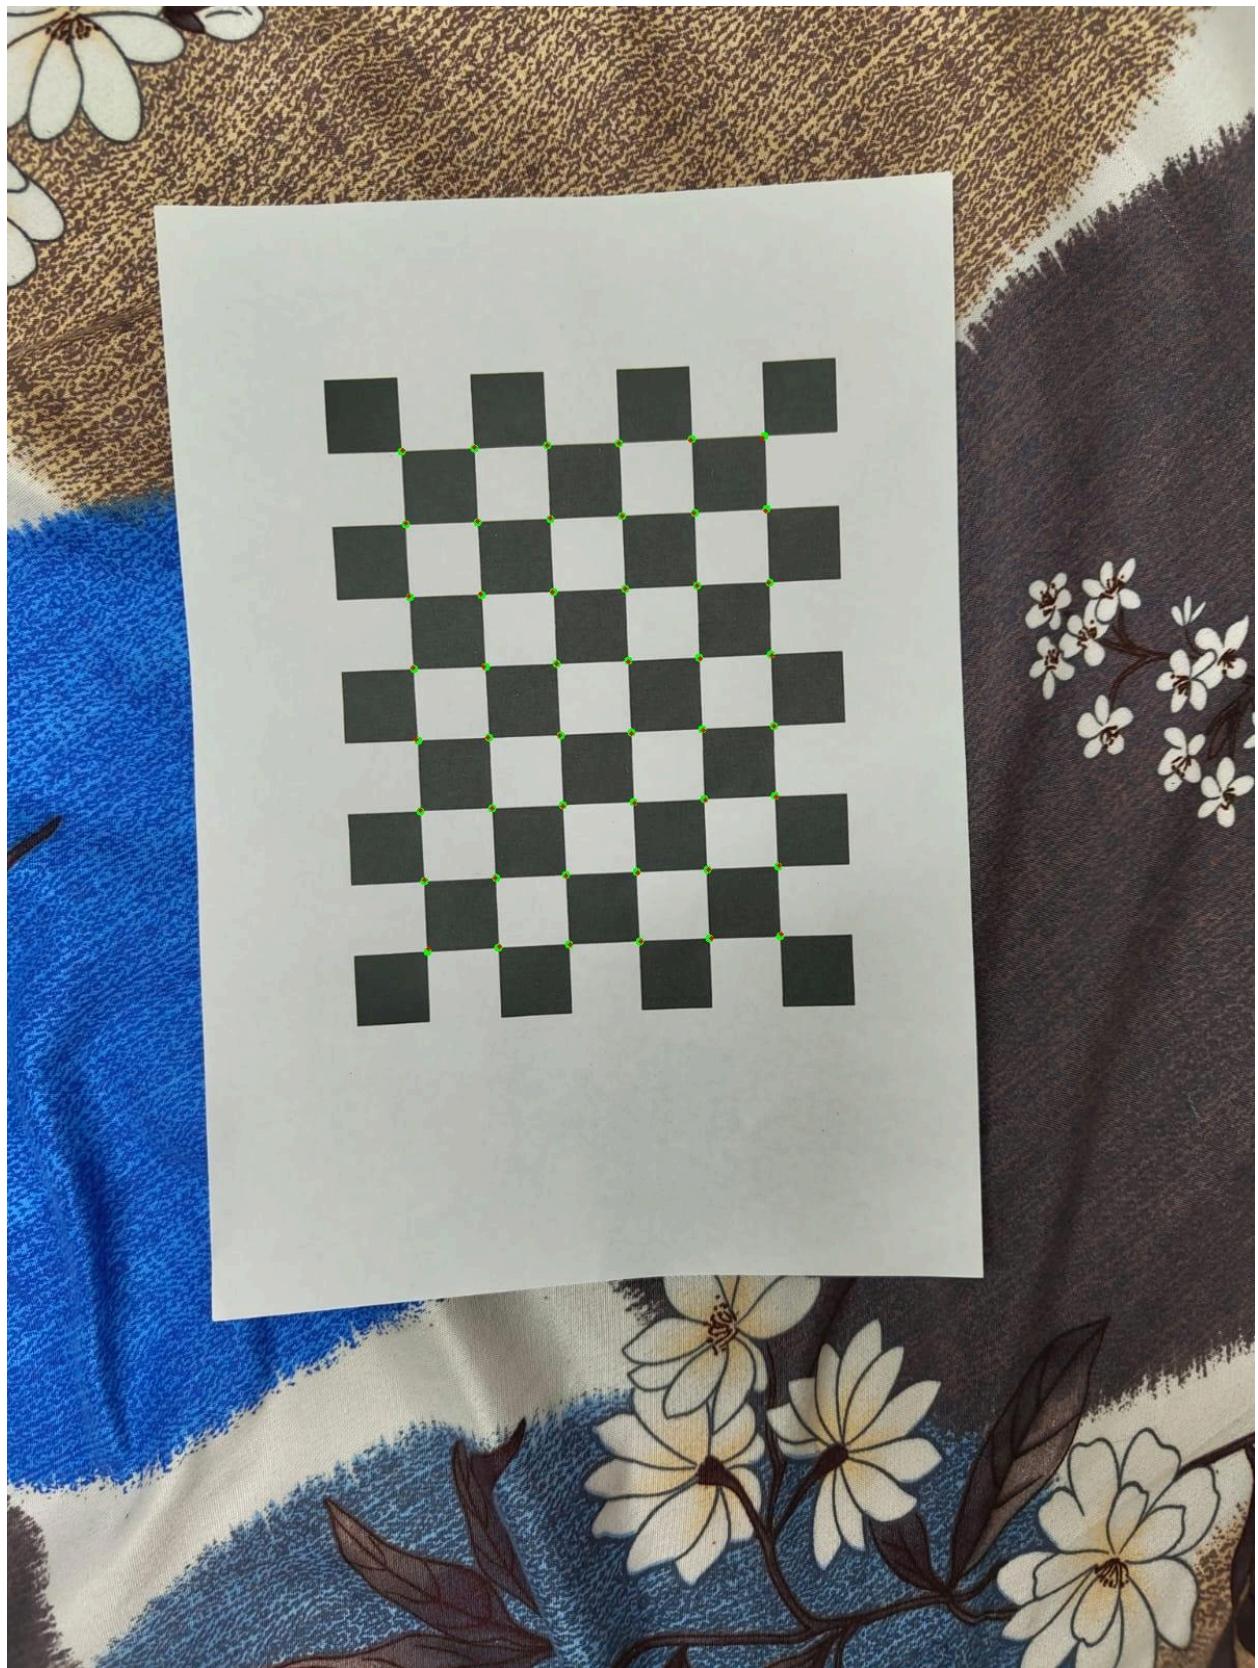


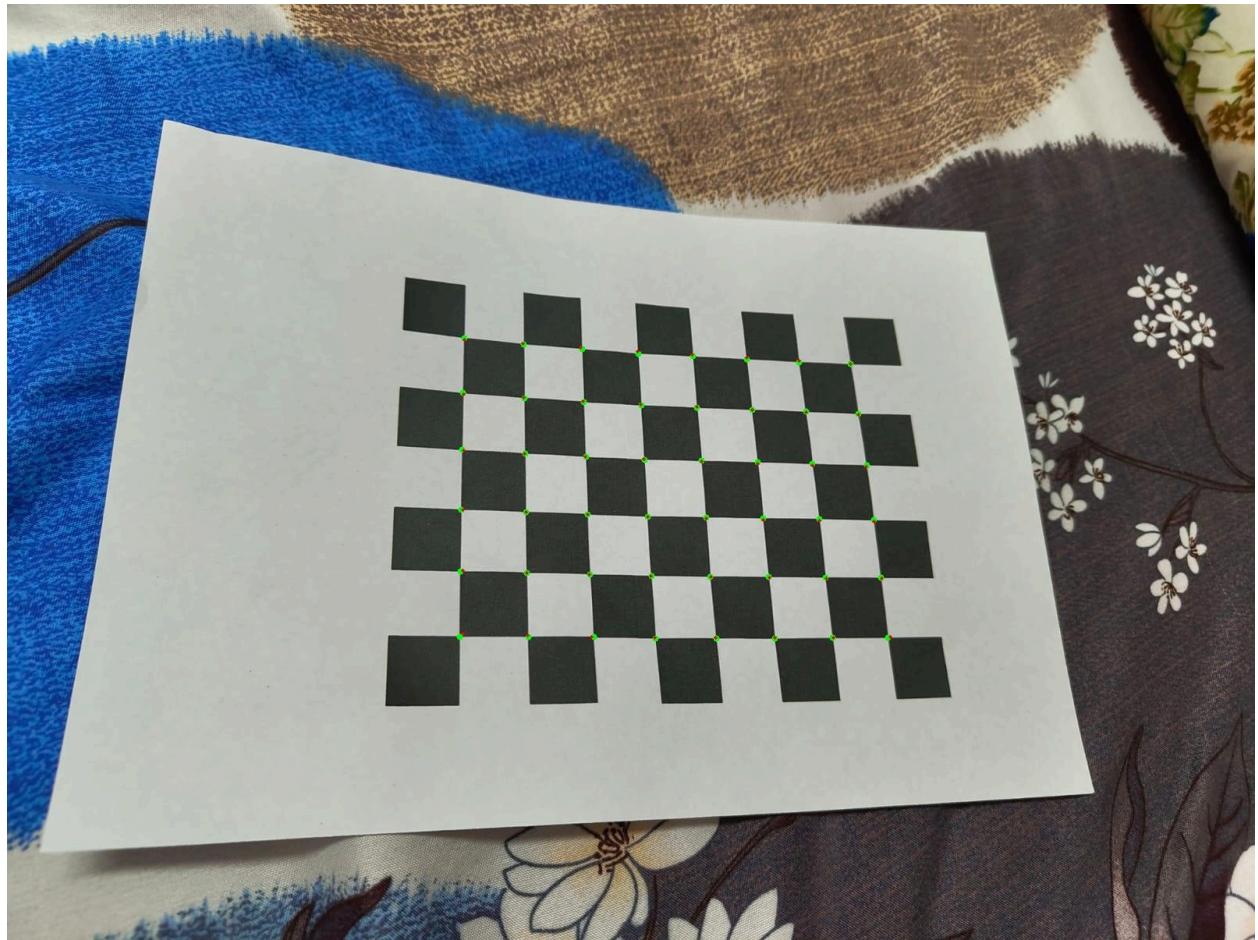


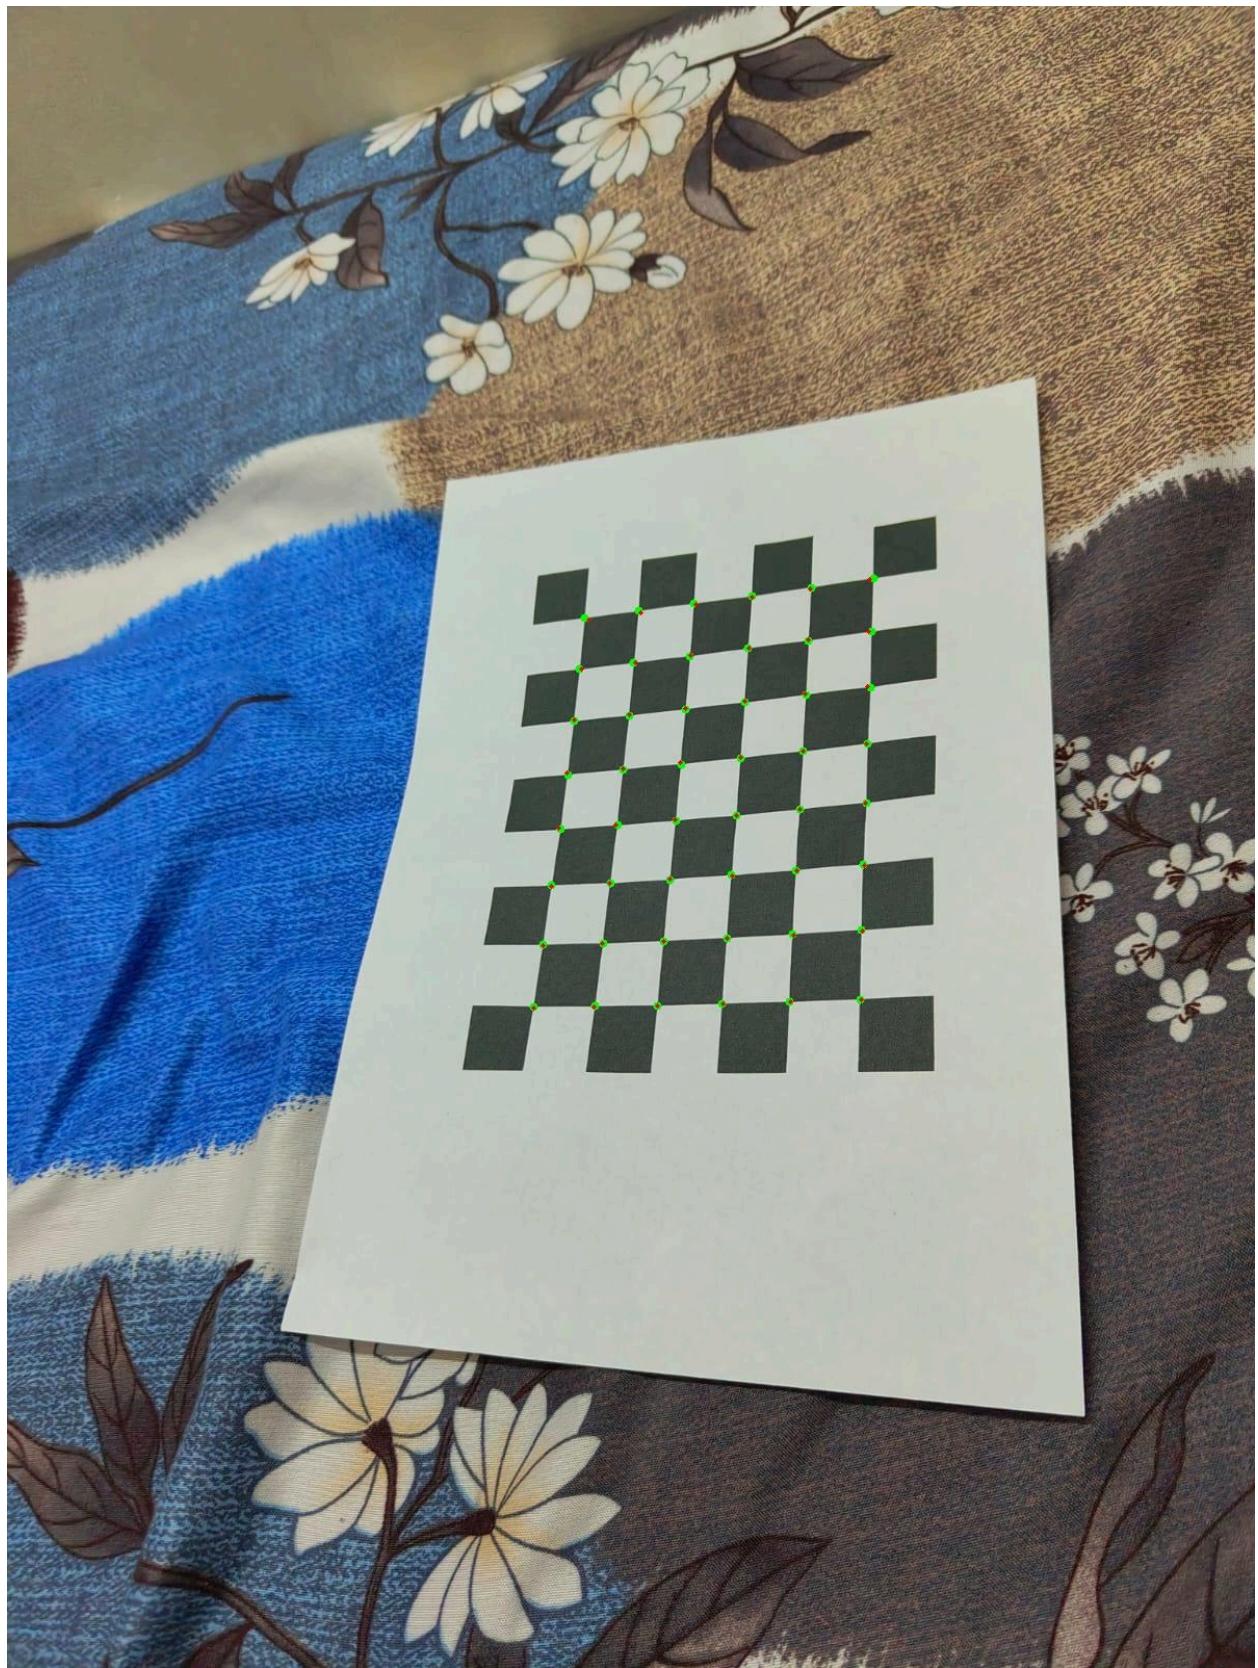


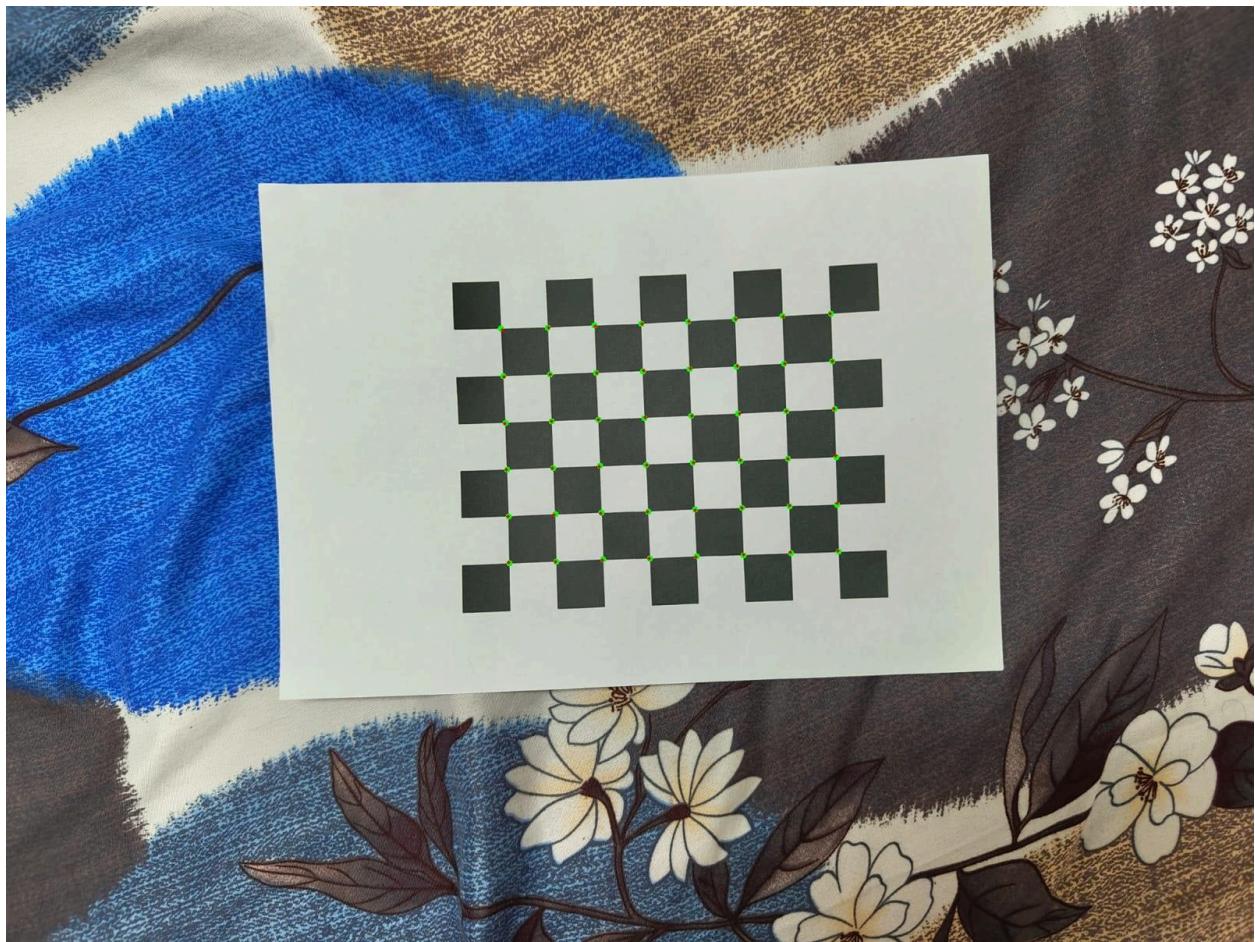




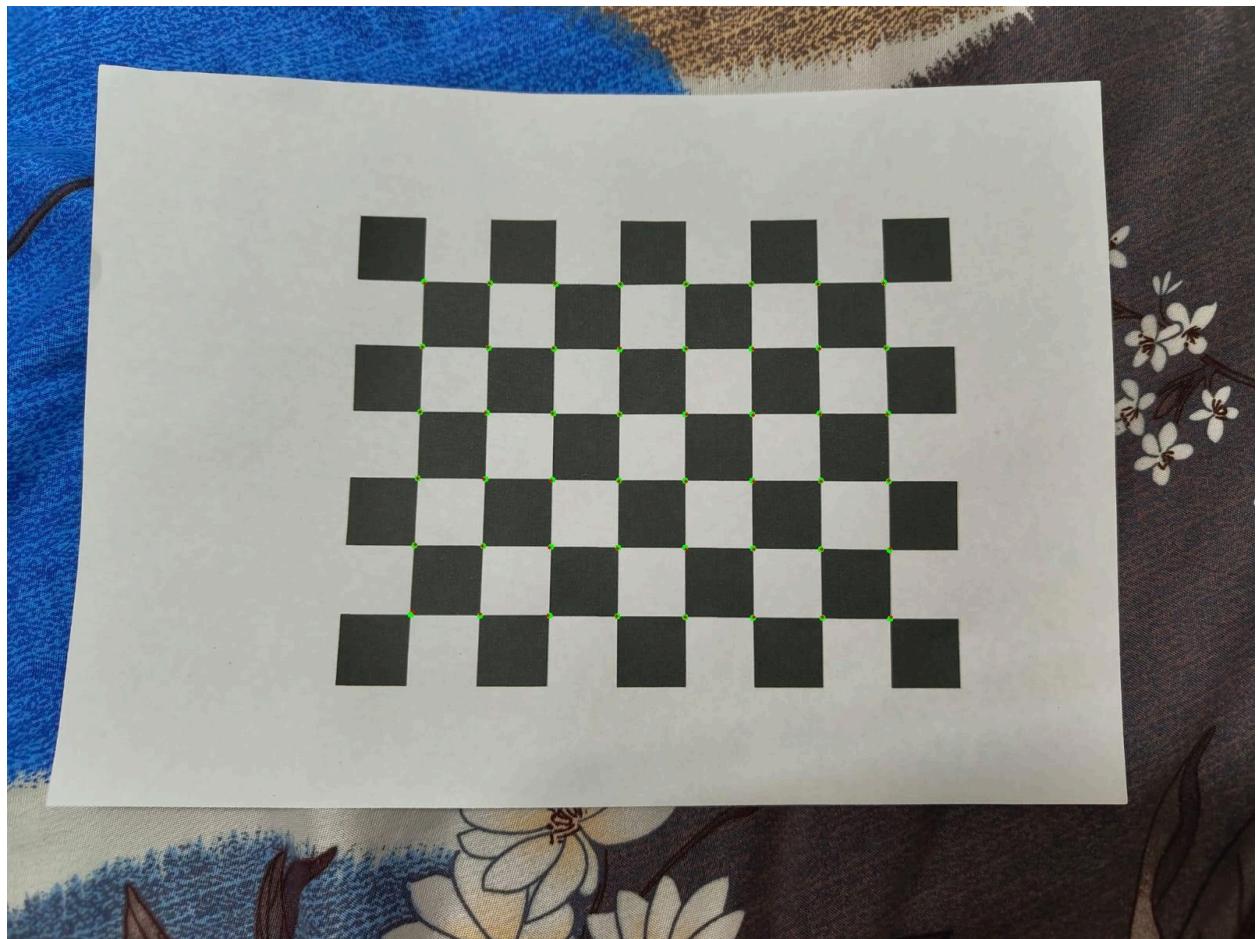




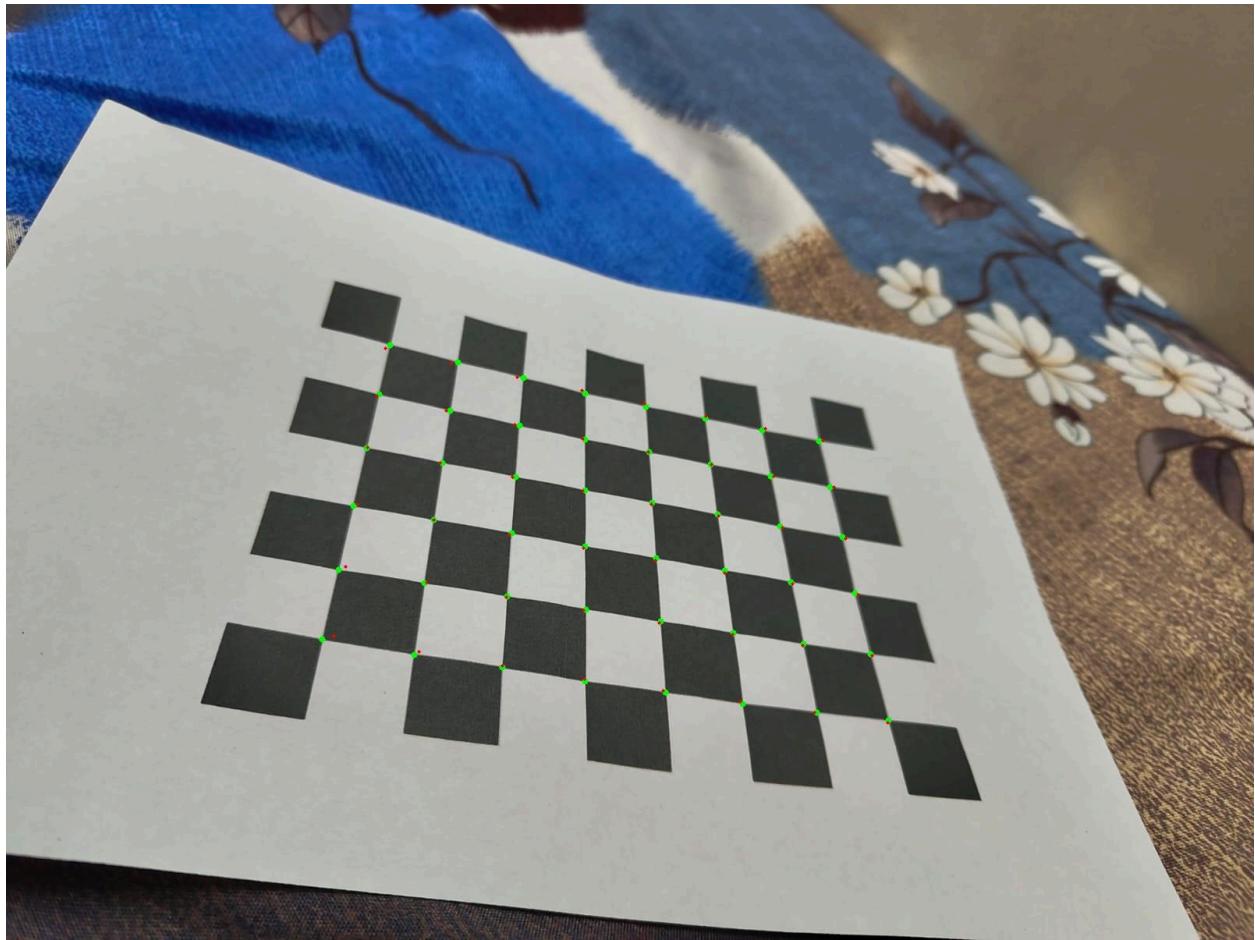


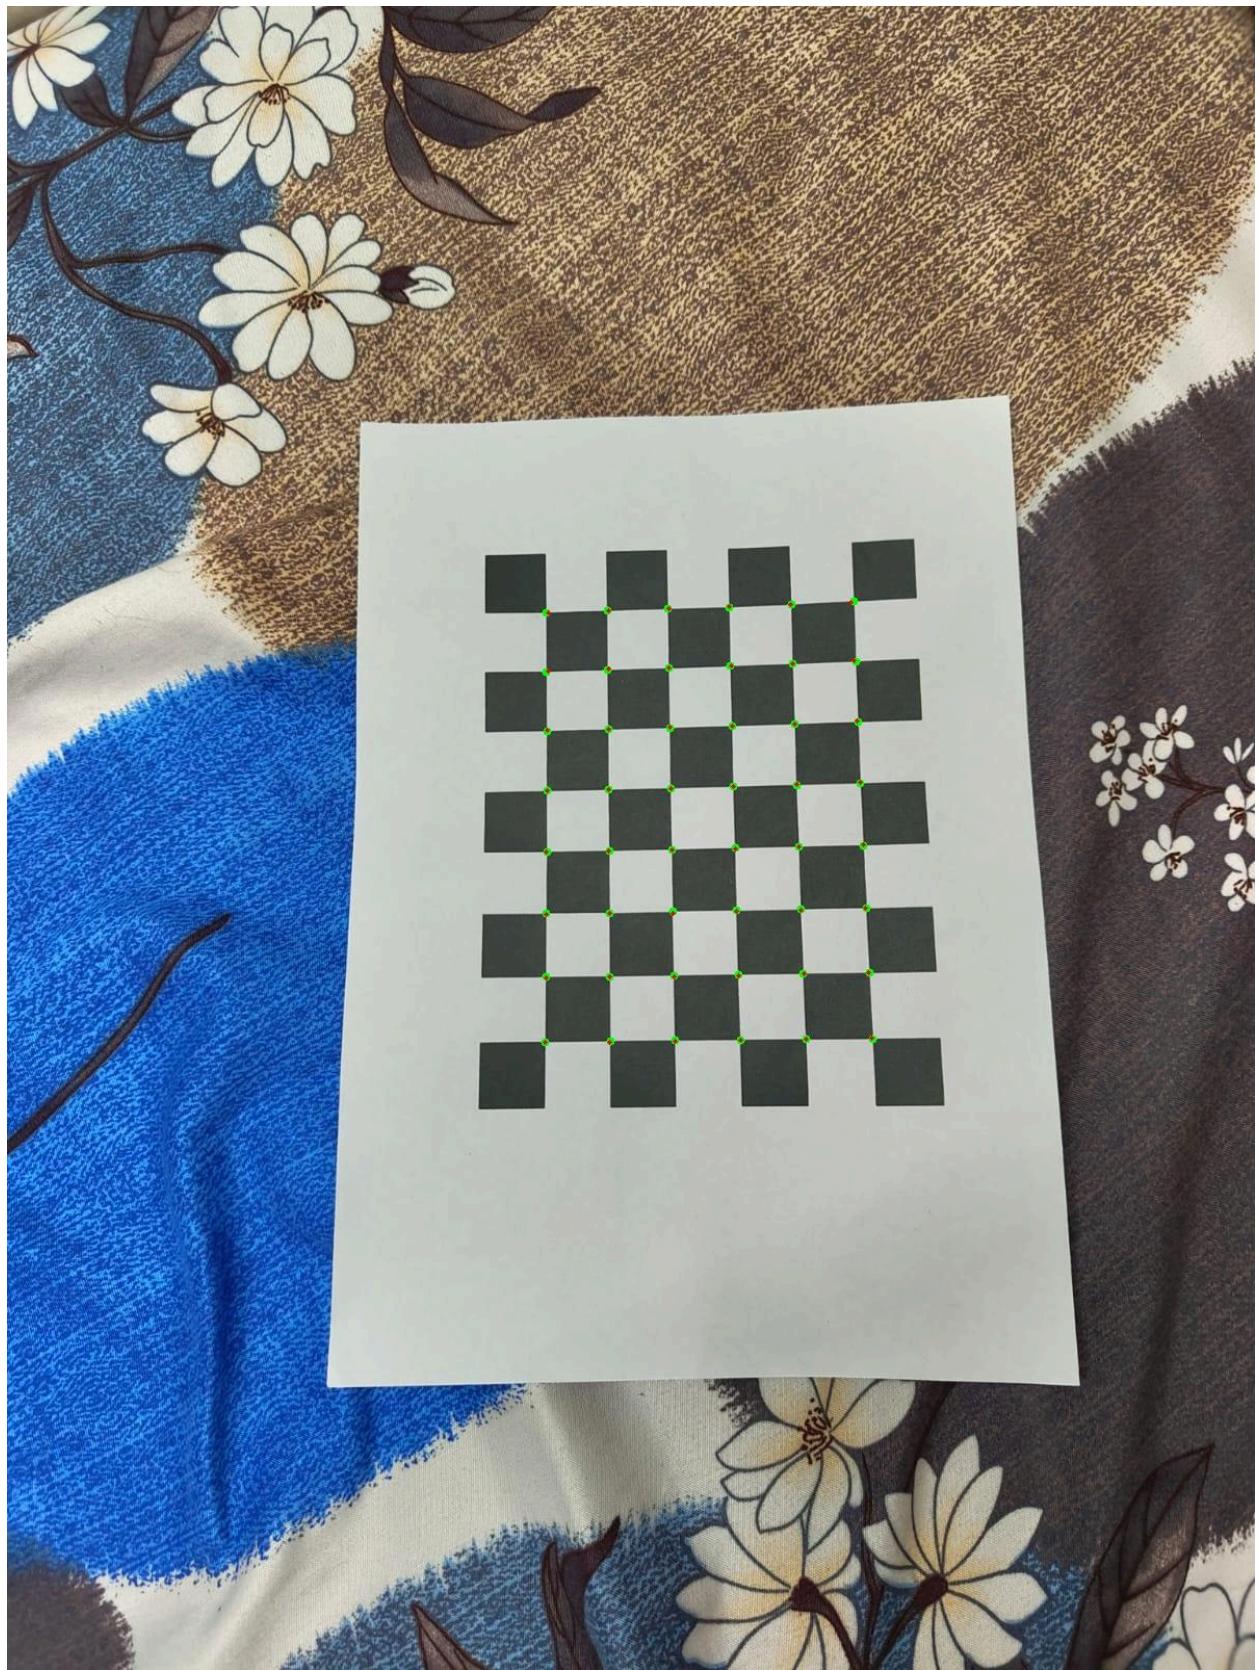


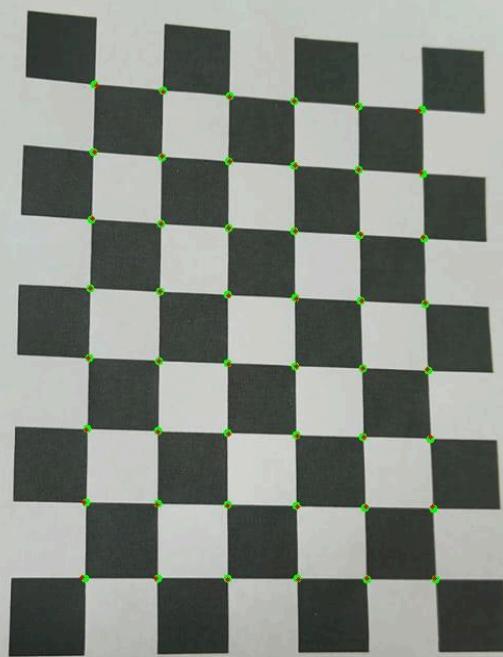






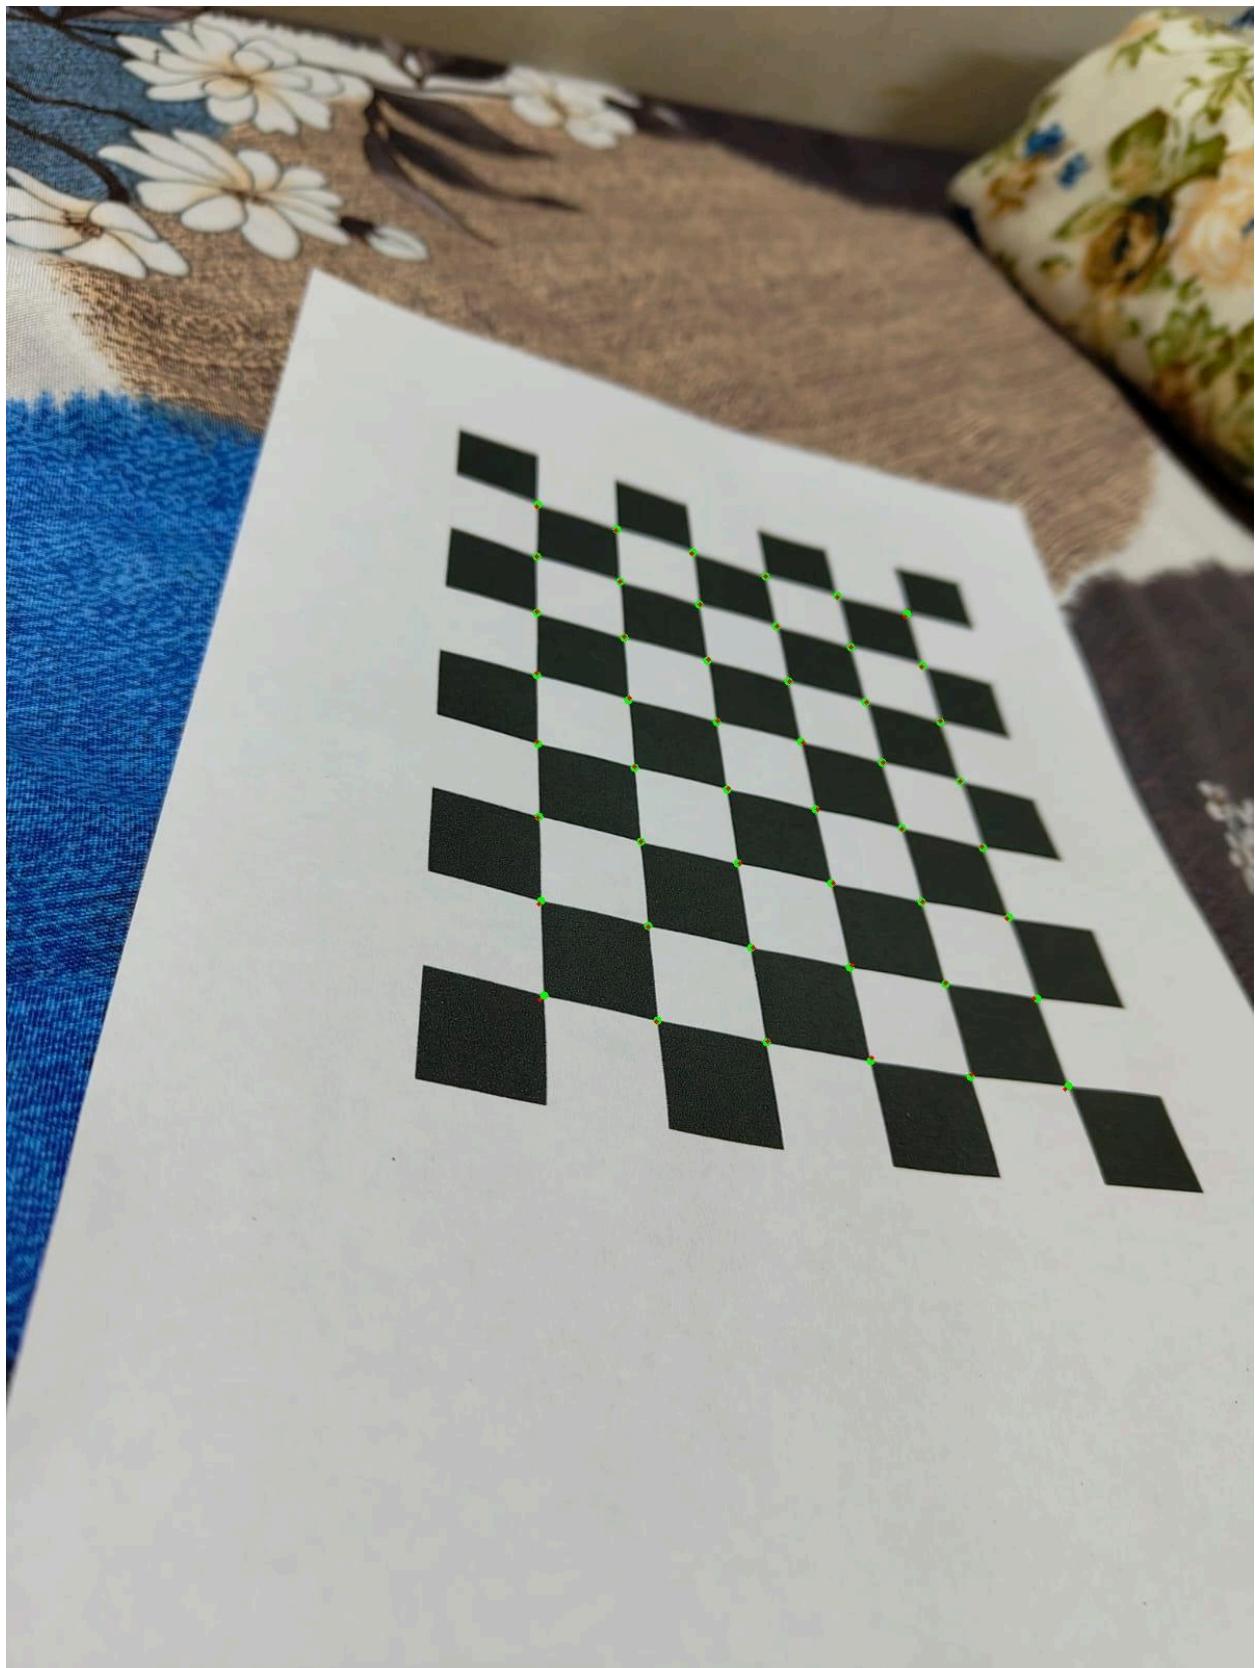


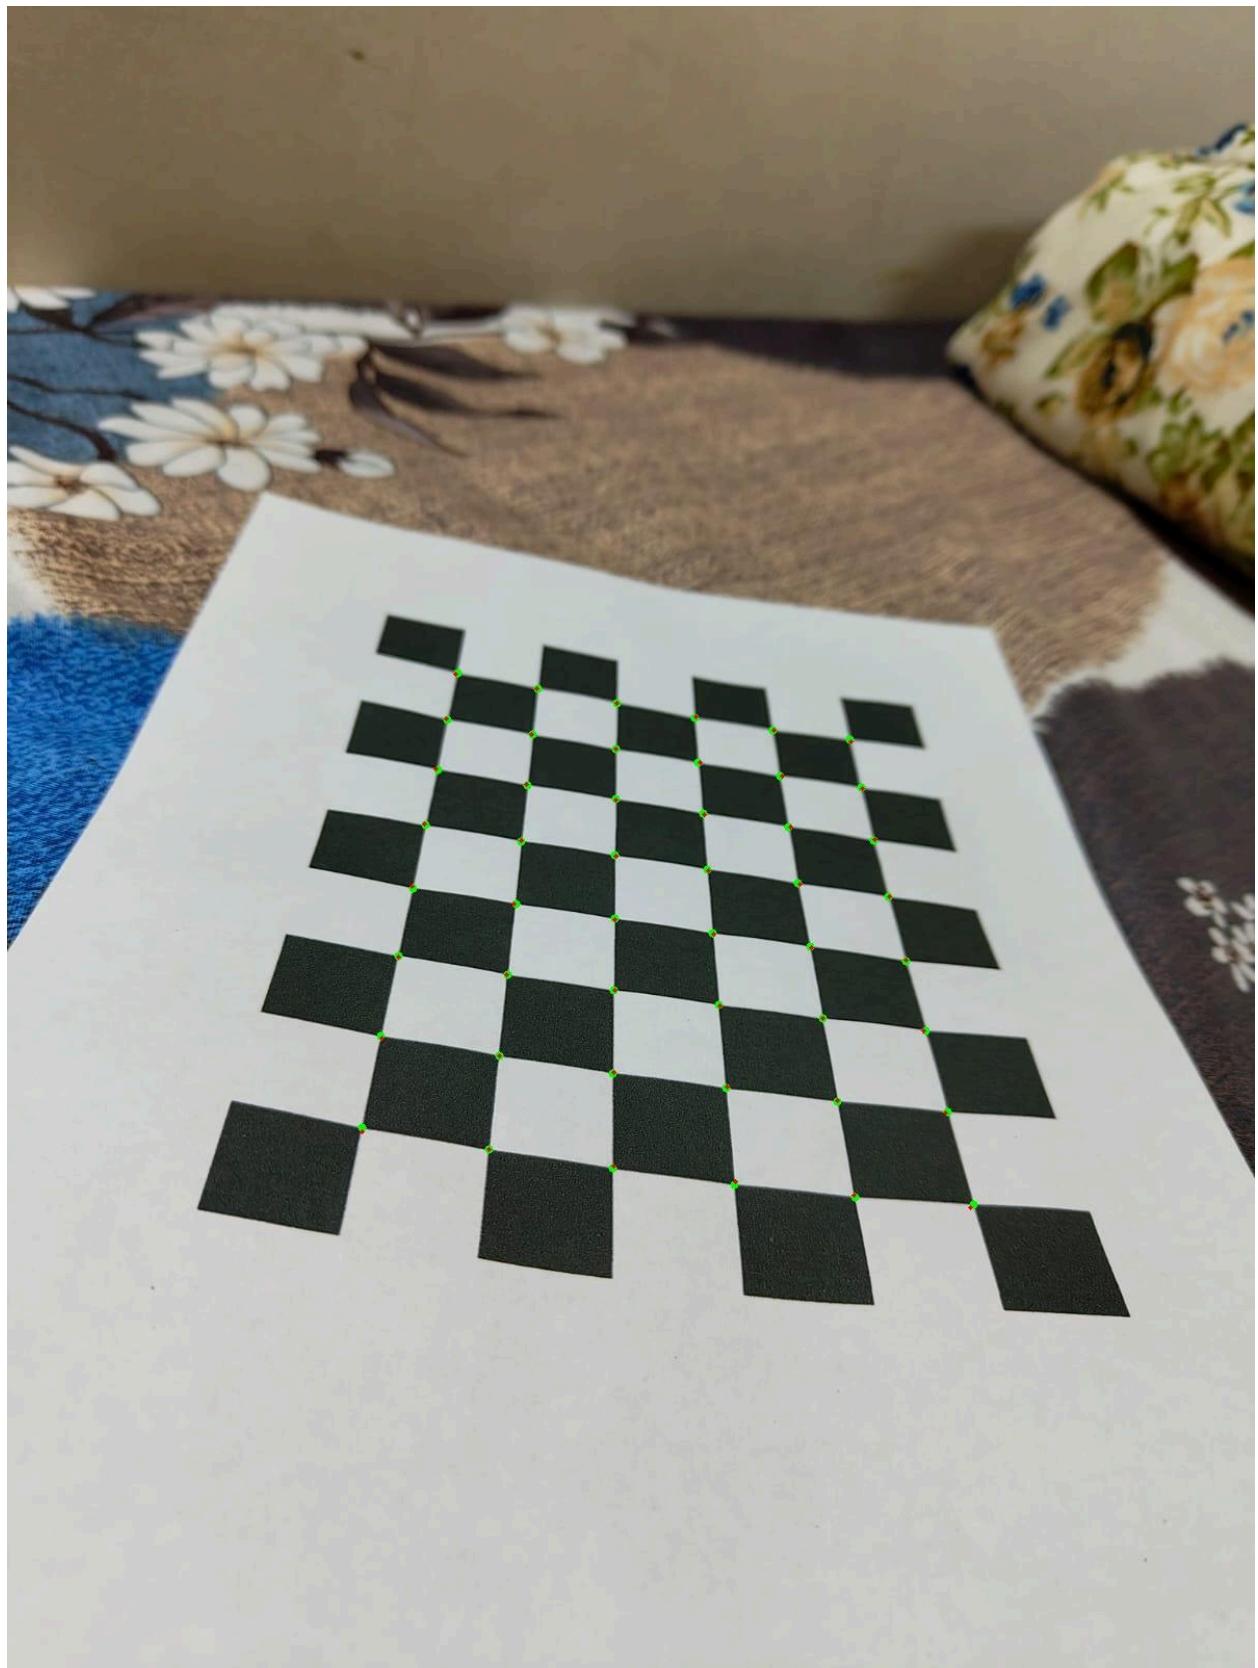












For each calibration image, the **3D world coordinates** of the checkerboard corners are known (in real-world scale).

These 3D points are **projected back onto the 2D image plane** using the estimated **intrinsic** (K) and **extrinsic** (R, t) parameters.

The **Euclidean distance** (L2 norm) is calculated between the projected 2D points and the **actual detected corner points** in the image.

The average of these distances across all corners in an image gives the **per-image reprojection error**.

This value quantifies the calibration accuracy — the lower the error, the better the model fits the observed data.

6. Checkerboard Plane Normals in Camera Frame

- The plane normal $n_{Ci} \in \mathbb{R}^3$ where n_{Ci} is the third column of the **rotation matrix** $R_i R_i^T$ obtained during camera calibration.
- This vector represents the **z-axis of the checkerboard plane** in the **camera's coordinate frame** and indicates its **orientation** relative to the camera.

Image 01 - Plane Normal (Camera Frame): [0.04079217 0.53648296 0.84292469]

Image 02 - Plane Normal (Camera Frame): [-0.06087978 0.59165706 0.80388779]

Image 03 - Plane Normal (Camera Frame): [0.37785269 0.13268537 0.91630887]

Image 04 - Plane Normal (Camera Frame): [-0.43568572 0.62526014 0.64747797]

Image 05 - Plane Normal (Camera Frame): [-0.24641055 0.37073319 0.89545449]

Image 06 - Plane Normal (Camera Frame): [0.38033675 0.64320056 0.66455774]

Image 07 - Plane Normal (Camera Frame): [-0.05397362 0.20087913 0.97812802]

Image 08 - Plane Normal (Camera Frame): [0.27481746 0.71103838 0.64722468]

Image 09 - Plane Normal (Camera Frame): [-0.09064862 0.72586098 0.68184211]

Image 10 - Plane Normal (Camera Frame): [0.14050772 0.32569077 0.9349776]

Image 11 - Plane Normal (Camera Frame): [0.11688001 0.56613103 0.81598696]

Image 12 - Plane Normal (Camera Frame): [-0.0470609 0.03257829 0.99836062]

Image 13 - Plane Normal (Camera Frame): [-0.14823305 0.41211664 0.89899212]

Image 14 - Plane Normal (Camera Frame): [0.36275149 0.56828491 0.73855508]

Image 15 - Plane Normal (Camera Frame): [0.10157754 0.07916255 0.99167298]

Image 16 - Plane Normal (Camera Frame): [0.23857529 0.59602273 0.76670642]

Image 17 - Plane Normal (Camera Frame): [0.06265055 0.18668666 0.98041981]

Image 18 - Plane Normal (Camera Frame): [-0.32193488 0.51615939 0.79368597]

Image 19 - Plane Normal (Camera Frame): [-0.23444451 0.68422356 0.69056057]

Image 20 - Plane Normal (Camera Frame): [0.07009761 0.34120311 0.93737226]

Image 21 - Plane Normal (Camera Frame): [-0.21349236 0.34599446 0.91362402]

Image 22 - Plane Normal (Camera Frame): [0.16121956 0.05441965 0.98541705]

Image 23 - Plane Normal (Camera Frame): [0.15584567 0.8196933 0.55119418]

Image 24 - Plane Normal (Camera Frame): [-0.38253519 0.76517249 0.51785895]

Image 25 - Plane Normal (Camera Frame): [-0.21382321 0.87811181 0.42801785]

4. 1.



The keypoints are spread across **foreground and background**, including:

- Trees
- Water ripples
- Railings
- Buildings
- Sky texture (clouds)

The keypoints are displayed with **varying circle sizes** (scale)

Their **directionality** is captured (shown as orientation vectors)

This confirms that SIFT has correctly encoded **scale- and rotation-invariant** descriptors

2.



After extracting SIFT keypoints and descriptors from the first two images, we matched these features using two different approaches:

1. BruteForce Matcher

- Compares **each descriptor** in Image 1 with **all descriptors** in Image 2
- Retains matches based on **Lowe's Ratio Test** (0.75 threshold)

2. FLANN-Based Matcher

- Uses **approximate nearest neighbor search** for faster matching
- A hierarchical KD-Tree structure was used with:
 - `trees = 5`
 - `checks = 50`
- Lowe's Ratio Test was applied with a stricter threshold of **0.7**

Left: BruteForce Matcher — **667 matches**

Right: FLANN-Based Matcher — **608 matches**

Both methods result in **dense and geometrically consistent matches** across the overlapping region.

Matched keypoints align well along edges, water ripples, and buildings — confirming **correctness and reliability** of detected features.

FLANN is slightly faster and still produces **high-quality matches**, making it ideal for large-scale datasets or multi-image stitching.

3. To geometrically align the two images (`image1` and `image2`), we compute the **homography matrix** — a 3×3 transformation matrix that maps points from one image plane to another.

Methodology:

- Used **BruteForce-matched keypoints** (filtered with Lowe's ratio test) for establishing correspondences.
- **RANSAC (Random Sample Consensus)** was applied via `cv2.findHomography()` to robustly estimate the transformation:
 - RANSAC mitigates the impact of outliers
 - A reprojection threshold of `5.0` pixels was used

Homography matrix saved to `homography_matrix.csv`

```
[[ 8.99023648e-01 -1.31282351e-01 -1.64442160e+02]
```

```
[ 1.43181983e-01  9.90713800e-01 -6.94375875e+01]
```

```
[-1.27006065e-07  3.60666010e-07  1.00000000e+00]]
```

4.



Before stitching the images together, I first warp them into a **common perspective space** using the estimated **homography matrix**. This transformation adjusts the viewpoint of the images, allowing us to align and overlap them correctly for panorama creation.

- We used SIFT features and **BruteForce matching with Lowe's ratio test** to identify good feature correspondences.
- Using these matches, a **homography matrix** was computed via **RANSAC**.
- Both images were then warped into a **shared coordinate frame**:
 - **image1** was warped using the homography matrix
 - **image2** was placed in the same canvas without transformation
- A translation matrix was added to avoid negative pixel coordinates in the combined space.

Left: Warped **Image 1** using homography

Right: Unwarped **Image 2**, but shifted to align with **Image 1**

The overlapping content (bridge, tree, and water) appears **aligned**, verifying the success of the homography.

5.



Homography Matrix from earlier steps was reused to align the images.

[image1](#) was warped using this transformation, and [image2](#) was placed directly on the shared canvas.

To ensure a **clean and polished output**, we removed empty (black) borders using **binary thresholding + contour detection** for cropping.

Left: Direct stitching without cropping. Notice the visible black borders due to warping into a larger canvas.

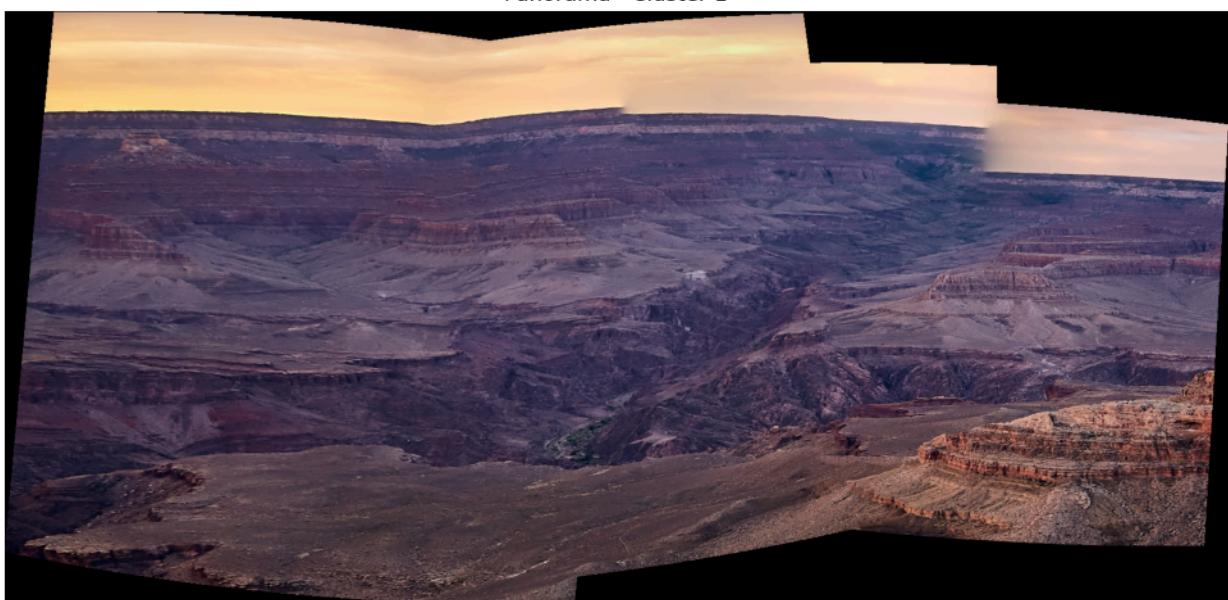
Right: Final panorama after cropping out unnecessary black regions. The final result looks visually cohesive and clean.

6.

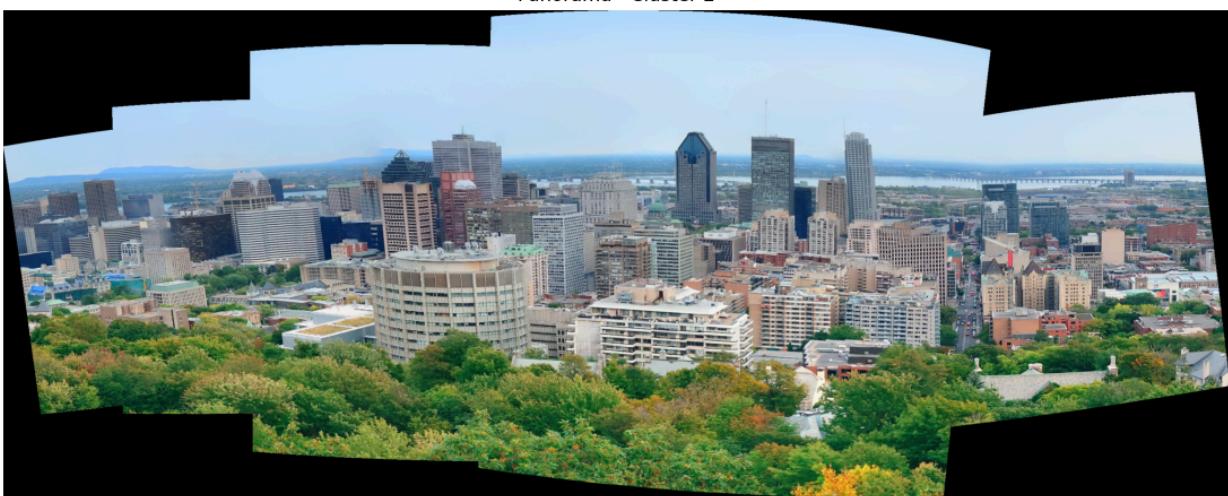
Panorama - Cluster 0



Panorama - Cluster 1



Panorama - Cluster 2



5. 1. I implemented the **Iterative Closest Point (ICP)** algorithm using the `open3d` library to align two consecutive point clouds, `pointcloud_0000.pcd` and `pointcloud_0001.pcd`, recorded using a 3D LiDAR mounted on a TurtleBot.

- We selected **Point-to-Point ICP** as our registration strategy.
- The point clouds were **voxel downsampled** for faster processing.
- An **initial transformation matrix** was randomly generated using a valid orthonormal rotation matrix and small translation offsets.
- The ICP algorithm was then executed to refine this transformation and align the source point cloud with the target.

Results

Metric	Initial Transformation	After ICP
Fitness	0.0311	0.0360
Inlier RMSE	0.0690	0.0651

- The fitness score slightly improved after ICP, indicating a better overlap between the two point clouds.
- The RMSE reduced, reflecting improved point correspondence accuracy.

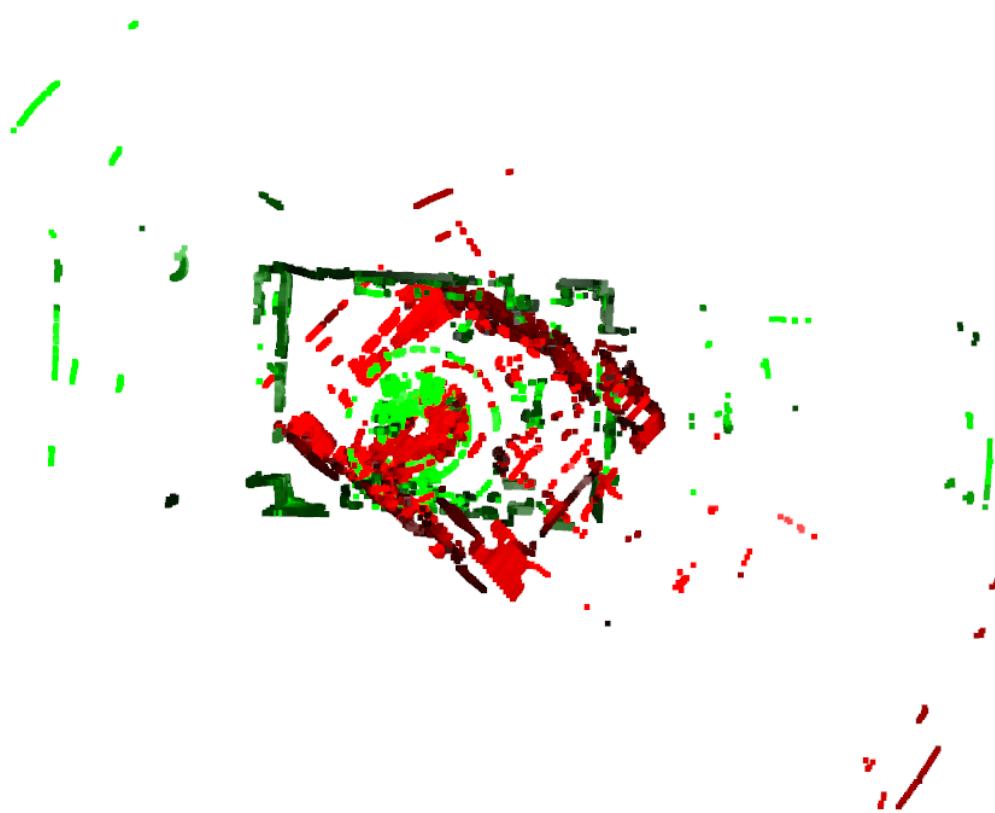
Estimated Transformation Matrix (After ICP):

```
[ 0.66587624  0.60855351  0.43159176  0.17543967]
```

```
[-0.28463572  0.74195671 -0.6070278 -0.00261867]
```

```
[ 0.68963131 -0.28135896 -0.66726741 -0.00356774]
```

```
[ 0.      0.      0.      1.    ]]
```



This figure shows the source point cloud (in **red**) and the target point cloud (in **green**) after registration.

2. To improve registration accuracy and robustness, I conducted a series of experiments by varying:

- The **initial transformation matrix** (identity, random orthonormal, and RANSAC-based),
- The **ICP method** (point-to-point vs point-to-plane),
- The **maximum correspondence threshold**.

Results

Init Guess	Method	Threshold	Fitness	Inlier RMSE	$\ T_{est} - T_{init}\ $
Identity	Point-to-Point	0.100	0.9937	0.0206	0.0044
Random	Point-to-Point	0.100	0.0474	0.0653	0.1564
Random	Point-to-Plane	0.100	0.0462	0.0651	0.2719
Identity	Point-to-Plane	0.050	0.9674	0.0179	0.0005
RANSAC	Point-to-Point	0.100	0.9937	0.0206	0.0118

Best configuration: *Identity matrix + Point-to-Plane ICP + 0.05 threshold*

This configuration yielded the **lowest inlier RMSE (0.0179)** and **minimum transformation correction ($\|T_{est} - T_{init}\| = 0.0005$)**, indicating a highly accurate alignment with minimal deviation from the initial guess.

Best Estimated Transformation Matrix

The learned transformation matrix from the best configuration (Identity Init + Point-to-Plane ICP, threshold = 0.05) is:

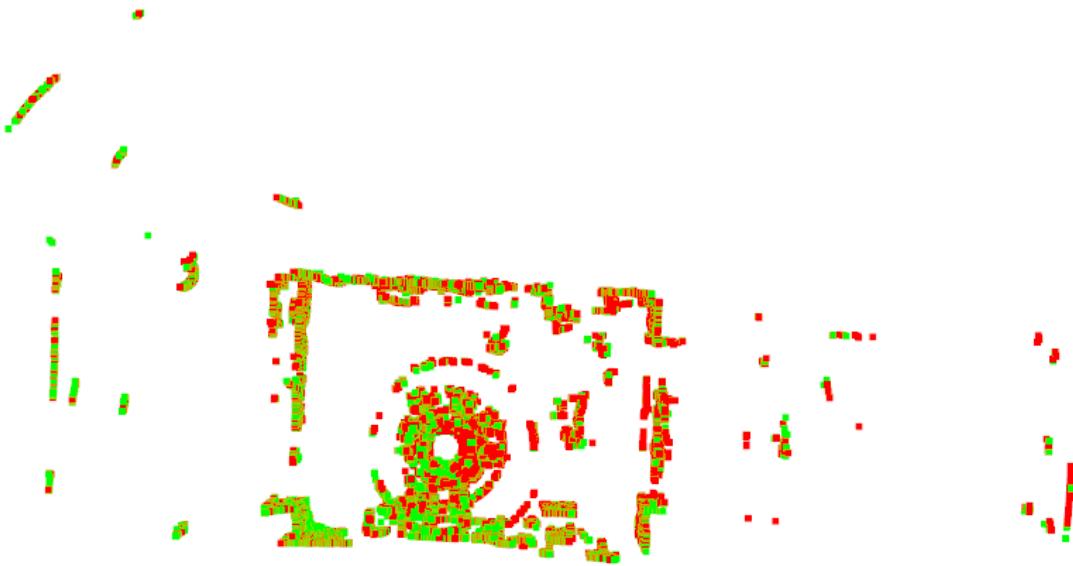
```
[ 9.9999995e-01 -5.12637423e-05 -8.79082824e-05 -4.56397355e-04]
[ 5.12554576e-05  9.9999994e-01 -9.42423543e-05  2.32611106e-05]
[ 8.79131131e-05  9.42378481e-05  9.9999992e-01 -6.31491481e-05]
[ 0.00000000e+00  0.00000000e+00  0.00000000e+00  1.00000000e+00]]
```

3. Based on the hyperparameter tuning in Part 2, the best result was achieved using the following configuration:

- **Initialization:** Identity matrix
- **ICP Type:** Point-to-Plane
- **Correspondence Threshold:** 0.05
- **Fitness:** 0.9674
- **Inlier RMSE:** 0.0179
- **$\|T_{\text{est}} - T_{\text{init}}\|$:** 0.0005

Using the above configuration, we applied the estimated transformation matrix to the original source point cloud (`pointcloud_000.pcd`). The matrix used is:

```
[ 9.9999995e-01 -5.12637423e-05 -8.79082824e-05 -4.56397355e-04]  
[ 5.12554576e-05  9.9999994e-01 -9.42423543e-05  2.32611106e-05]  
[ 8.79131131e-05  9.42378481e-05  9.9999992e-01 -6.31491481e-05]  
[ 0.00000000e+00  0.00000000e+00  0.00000000e+00  1.00000000e+00]]
```



This figure shows the transformed source point cloud (**red**) overlaid on the target point cloud (**green**).

- The near-perfect alignment seen in the visualization validates the numerical results.
- The **high fitness score (0.967)** indicates that a large proportion of the source points found close correspondences in the target cloud.
- The **low RMSE (0.0179)** implies minimal alignment error between corresponding points.
- The success of the identity initialization suggests that the point clouds were already closely aligned, and the point-to-plane ICP only needed to make fine adjustments.

4. To evaluate the global performance of point cloud registration, I applied **point-to-point ICP** iteratively across a sequence of **100 point cloud frames** (`pointcloud_0000.pcd` to `pointcloud_0396.pcd`).

- **Pairwise ICP** was applied between every two consecutive frames (4-frame skips).
- The **estimated transformation** from each ICP step was accumulated to compute a **global transformation** for each frame.
- Each point cloud was transformed using its global pose and merged to create a **single registered map**.

- The **translation component** of each global transformation matrix was recorded to estimate the **3D trajectory** of the TurtleBot.

Registered pointcloud_0000.pcd -> pointcloud_0004.pcd | Fitness: 0.9937

Registered pointcloud_0004.pcd -> pointcloud_0008.pcd | Fitness: 0.9915

Registered pointcloud_0008.pcd -> pointcloud_0012.pcd | Fitness: 0.9912

Registered pointcloud_0012.pcd -> pointcloud_0016.pcd | Fitness: 0.9909

Registered pointcloud_0016.pcd -> pointcloud_0020.pcd | Fitness: 0.9917

Registered pointcloud_0020.pcd -> pointcloud_0024.pcd | Fitness: 0.9887

Registered pointcloud_0024.pcd -> pointcloud_0028.pcd | Fitness: 0.9881

Registered pointcloud_0028.pcd -> pointcloud_0032.pcd | Fitness: 0.9918

Registered pointcloud_0032.pcd -> pointcloud_0036.pcd | Fitness: 0.9844

Registered pointcloud_0036.pcd -> pointcloud_0040.pcd | Fitness: 0.9600

Registered pointcloud_0040.pcd -> pointcloud_0044.pcd | Fitness: 0.9683

Registered pointcloud_0044.pcd -> pointcloud_0048.pcd | Fitness: 0.9671

Registered pointcloud_0048.pcd -> pointcloud_0052.pcd | Fitness: 0.9781

Registered pointcloud_0052.pcd -> pointcloud_0056.pcd | Fitness: 0.9785

Registered pointcloud_0056.pcd -> pointcloud_0060.pcd | Fitness: 0.9677

Registered pointcloud_0060.pcd -> pointcloud_0064.pcd | Fitness: 0.9749

Registered pointcloud_0064.pcd -> pointcloud_0068.pcd | Fitness: 0.9692

Registered pointcloud_0068.pcd -> pointcloud_0072.pcd | Fitness: 0.9831

Registered pointcloud_0072.pcd -> pointcloud_0076.pcd | Fitness: 0.9689

Registered pointcloud_0076.pcd -> pointcloud_0080.pcd | Fitness: 0.9807

Registered pointcloud_0080.pcd -> pointcloud_0084.pcd | Fitness: 0.9754

Registered pointcloud_0084.pcd -> pointcloud_0088.pcd | Fitness: 0.9842

Registered pointcloud_0088.pcd -> pointcloud_0092.pcd | Fitness: 0.9758

Registered pointcloud_0092.pcd -> pointcloud_0096.pcd | Fitness: 0.9881
Registered pointcloud_0096.pcd -> pointcloud_0100.pcd | Fitness: 0.9781
Registered pointcloud_0100.pcd -> pointcloud_0104.pcd | Fitness: 0.9863
Registered pointcloud_0104.pcd -> pointcloud_0108.pcd | Fitness: 0.9804
Registered pointcloud_0108.pcd -> pointcloud_0112.pcd | Fitness: 0.9875
Registered pointcloud_0112.pcd -> pointcloud_0116.pcd | Fitness: 0.9696
Registered pointcloud_0116.pcd -> pointcloud_0120.pcd | Fitness: 0.9849
Registered pointcloud_0120.pcd -> pointcloud_0124.pcd | Fitness: 0.9773
Registered pointcloud_0124.pcd -> pointcloud_0128.pcd | Fitness: 0.9878
Registered pointcloud_0128.pcd -> pointcloud_0132.pcd | Fitness: 0.9782
Registered pointcloud_0132.pcd -> pointcloud_0136.pcd | Fitness: 0.9874
Registered pointcloud_0136.pcd -> pointcloud_0140.pcd | Fitness: 0.9803
Registered pointcloud_0140.pcd -> pointcloud_0144.pcd | Fitness: 0.9873
Registered pointcloud_0144.pcd -> pointcloud_0148.pcd | Fitness: 0.9834
Registered pointcloud_0148.pcd -> pointcloud_0152.pcd | Fitness: 0.9847
Registered pointcloud_0152.pcd -> pointcloud_0156.pcd | Fitness: 0.9835
Registered pointcloud_0156.pcd -> pointcloud_0160.pcd | Fitness: 0.9888
Registered pointcloud_0160.pcd -> pointcloud_0164.pcd | Fitness: 0.9864
Registered pointcloud_0164.pcd -> pointcloud_0168.pcd | Fitness: 0.9909
Registered pointcloud_0168.pcd -> pointcloud_0172.pcd | Fitness: 0.9870
Registered pointcloud_0172.pcd -> pointcloud_0176.pcd | Fitness: 0.9846
Registered pointcloud_0176.pcd -> pointcloud_0180.pcd | Fitness: 0.9826
Registered pointcloud_0180.pcd -> pointcloud_0184.pcd | Fitness: 0.9885
Registered pointcloud_0184.pcd -> pointcloud_0188.pcd | Fitness: 0.9891

Registered pointcloud_0188.pcd -> pointcloud_0192.pcd | Fitness: 0.9819
Registered pointcloud_0192.pcd -> pointcloud_0196.pcd | Fitness: 0.9887
Registered pointcloud_0196.pcd -> pointcloud_0200.pcd | Fitness: 0.9844
Registered pointcloud_0200.pcd -> pointcloud_0204.pcd | Fitness: 0.9878
Registered pointcloud_0204.pcd -> pointcloud_0208.pcd | Fitness: 0.9838
Registered pointcloud_0208.pcd -> pointcloud_0212.pcd | Fitness: 0.9891
Registered pointcloud_0212.pcd -> pointcloud_0216.pcd | Fitness: 0.9860
Registered pointcloud_0216.pcd -> pointcloud_0220.pcd | Fitness: 0.9864
Registered pointcloud_0220.pcd -> pointcloud_0224.pcd | Fitness: 0.9805
Registered pointcloud_0224.pcd -> pointcloud_0228.pcd | Fitness: 0.9917
Registered pointcloud_0228.pcd -> pointcloud_0232.pcd | Fitness: 0.9853
Registered pointcloud_0232.pcd -> pointcloud_0236.pcd | Fitness: 0.9860
Registered pointcloud_0236.pcd -> pointcloud_0240.pcd | Fitness: 0.9901
Registered pointcloud_0240.pcd -> pointcloud_0244.pcd | Fitness: 0.9886
Registered pointcloud_0244.pcd -> pointcloud_0248.pcd | Fitness: 0.9892
Registered pointcloud_0248.pcd -> pointcloud_0252.pcd | Fitness: 0.9898
Registered pointcloud_0252.pcd -> pointcloud_0256.pcd | Fitness: 0.9887
Registered pointcloud_0256.pcd -> pointcloud_0260.pcd | Fitness: 0.9888
Registered pointcloud_0260.pcd -> pointcloud_0264.pcd | Fitness: 0.9879
Registered pointcloud_0264.pcd -> pointcloud_0268.pcd | Fitness: 0.9842
Registered pointcloud_0268.pcd -> pointcloud_0272.pcd | Fitness: 0.9869
Registered pointcloud_0272.pcd -> pointcloud_0276.pcd | Fitness: 0.9900
Registered pointcloud_0276.pcd -> pointcloud_0280.pcd | Fitness: 0.9886
Registered pointcloud_0280.pcd -> pointcloud_0284.pcd | Fitness: 0.9866

Registered pointcloud_0284.pcd -> pointcloud_0288.pcd | Fitness: 0.9881
Registered pointcloud_0288.pcd -> pointcloud_0292.pcd | Fitness: 0.9906
Registered pointcloud_0292.pcd -> pointcloud_0296.pcd | Fitness: 0.9888
Registered pointcloud_0296.pcd -> pointcloud_0300.pcd | Fitness: 0.9880
Registered pointcloud_0300.pcd -> pointcloud_0304.pcd | Fitness: 0.9897
Registered pointcloud_0304.pcd -> pointcloud_0308.pcd | Fitness: 0.9894
Registered pointcloud_0308.pcd -> pointcloud_0312.pcd | Fitness: 0.9889
Registered pointcloud_0312.pcd -> pointcloud_0316.pcd | Fitness: 0.9881
Registered pointcloud_0316.pcd -> pointcloud_0320.pcd | Fitness: 0.9864
Registered pointcloud_0320.pcd -> pointcloud_0324.pcd | Fitness: 0.9925
Registered pointcloud_0324.pcd -> pointcloud_0328.pcd | Fitness: 0.9911
Registered pointcloud_0328.pcd -> pointcloud_0332.pcd | Fitness: 0.9853
Registered pointcloud_0332.pcd -> pointcloud_0336.pcd | Fitness: 0.9861
Registered pointcloud_0336.pcd -> pointcloud_0340.pcd | Fitness: 0.9896
Registered pointcloud_0340.pcd -> pointcloud_0344.pcd | Fitness: 0.9883
Registered pointcloud_0344.pcd -> pointcloud_0348.pcd | Fitness: 0.9846
Registered pointcloud_0348.pcd -> pointcloud_0352.pcd | Fitness: 0.9849
Registered pointcloud_0352.pcd -> pointcloud_0356.pcd | Fitness: 0.9896
Registered pointcloud_0356.pcd -> pointcloud_0360.pcd | Fitness: 0.9900
Registered pointcloud_0360.pcd -> pointcloud_0364.pcd | Fitness: 0.9872
Registered pointcloud_0364.pcd -> pointcloud_0368.pcd | Fitness: 0.9854
Registered pointcloud_0368.pcd -> pointcloud_0372.pcd | Fitness: 0.9884
Registered pointcloud_0372.pcd -> pointcloud_0376.pcd | Fitness: 0.9901
Registered pointcloud_0376.pcd -> pointcloud_0380.pcd | Fitness: 0.9832

Registered pointcloud_0380.pcd -> pointcloud_0384.pcd | Fitness: 0.9874

Registered pointcloud_0384.pcd -> pointcloud_0388.pcd | Fitness: 0.9847

Registered pointcloud_0388.pcd -> pointcloud_0392.pcd | Fitness: 0.9891

Registered pointcloud_0392.pcd -> pointcloud_0396.pcd | Fitness: 0.9828

Trajectory:

```
x,y,z  
0,0,0  
-0.0043660826205065215,-0.00016952890859987547,5.251785200055065e-06  
-0.009152816027674277,-0.0008648978095404534,2.3424807349124813e-06  
-0.005039400761475491,-0.0025294558970605863,-1.0675960874680467e-06  
-0.007498011966125214,-0.00019508470876557823,3.9504206313727394e-05  
-0.008053494479506379,-0.0008949925877909074,0.00010525243392392223  
-0.012002413938078065,-0.0009755353565196759,3.5120991286818895e-05  
-0.007637615242239352,0.0009430250748799069,0.00010406167463968856  
-0.005914193149483954,0.003458054639122488,0.00015679659398876096  
0.0001193008608031329,0.0025409629149067622,0.0003412504290226131  
0.0017961997647422054,-0.0004255855732969373,0.0004896642586124315  
0.0011862999635071985,-0.00400555599639389,0.0005362236780599271  
0.0009182165493810945,-0.006257040866450254,0.0005657850699405314  
-7.819803541445427e-05,-0.009776177441207267,0.0006465700461654895  
-0.0012029680180015879,-0.012911044760771047,0.000644442987387952
```

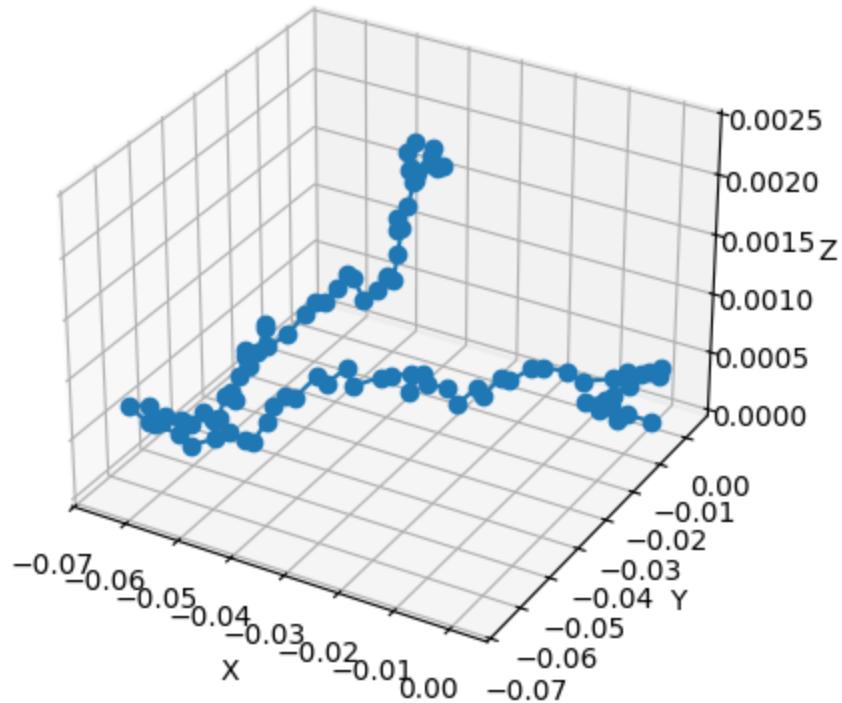
-0.003814010823541326,-0.01845770554664803,0.0007154245187361631
-0.0053948029655769185,-0.020933339298173988,0.0008319856612639093
-0.008429349711718298,-0.024100352609585336,0.0008963012248390961
-0.009881345603323407,-0.02546059022776425,0.0008969908603519333
-0.011931587796477836,-0.028809432126608615,0.0008413385981335652
-0.01297193999198645,-0.03028661631792569,0.0008831662085904139
-0.014285458328856993,-0.03347535245294288,0.0007951007141514124
-0.014964432013264932,-0.034778988227149574,0.0008746869643309652
-0.017139159216700044,-0.03733791757199622,0.0007637296792556474
-0.018479164691301207,-0.03809211690955326,0.0009055274167620322
-0.020619564877452402,-0.04104560483633177,0.0009708387446446787
-0.020862456670043397,-0.04177630896398181,0.001074530713557033
-0.022677266444912657,-0.043569881186280225,0.0009450717128086256
-0.0226095373514193,-0.04264206334845639,0.001074312143693509
-0.02591414593817742,-0.04376471985898808,0.0010288377930601715
-0.026693539494835425,-0.04551443431193998,0.001051514953616136
-0.030278427730626136,-0.04838863465009295,0.0009904268774796247
-0.030900176578996224,-0.04914247916154738,0.001152692271255882
-0.03342658435500966,-0.05107629526109681,0.0010360670924483195
-0.034510712411705044,-0.05232560211698997,0.0011089092575289527
-0.03767125056458774,-0.05434942985540462,0.0009279049042787266
-0.03795281294785317,-0.056852350474993985,0.0009946156597578884
-0.039024333095868245,-0.058957130446384697,0.0009439124005942654

-0.03954430003461621,-0.05992882985743454,0.0008400956628468162
-0.04124420622244345,-0.061210178728521705,0.0006765547937718071
-0.04165641105288191,-0.06288563643830392,0.0007248120504238393
-0.04452443563681413,-0.06337791419236145,0.0007620452715278209
-0.046967470290636165,-0.06384190829774845,0.0006820034020514056
-0.050019417844588056,-0.06599290130609108,0.0006340286590258012
-0.05318663513579704,-0.06456526956177627,0.0006512446818787676
-0.05939918557488429,-0.064133610330228,0.0006528998307570948
-0.06553674456478521,-0.060669743876200666,0.0006277966102540579
-0.06338445487919692,-0.057969044741430324,0.0006042898570725258
-0.06181776118509736,-0.058209210778614844,0.0004696149207912085
-0.060835757682202556,-0.05847748285339164,0.0005077147656849711
-0.06015556303577791,-0.058797501056072184,0.0005545460558230201
-0.05839431337424539,-0.060593372558217826,0.0006456724315926454
-0.05570353228171102,-0.06026325837979646,0.0006297171195764498
-0.055069266398081684,-0.060242461855761216,0.0006200309490380298
-0.053686879149599145,-0.06033326548058963,0.000624301988052966
-0.052399688257393266,-0.058590830496124444,0.0007150116368076724
-0.051560710794059196,-0.05712261157151264,0.0006194085393818942
-0.05086431817928784,-0.05740996523672961,0.0006429244625099639
-0.0489824006869683,-0.0591169234107433,0.0007169500141541872
-0.04839633542660667,-0.05828604846144124,0.0008915729563104534
-0.04814678916047527,-0.05664913707051733,0.0008922285665467027

-0.04693779272207393,-0.0573461857921602,0.0008612826737263824
-0.04569123378176049,-0.05800376887148931,0.0010884857904380986
-0.045411313967878085,-0.05626426486731831,0.0012680856570418749
-0.045470154780877,-0.05646312860866613,0.0012196960916872265
-0.04599077655213694,-0.05415003687163548,0.0010823543508698102
-0.045515362153490786,-0.05271633254432825,0.001164170657658072
-0.044726870522402914,-0.05122716572761566,0.0013433296048011082
-0.045660097292985376,-0.04964334673918672,0.0013286992093026829
-0.04517739660426816,-0.04987988090122164,0.0011554471973217284
-0.04271842759961138,-0.047810295486488914,0.0012410197669138938
-0.04128268284228807,-0.04441446522501568,0.0013564339516343488
-0.04105565474829679,-0.041934581448628716,0.0014143482294145544
-0.03960274330418402,-0.041113631798256174,0.0014151434416066764
-0.037586218613205505,-0.04025006968785214,0.0015286735154784261
-0.03644056591799288,-0.039427838088668364,0.0016370535458498899
-0.03555141193948189,-0.03876623766430633,0.0016065438141000545
-0.03364907403464296,-0.03866222207247682,0.0014512251425556073
-0.03181398867119866,-0.03730645920622639,0.0015215885455604561
-0.030574223145178395,-0.036690489584590744,0.0016368476533528615
-0.03076867353927092,-0.036004757375339896,0.0016030138605045683
-0.02972575774249185,-0.03590400780719109,0.001596111463276597
-0.029553123135695428,-0.03472800323482742,0.001794358484670498
-0.028979058569657195,-0.03436201374374597,0.0020098729591888583

-0.029178589977891054,-0.03541378340081337,0.0021085584574104858
-0.029405448804045935,-0.0354363741865469,0.0020098594542242733
-0.028673013791999408,-0.03298495055897715,0.002160695135997903
-0.028584518991962283,-0.030552914101489585,0.002355254942130736
-0.029171752535482917,-0.029916266306177407,0.002292465711150108
-0.03118693825808143,-0.028230428495087747,0.002320921576030658
-0.0322177599952533,-0.026828669479891556,0.0024253007259555817
-0.031826794200093614,-0.025273429055791604,0.002476474543949314
-0.03257740783446562,-0.025061496641423174,0.002255885074259409
-0.032425159225295826,-0.024611252263183468,0.0021312529441158898
-0.03284706788482216,-0.022141120045027496,0.0021616229090021728
-0.03213055245804159,-0.0184433630878969,0.0023028524099868823
-0.033077923746051084,-0.017239388791452478,0.002194656710041377
-0.03344698504363322,-0.01693033319258584,0.0021121490097124
-0.032889276567064724,-0.015697550767446788,0.0020626350888426235
-0.0324399616956926,-0.014267522963597292,0.0020549767501461765

Estimated 3D Trajectory of TurtleBot



References:

3.

1. **OpenCV Camera Calibration Tutorial**
[OpenCV: Camera Calibration](#)
2. **OpenCV `cv2.findChessboardCorners` Documentation**
[Camera Calibration and 3D Reconstruction](#)
3. **OpenCV `cv2.cornerSubPix` Documentation**
[Detecting corners location in subpixels](#)
4. **OpenCV `cv2.calibrateCamera` Documentation**
[Camera Calibration and 3D Reconstruction](#)
5. **LearnOpenCV: Camera Calibration Using OpenCV**
<https://learnopencv.com/camera-calibration-using-opencv/>
6. **BoofCV Tutorial on Camera Calibration**
[Tutorial Camera Calibration - BoofCV](#)
7. **OpenCV API: `cv2.projectPoints` for Reprojection**
[Camera Calibration and 3D Reconstruction](#)
8. **OpenCvSharp `Cv2.CalibrateCamera` Documentation**
[Cv2.Undistort Method](#)

4.

1. SIFT Keypoint Detection and Description
[OpenCV: Introduction to SIFT \(Scale-Invariant Feature Transform\)](#)
2. BruteForce and FLANN Feature Matching
[OpenCV: Feature Matching](#)

[cv::FlannBasedMatcher Class Reference](#)

3. Homography Estimation with RANSAC

[Feature Matching + Homography to find Objects](#)

[RANSAC | LearnOpenCV](#)

4. Perspective Warping and Stitching

[cv2.warpPerspective\(\) | TheAILearner](#)

5. Image Clustering using KMeans (for multi-stitching)

[2.3. Clustering — scikit-learn 1.6.1 documentation](#)

- 5.

1. Open3D ICP Registration Tutorial

https://www.open3d.org/docs/release/tutorial/pipelines/icp_registration.html

2. Scipy `ortho_group` Documentation

https://docs.scipy.org/doc/scipy/reference/generated/scipy.stats.ortho_group.html

3. Open3D Normal Estimation (v0.7.0)

https://www.open3d.org/docs/0.7.0/python_api/open3d.geometry.estimate_normals.html